



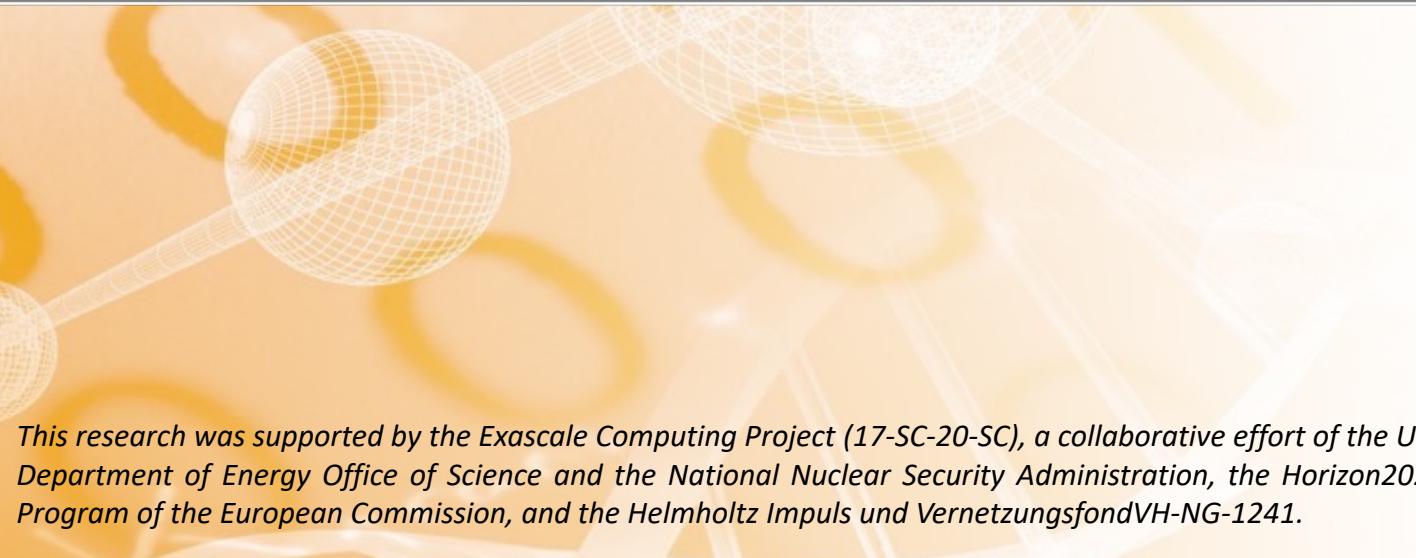
THE UNIVERSITY OF
TENNESSEE
KNOXVILLE



Lossy Compression and Mixed Precision Strategies for Memory-Bound Linear Algebra

PPAM 2022
Gdansk, Poland

Hartwig Anzt
Innovative Computing Lab, University of Tennessee



This research was supported by the Exascale Computing Project (17-SC-20-SC), a collaborative effort of the U.S. Department of Energy Office of Science and the National Nuclear Security Administration, the Horizon2020 Program of the European Commission, and the Helmholtz Impuls und Vernetzungsfond VH-NG-1241.



Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) **NVIDIA V100 GPU**

Double precision GMRES

Initial residual norm **Relative residual $\sim 10^{-12}$**

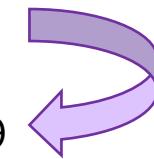
$\text{sqrt}(r^T r)$: 9670.36

Final residual norm

$\text{sqrt}(r^T r)$: 9.6639e-09

GMRES iteration count: 23271

GMRES execution time: 43801 ms



Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) **NVIDIA V100 GPU**

Double precision GMRES

Initial residual norm **Relative residual $\sim 10^{-12}$**

$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 9.6639e-09$

GMRES iteration count: 23271

GMRES execution time: 43801 ms

Single precision GMRES

Initial residual norm **Relative residual $\sim 10^{-7}$**

$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 0.00175464$

GMRES iteration count: 25000

GMRES execution time: 27376 ms

Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) NVIDIA V100 GPU

Double precision GMRES

Initial residual norm Relative residual $\sim 10^{-12}$

$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 9.6639e-09$

GMRES iteration count: 23271

GMRES execution time: 43801 ms

Single precision GMRES

Initial residual norm Relative residual $\sim 10^{-7}$

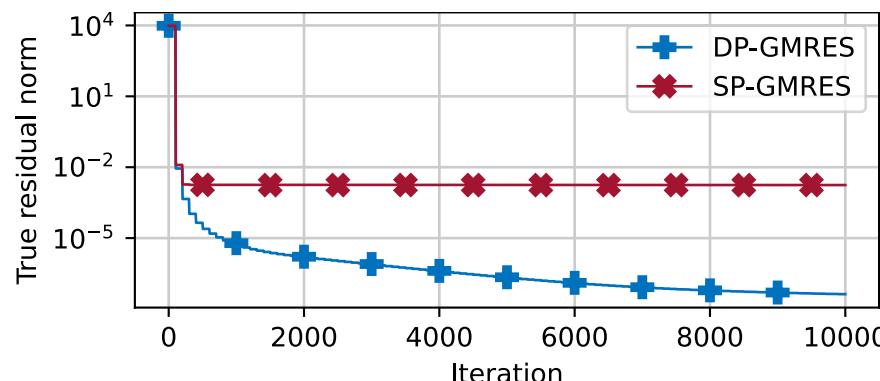
$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 0.00175464$

GMRES iteration count: 25000

GMRES execution time: 27376 ms



Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) NVIDIA V100 GPU

Double precision GMRES

Initial residual norm Relative residual $\sim 10^{-12}$

$\text{sqrt}(r^T r)$: 9670.36

Final residual norm

$\text{sqrt}(r^T r)$: $9.6639e-09$

GMRES iteration count: 23271

GMRES execution time: 43801 ms

Single precision GMRES

Initial residual norm Relative residual $\sim 10^{-7}$

$\text{sqrt}(r^T r)$: 9670.36

Final residual norm

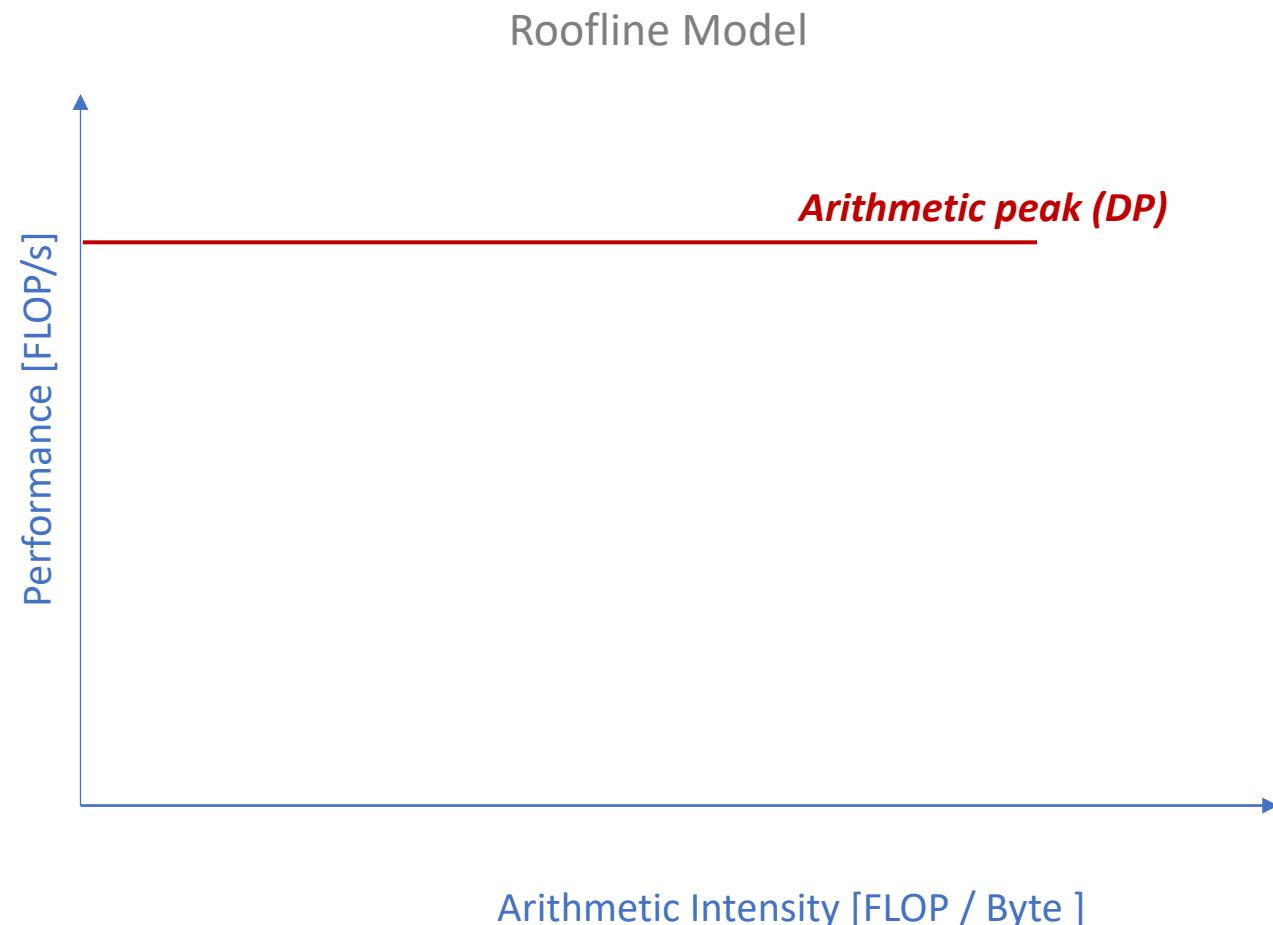
$\text{sqrt}(r^T r)$: 0.00175464

GMRES iteration count: 25000

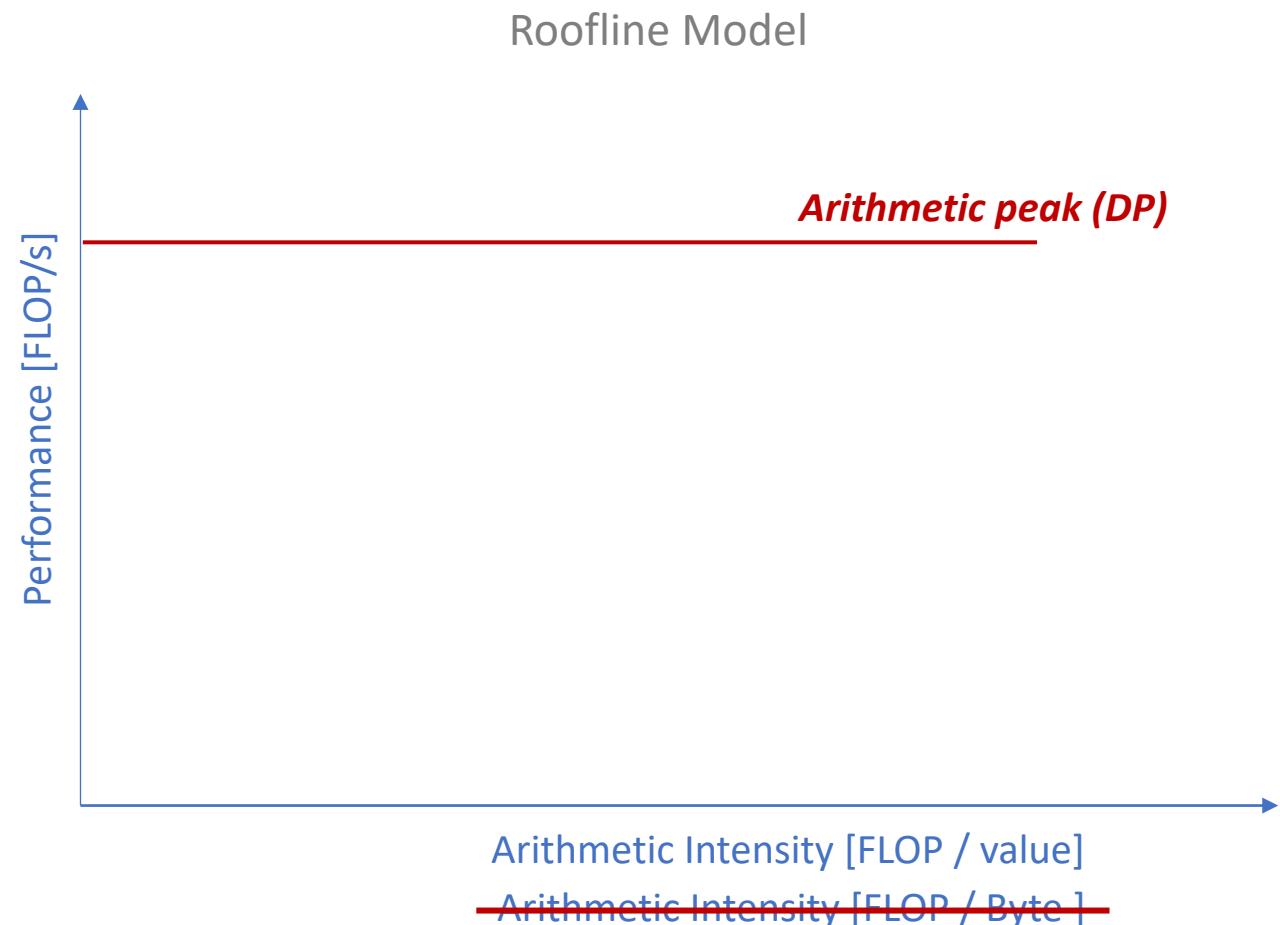
GMRES execution time: 27376 ms

~2x faster!

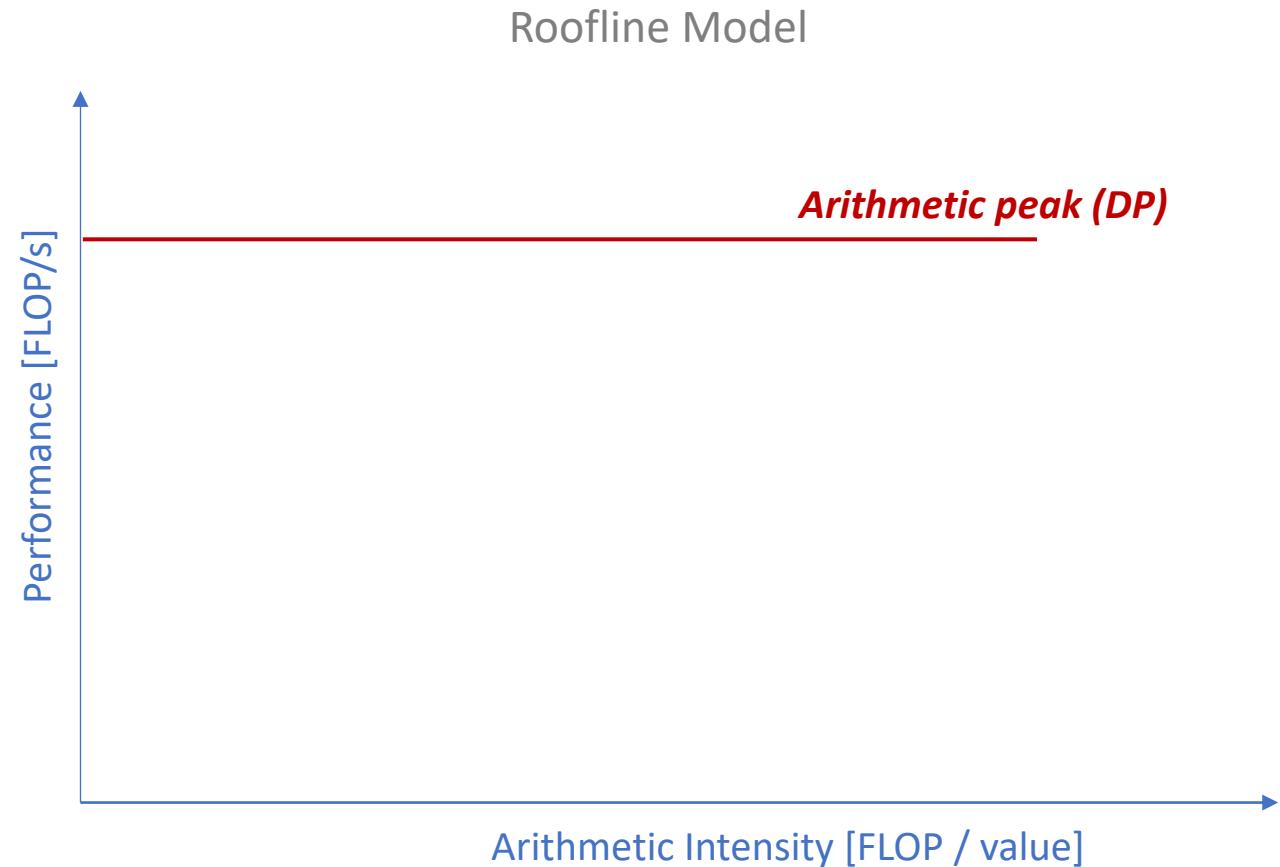
Why are we faster with a single precision algorithm?



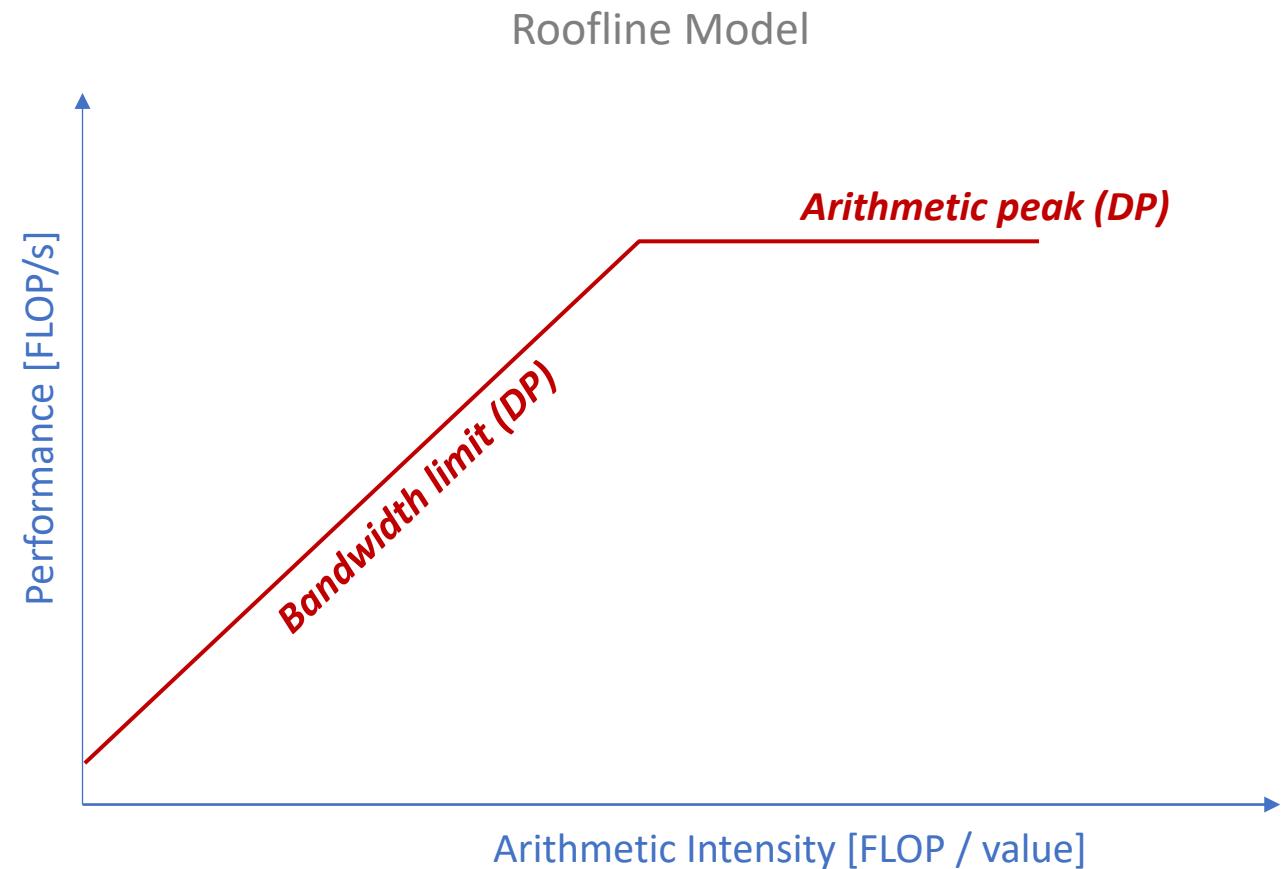
Why are we faster with a single precision algorithm?



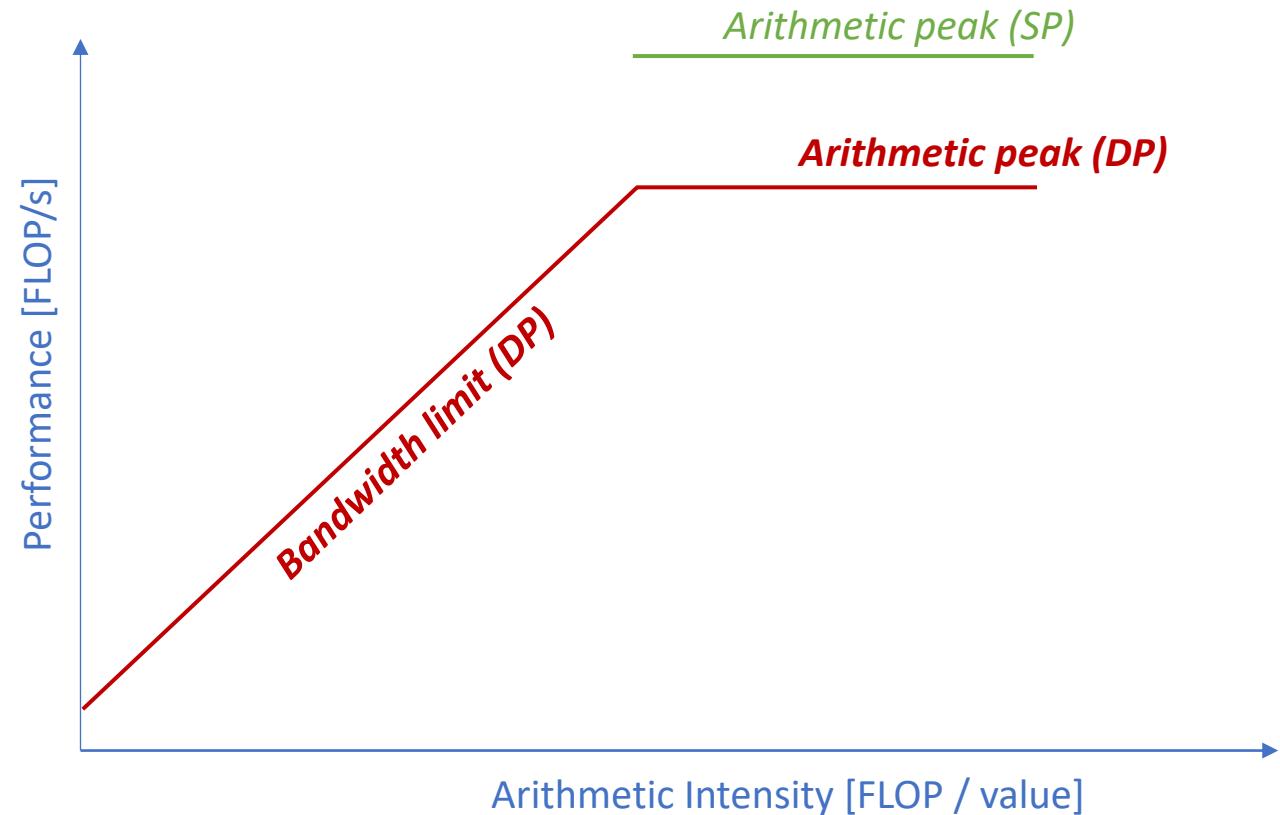
Why are we faster with a single precision algorithm?



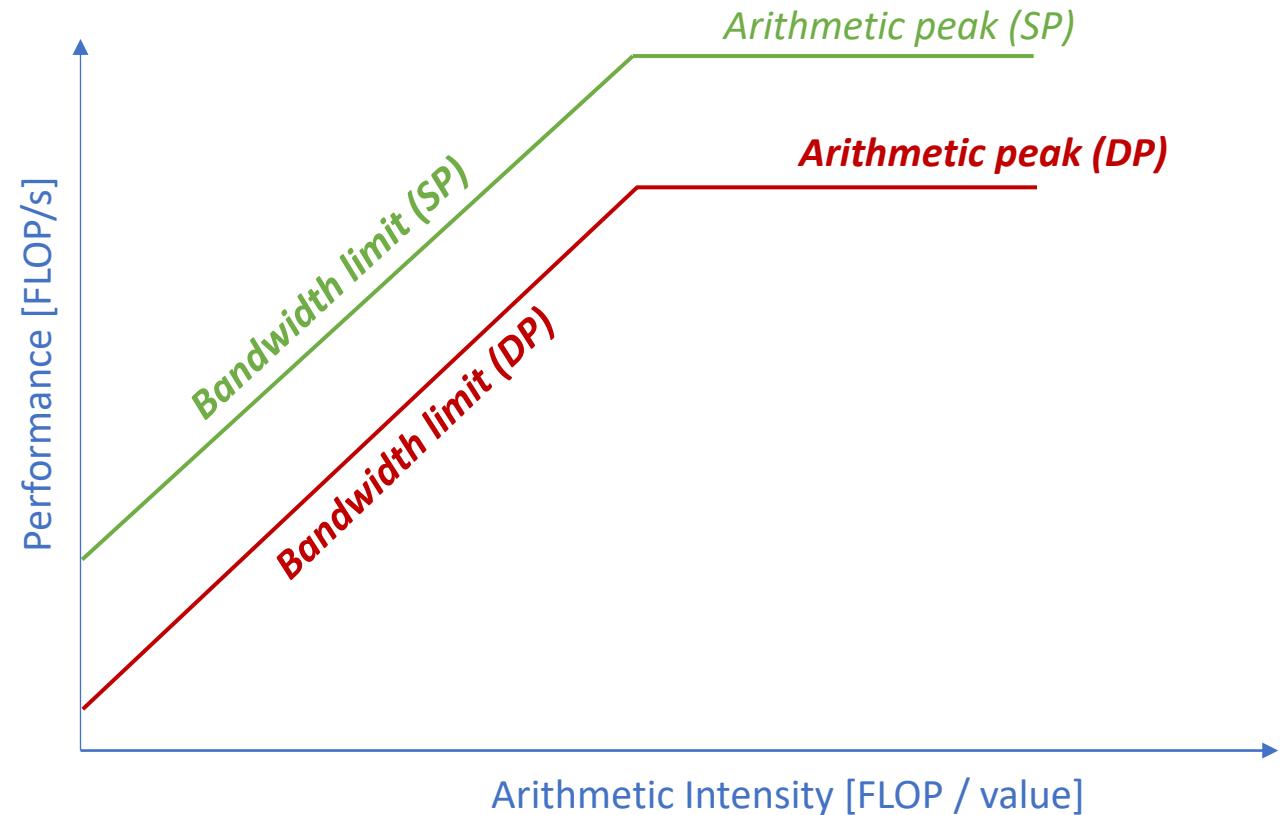
Why are we faster with a single precision algorithm?



Why are we faster with a single precision algorithm?



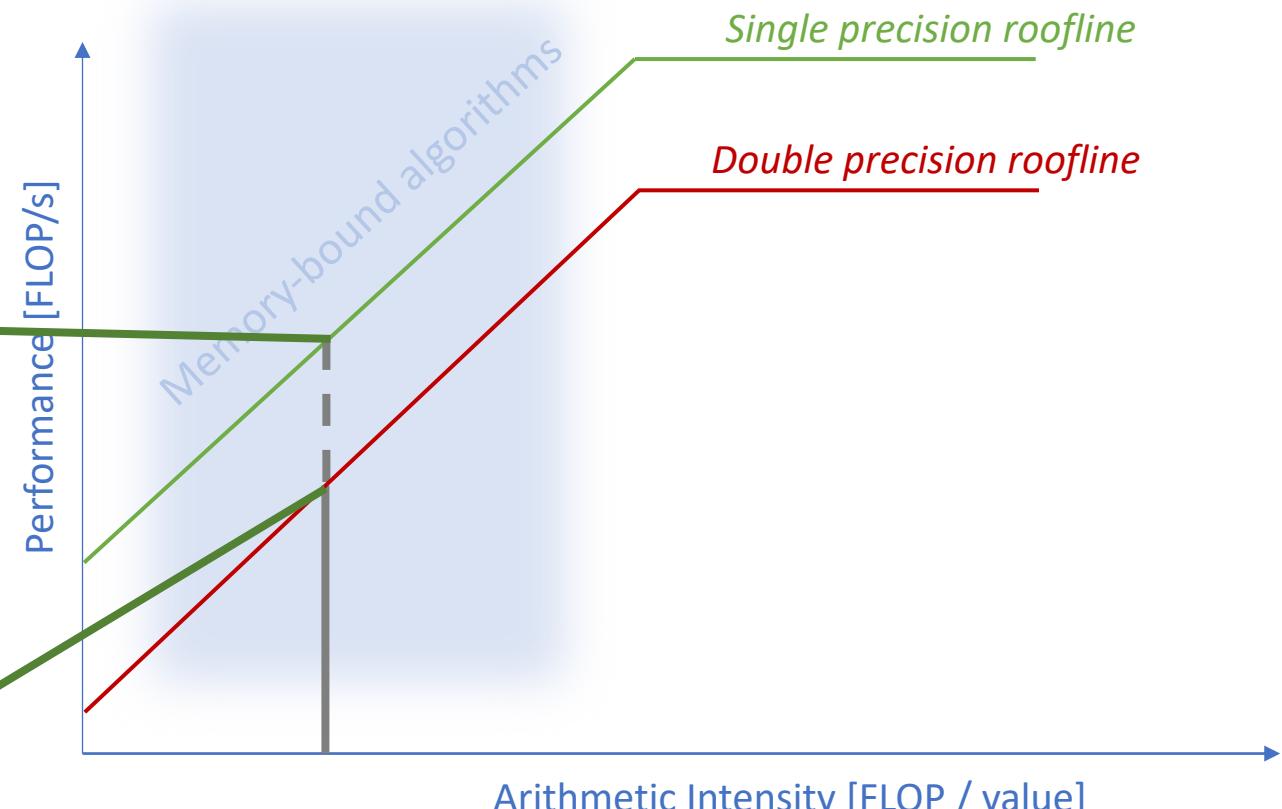
Why are we faster with a single precision algorithm?



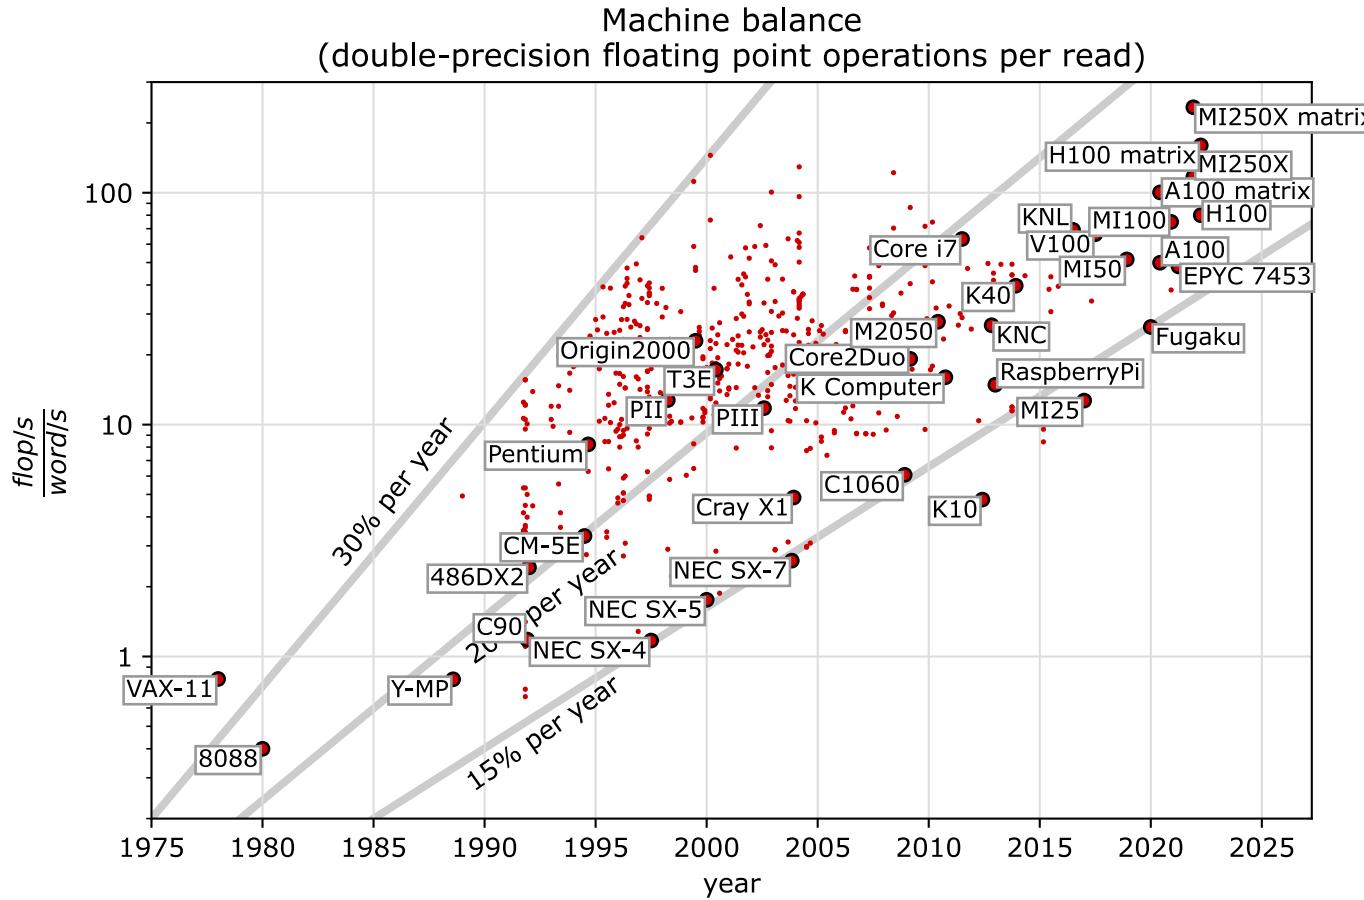
Why are we faster with a single precision algorithm?

Single precision GMRES
Initial residual norm
 $\text{sqrt}(r^T r)$: 9670.36
Final residual norm
 $\text{sqrt}(r^T r)$: 0.00175464
GMRES iteration count: 25000
GMRES execution time: 27376 ms

Double precision GMRES
Initial residual norm
 $\text{sqrt}(r^T r)$: 9670.36
Final residual norm
 $\text{sqrt}(r^T r)$: 9.6639e-09
GMRES iteration count: 23271
GMRES execution time: 43801 ms



The memory wall



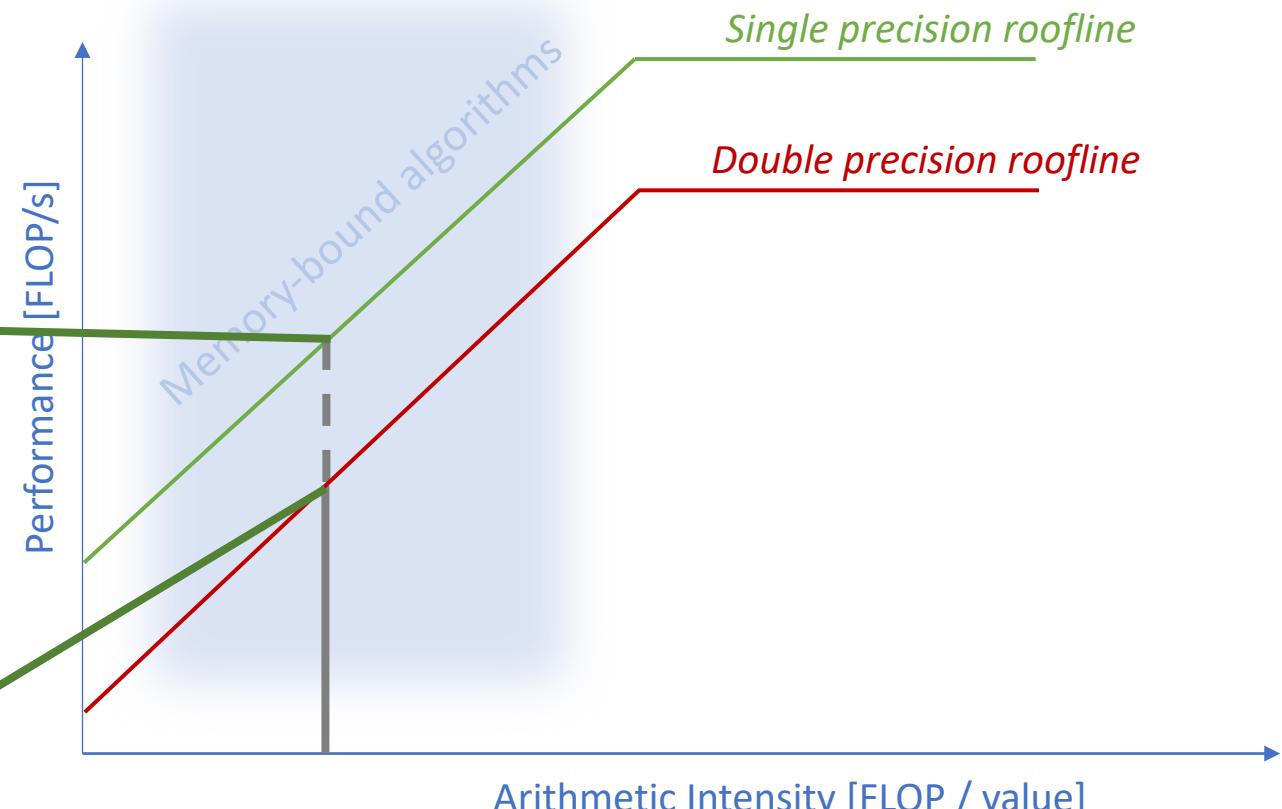
M. Gates

Compute performance grows faster than memory bandwidth.

Why are we faster with a single precision algorithm?

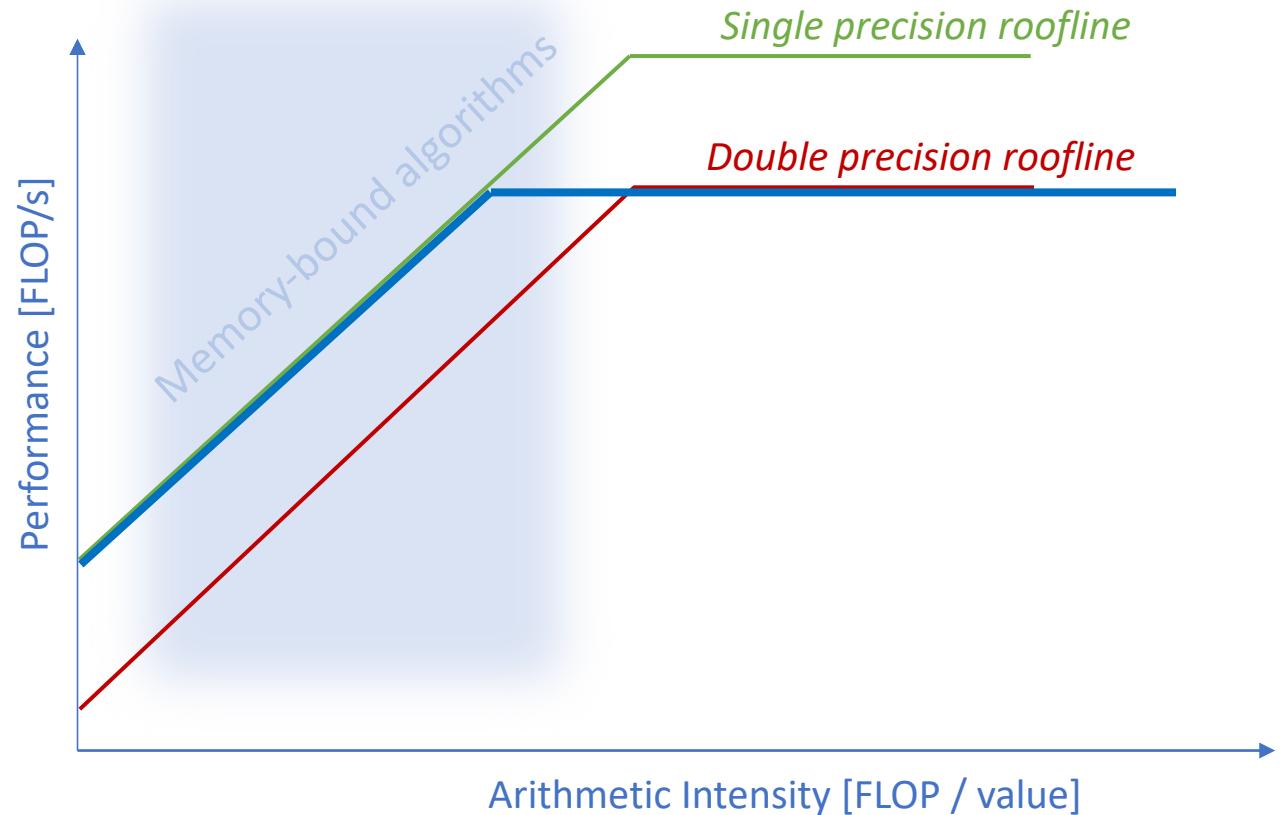
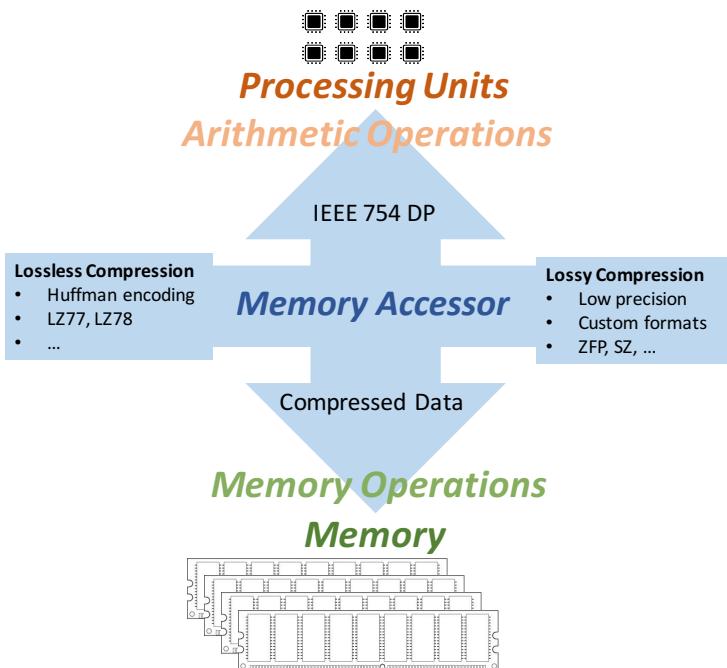
Single precision GMRES
Initial residual norm
 $\text{sqrt}(r^T r)$: 9670.36
Final residual norm
 $\text{sqrt}(r^T r)$: 0.00175464
GMRES iteration count: 25000
GMRES execution time: 27376 ms

Double precision GMRES
Initial residual norm
 $\text{sqrt}(r^T r)$: 9670.36
Final residual norm
 $\text{sqrt}(r^T r)$: 9.6639e-09
GMRES iteration count: 23271
GMRES execution time: 43801 ms

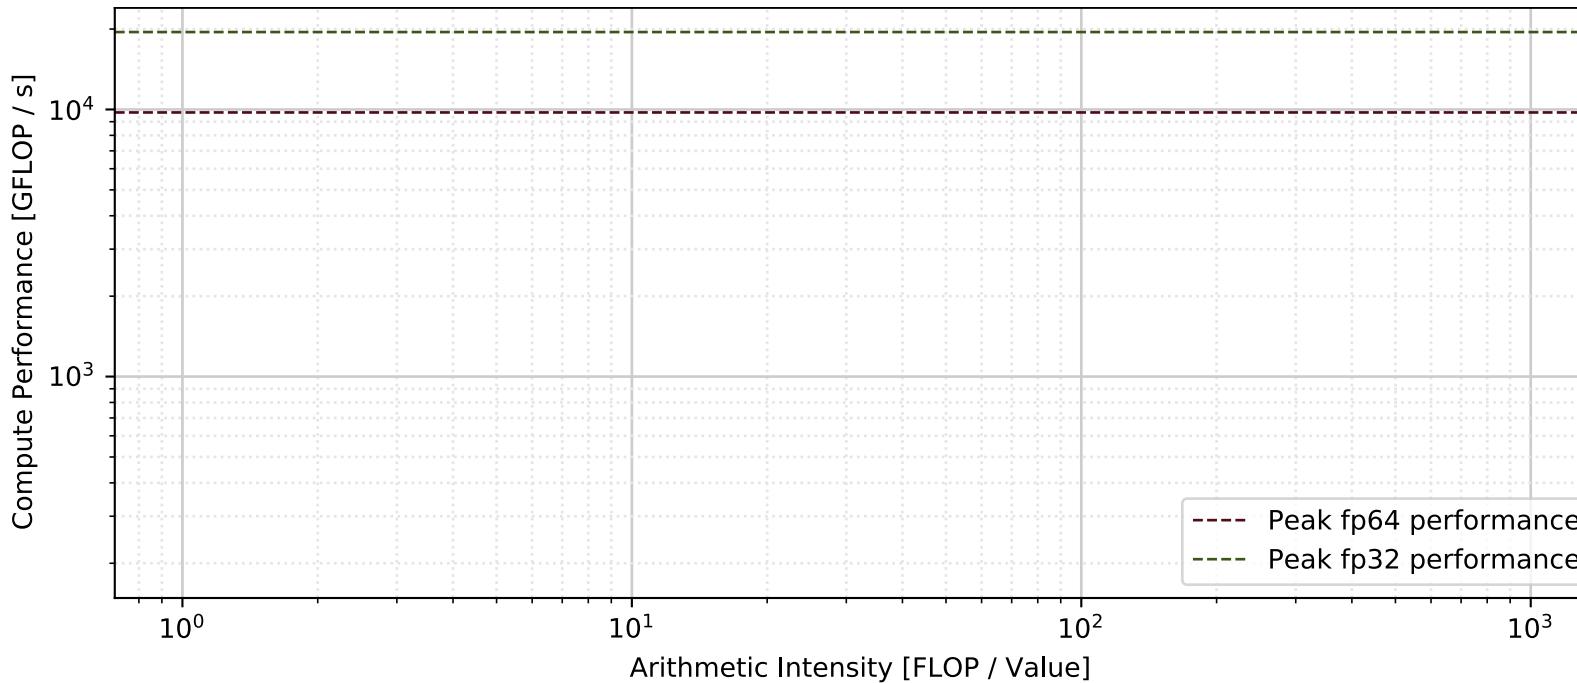


Can we get the best of both worlds?

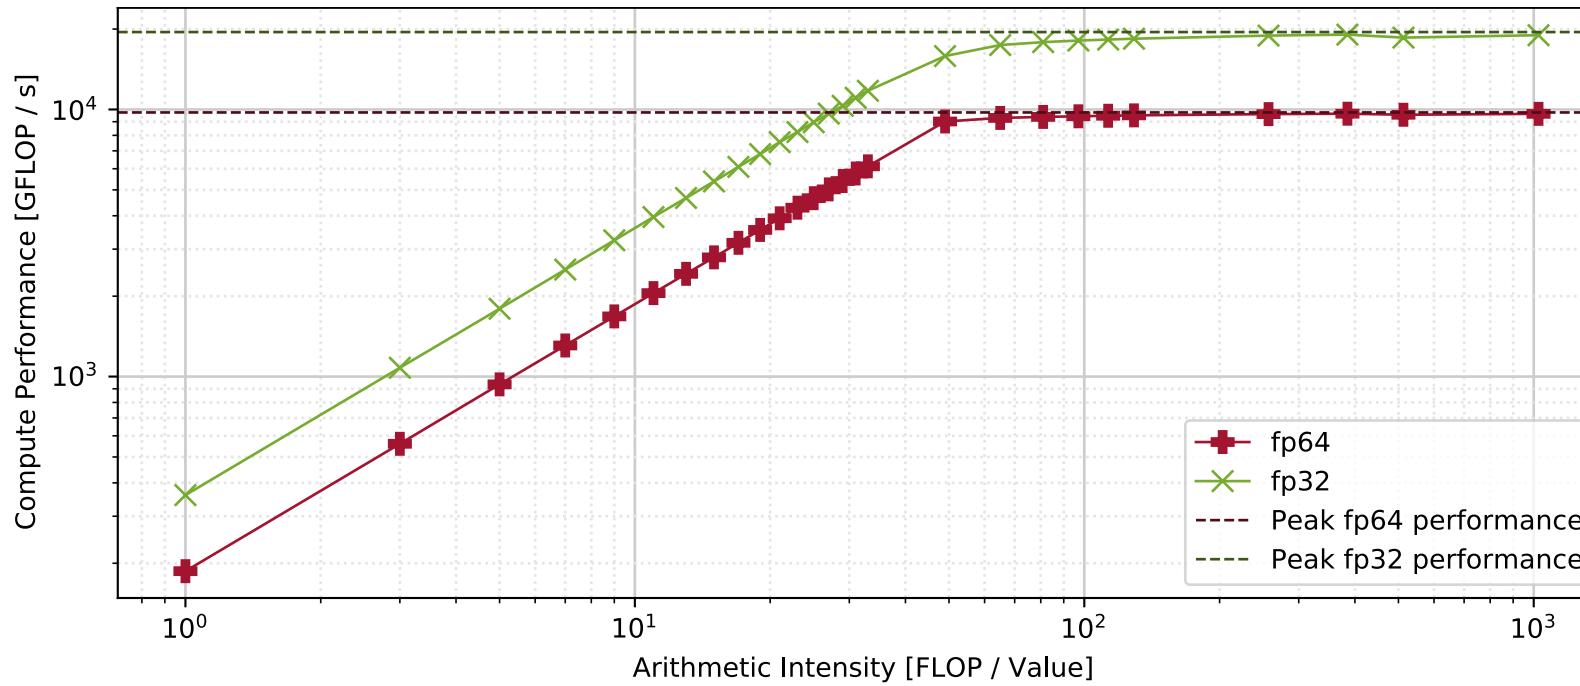
- For memory-bound algorithms, the **arithmetic operations** are free, can use **high precision formats**.
- **Data access** should be as cheap as possible, use **reduced precision**.
- **In-Register Compression**



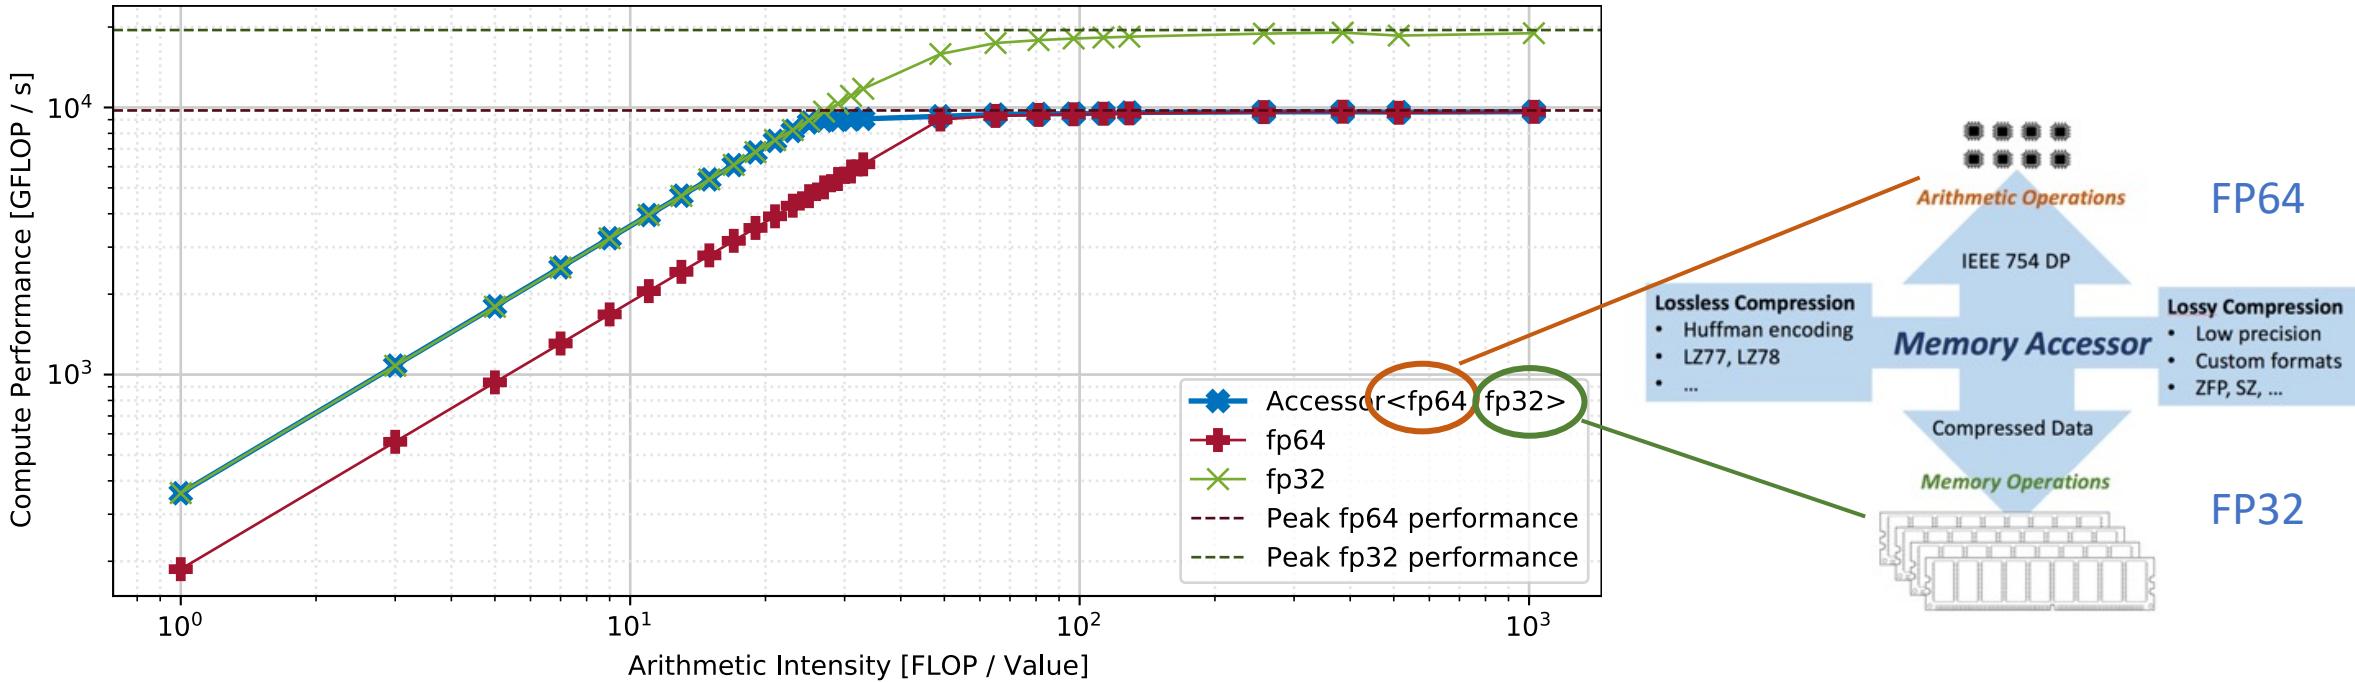
Memory Accessor for NVIDIA A100 GPU



Memory Accessor for NVIDIA A100 GPU



Memory Accessor for NVIDIA A100 GPU



T. Grützmacher

Use the memory accessor to boost performance

- *Start from double precision algorithm*
- *Use memory accessor to store intermediate data in compressed form*
- *Require double precision output accuracy*

Use the memory accessor to boost performance

- Start from double precision algorithm
- Use memory accessor to store intermediate data in compressed form
- Require double precision output accuracy
- Preconditioning

[1] G Flegar, H Anzt, T Cojean, ES Quintana-Ortí, "Adaptive precision block-Jacobi for high performance preconditioning in the Ginkgo linear algebra software," ACM Transactions on Mathematical Software (TOMS) 47 (2), 1-28.

[2] F Göbel, T Grütmacher, T Ribizel, H Anzt, "Mixed precision incomplete and factorized sparse approximate inverse preconditioning on GPUs," European Conference on Parallel Processing, 550-564.

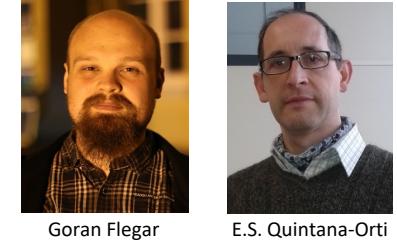
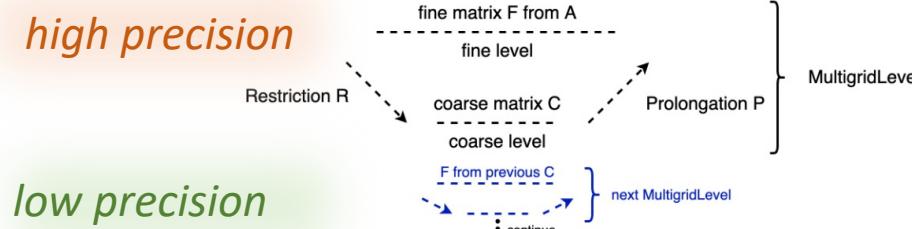
- Multigrid

[3] Stephen F. McCormick, Joseph Benzaken, Rasmus Tamstorf "Algebraic error analysis for mixed-precision multigrid solvers", <https://arxiv.org/abs/2007.06614>

[4] M Tsai, N Beams, H Anzt, "Mixed precision algebraic multigrid on GPUs," PPAM 2022.

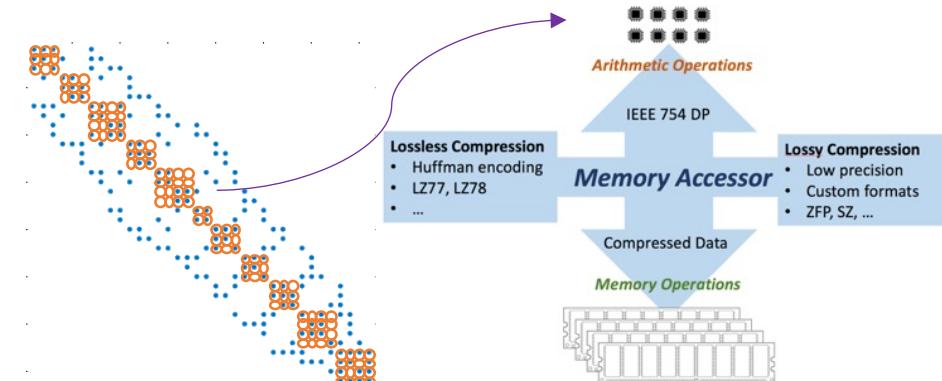
- Krylov solvers

[5] J Aliaga, H Anzt, T Grütmacher, E S Quintana-Ortí, A Tomás, "Compressed basis GMRES on high-performance graphics processing units", IJHPCA 2022



Goran Flegar

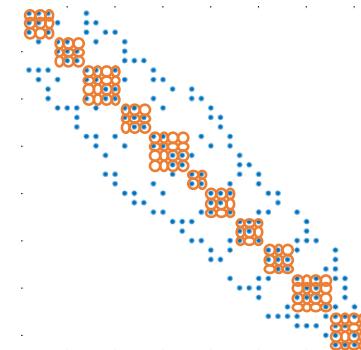
E.S. Quintana-Ortí



Mike Tsai

Using the memory accessor for mixed precision preconditioning

- **Preconditioning iterative solvers.**
 - Idea: Approximate inverse of system matrix to make the system “easier to solve”: $P^{-1} \approx A^{-1}$ and solve $Ax = b \Leftrightarrow P^{-1}Ax = P^{-1}b \Leftrightarrow \tilde{A}x = \tilde{b}$.
- **Block-Jacobi preconditioner** is based on **block-diagonal scaling**: $P = diag_B(A)$
 - Each block corresponds to one (small) linear system.
 - *Larger* blocks typically improve convergence.
 - *Larger* blocks make block-Jacobi more expensive.



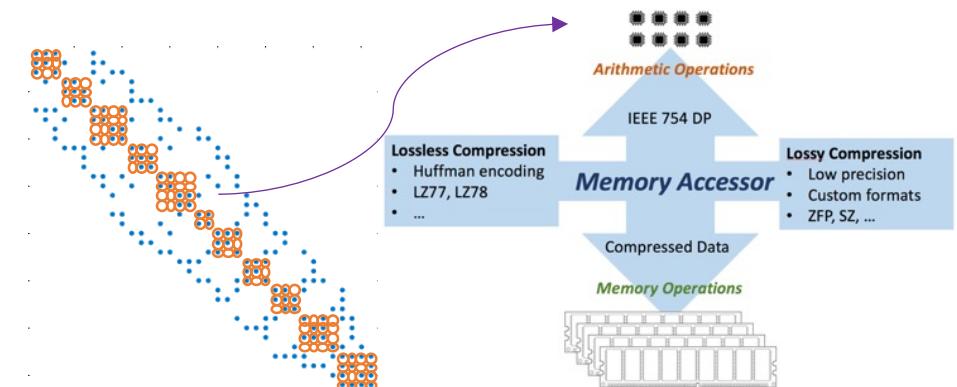
Using the memory accessor for mixed precision preconditioning

- Preconditioning iterative solvers.
 - Idea: Approximate inverse of system matrix to make the system “easier to solve”: $P^{-1} \approx A^{-1}$ and solve $Ax = b \Leftrightarrow P^{-1}Ax = P^{-1}b \Leftrightarrow \tilde{A}x = \tilde{b}$.
- Block-Jacobi preconditioner is based on block-diagonal scaling: $P = \text{diag}_B(A)$
 - Each block corresponds to one (small) linear system.
 - Larger blocks typically improve convergence.
 - Larger blocks make block-Jacobi more expensive.
- Why should we store the preconditioner matrix P^{-1} in full (high) precision?
- Use the accessor to store the inverted diagonal blocks in lower precision.
 - Be careful to preserve the regularity of each inverted diagonal block!



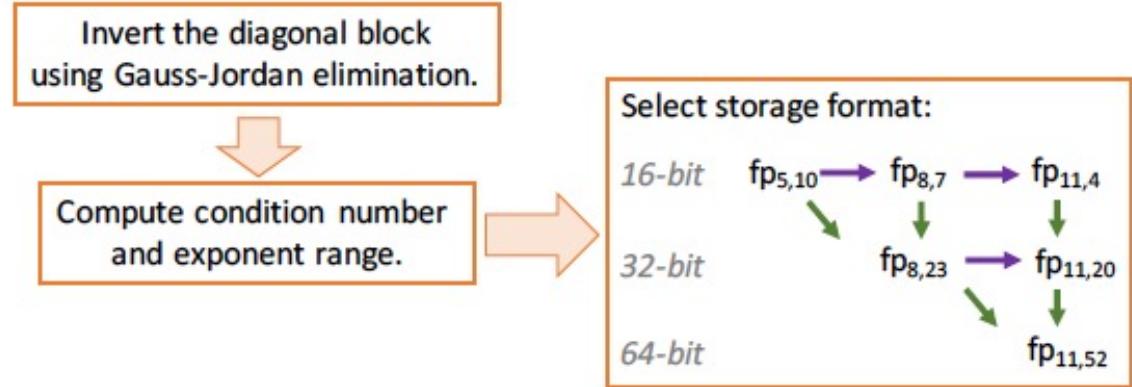
Goran Flegar

E.S. Quintana-Orti



Using the memory accessor for mixed precision preconditioning

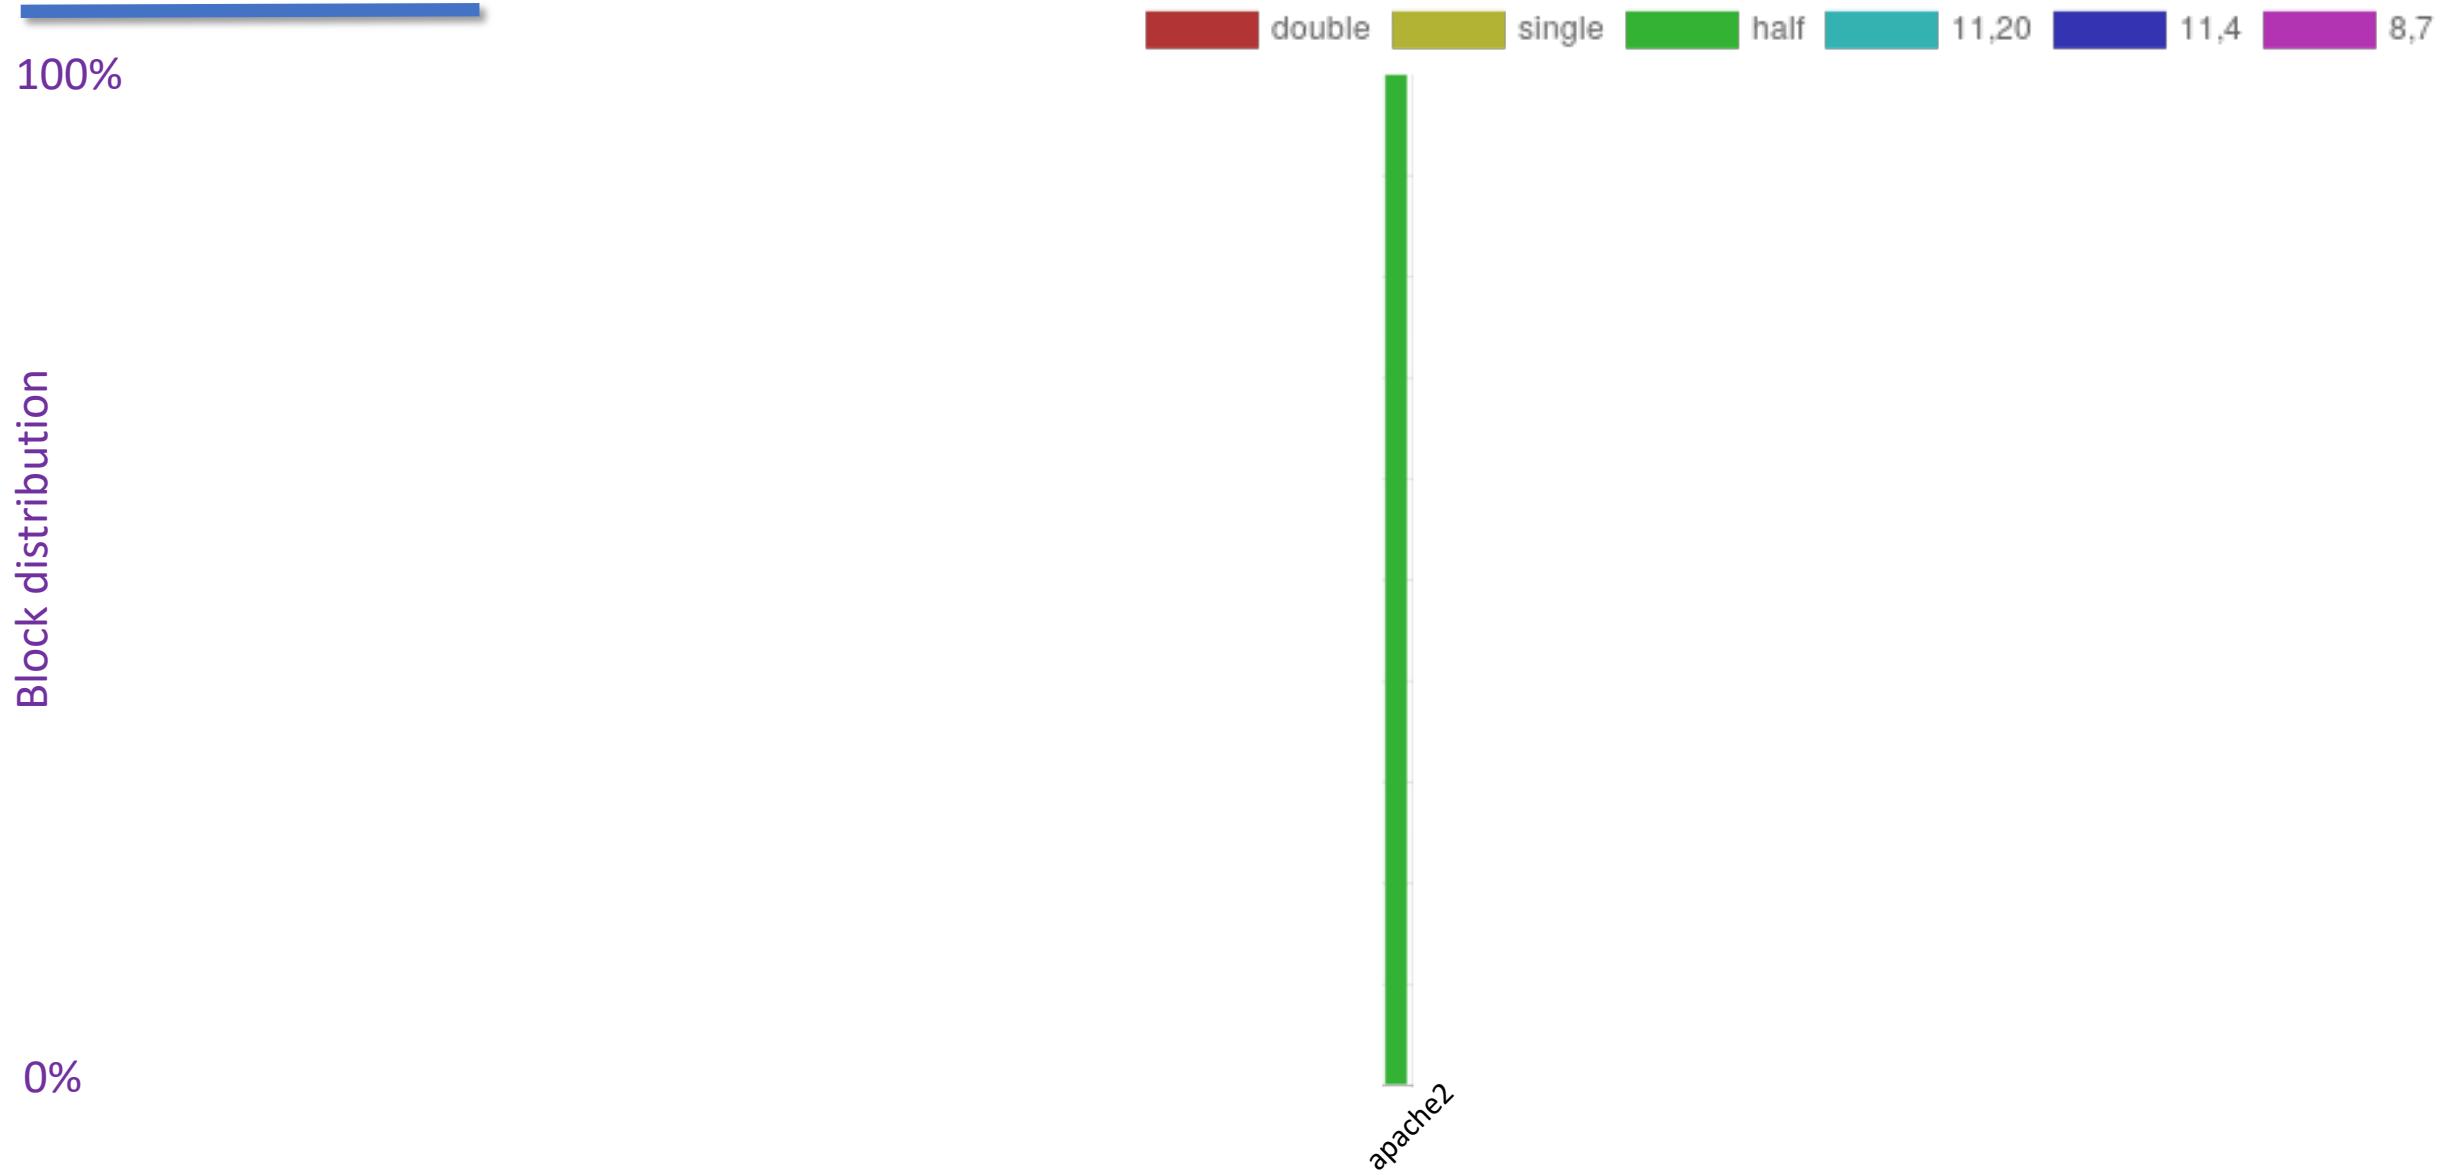
- Choose how much accuracy of the preconditioner should be preserved in the selection of the storage format.
- All computations use double precision, but store blocks in lower precision.



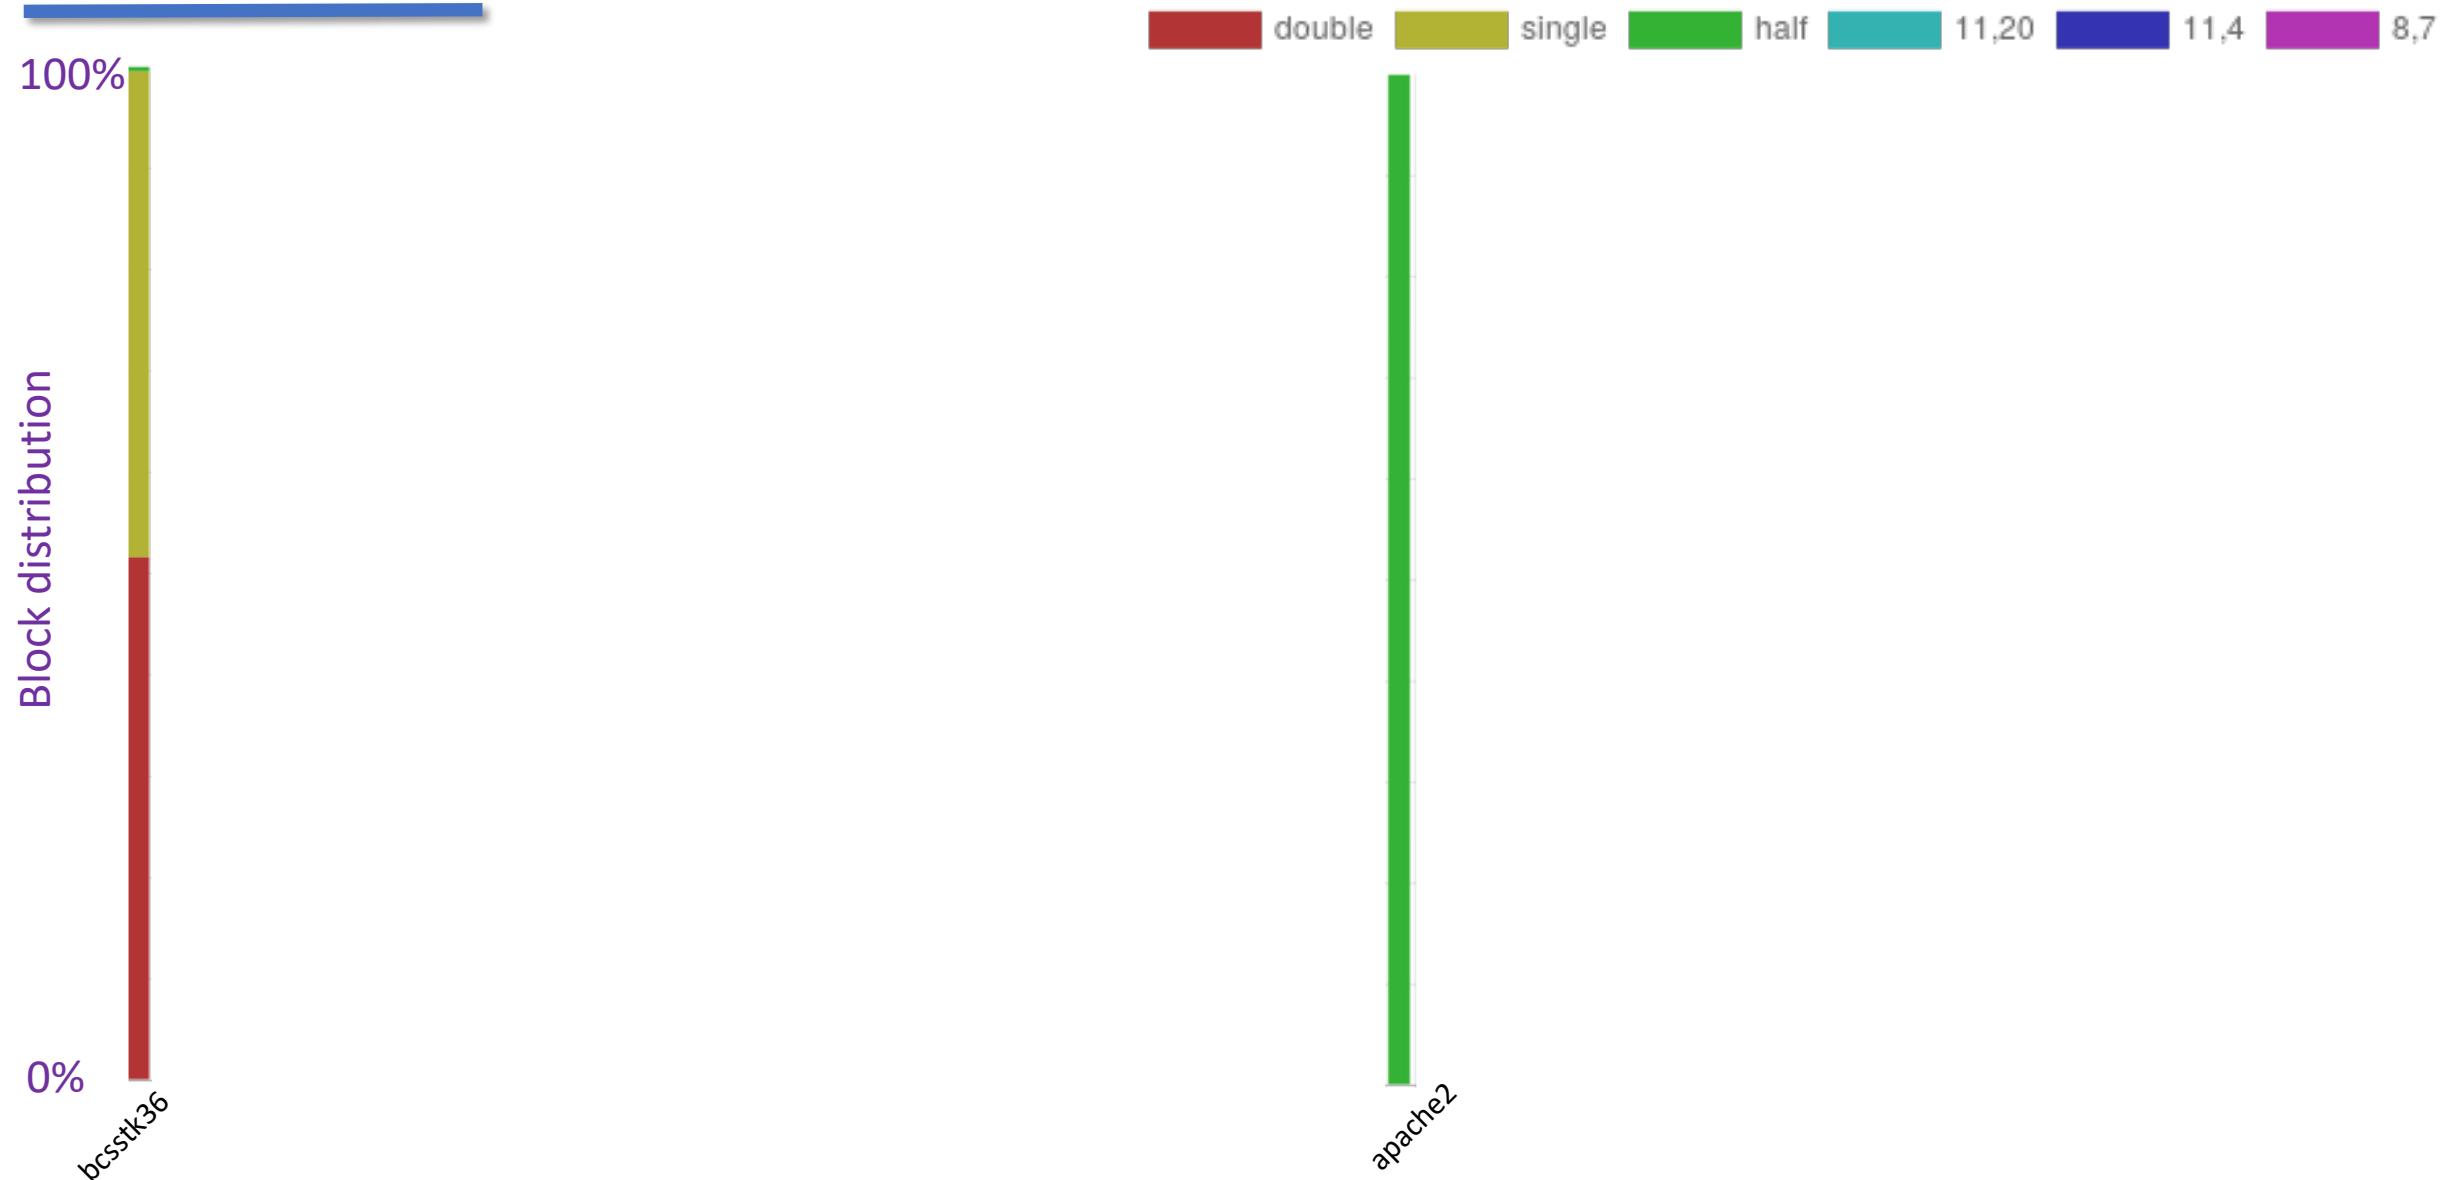
- + Regularity preserved;
- + Flexibility in the accuracy;
- + "Not a low precision preconditioner"
 - + Preconditioner is a constant operator;
 - + No flexible Krylov solver needed ;

- Overhead of the precision detection
(condition number calculation);
- Overhead from storing precision information
(need to additionally store/retrieve flag);
- Speedups / preconditioner quality problem-dependent;

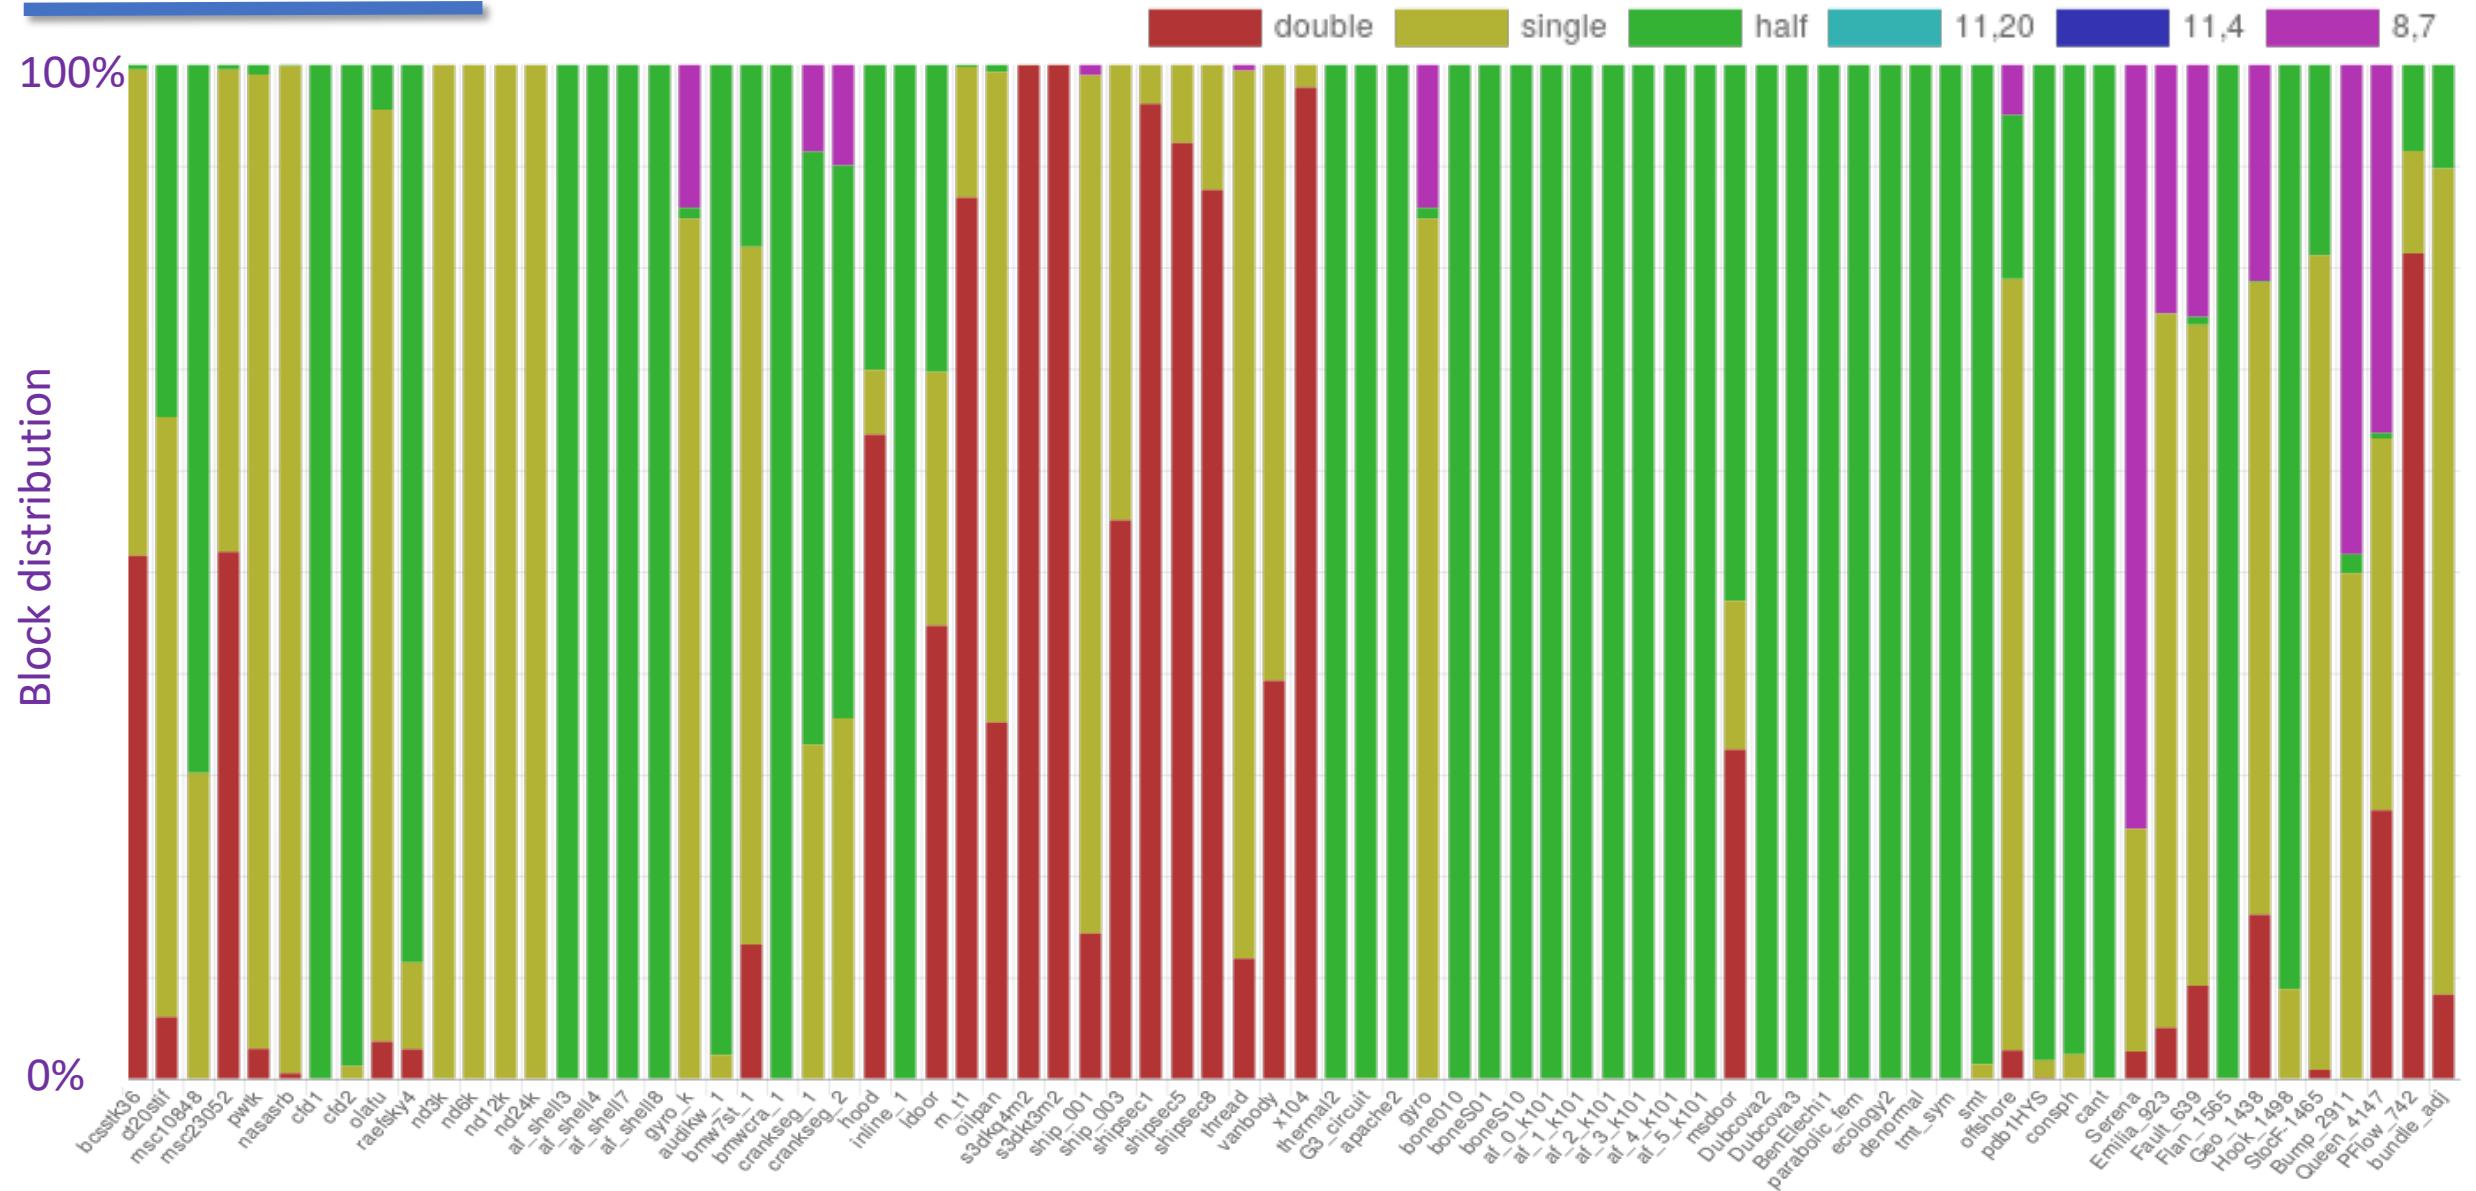
Using the memory accessor for mixed precision preconditioning



Using the memory accessor for mixed precision preconditioning



Using the memory accessor for mixed precision preconditioning



Using the memory accessor for mixed precision preconditioning



Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$ (apache2 from SuiteSparse) NVIDIA A100 GPU

Double Precision CG + Double Precision Preconditioner

```
Initial residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
1390.67  
Final residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
3.97985e-06 Accuracy improvement ~109  
CG iteration count: 4797  
CG execution time [ms]: 2971.18
```

Single Precision CG + Single Precision Preconditioner

```
Initial residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
1390.67  
Final residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
1588.77 No improvement  
CG iteration count: 8887  
CG execution time [ms]: 2972.46
```

Experiments based on the Ginkgo library <https://ginkgo-project.github.io/>

[ginkgo/examples/adaptiveprecision-blockjacobi/adaptiveprecision-blockjacobi.cpp](https://github.com/GinkgoProject/ginkgo/blob/main/examples/adaptiveprecision-blockjacobi/adaptiveprecision-blockjacobi.cpp)

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$ (apache2 from SuiteSparse) NVIDIA A100 GPU

Double Precision CG + Double Precision Preconditioner

```
Initial residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
1390.67
Final residual norm sqrt(r^T r):
%%MatrixMarket matrix array real general
1 1
3.97985e-06
CG iteration count:    4797
CG execution time [ms]: 2971.18
```

Double Precision CG + Mixed Precision Preconditioner

- *Preconditioner remains a constant operator*
- *Attainable accuracy of CG unaffected*
- *Faster because of less data movement*

Experiments based on the Ginkgo library <https://ginkgo-project.github.io/>

ginkgo/examples/adaptiveprecision-blockjacobi/adaptiveprecision-blockjacobi.cpp

Using the memory accessor for mixed precision preconditioning



Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$ (apache2 from SuiteSparse) NVIDIA A100 GPU

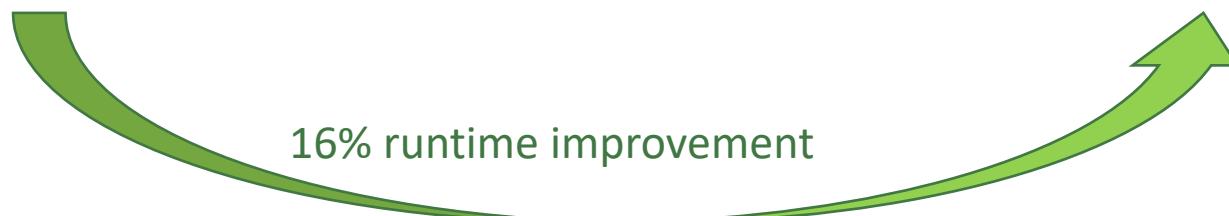
Double Precision CG + Double Precision Preconditioner

```
Initial residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
1390.67  
Final residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
3.97985e-06  
CG iteration count:  
CG execution time [ms]: 4797 2971.18
```

Double Precision CG + Mixed Precision Preconditioner

```
Initial residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
1390.67  
Final residual norm sqrt(r^T r):  
%%MatrixMarket matrix array real general  
1 1  
3.98574e-06  
CG iteration count:  
CG execution time [ms]: 4794 2568.1
```

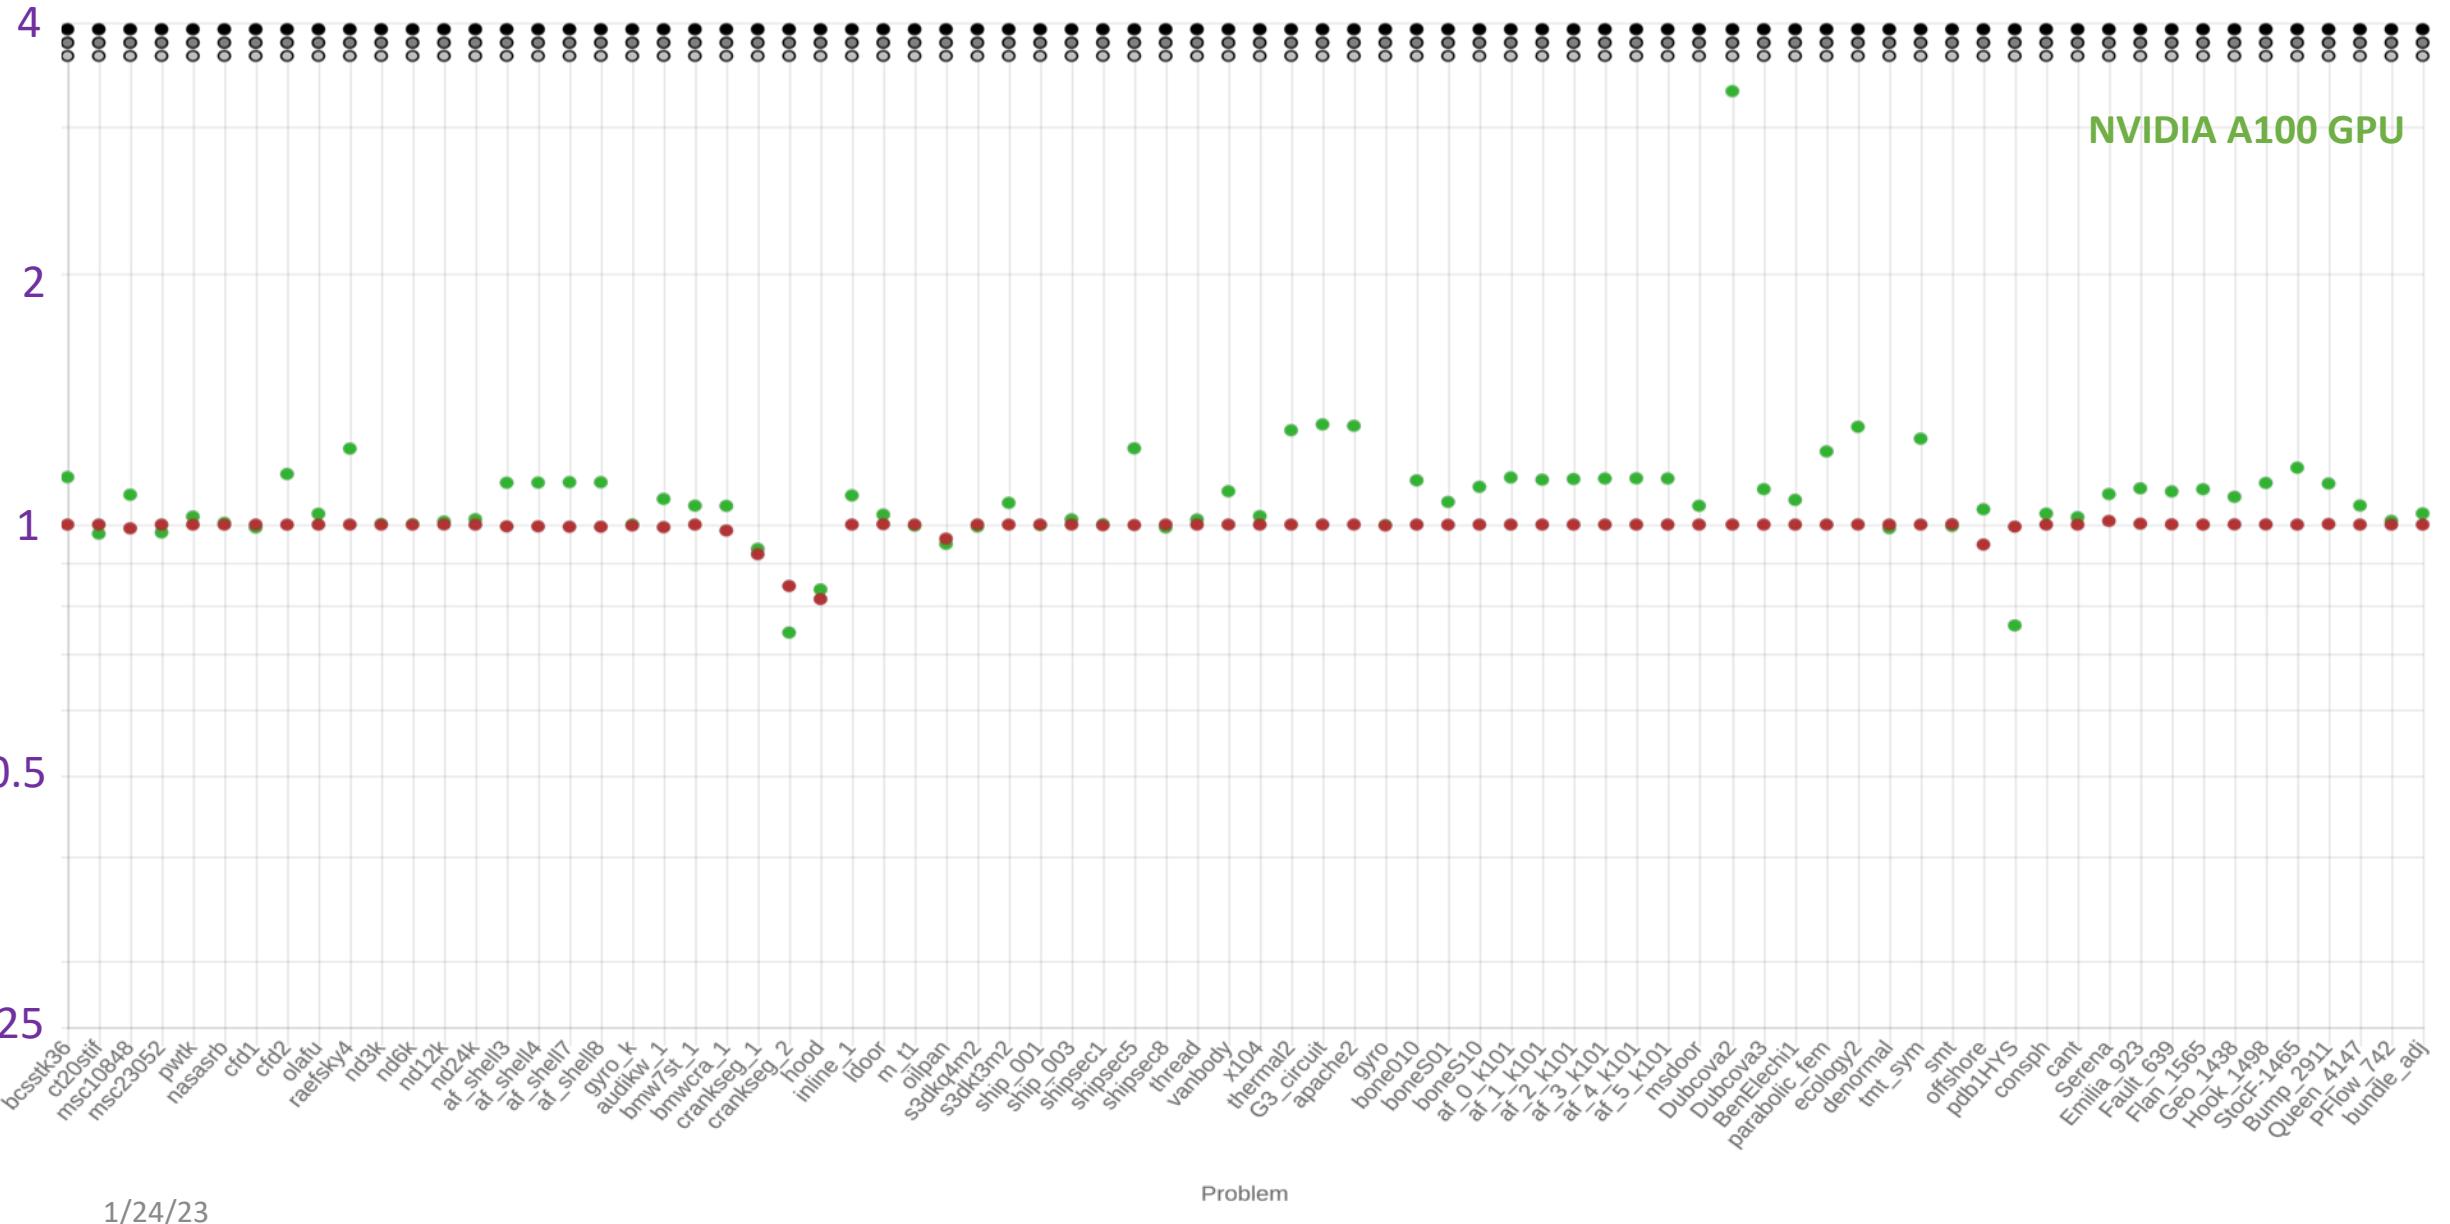
16% runtime improvement



Experiments based on the Ginkgo library <https://ginkgo-project.github.io/>

ginkgo/examples/adaptiveprecision-blockjacobi/adaptiveprecision-blockjacobi.cpp

Iterations (adaptive) Time (adaptive) CG converged? CG + Jacobi converged? CG + adaptive Jacobi converged?



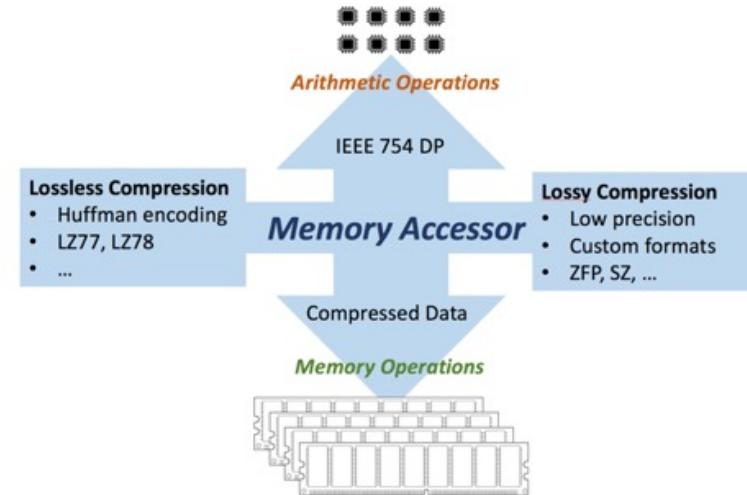
Use the memory accessor to boost performance

Can we use the memory accessor to accelerate the algorithm without changing the final result?

- Yes, if we can do all operations in registers and write the final result in high precision.
- Not in general, if we read/write intermediate date in low precision.
- We need to analyze the error propagation and adapt the algorithms to the application & data.

Possibilities in the context of solving linear systems:

- *Approximate linear operators / Preconditioners / Inner solvers;*
- *“Self-healing” iterative methods;*

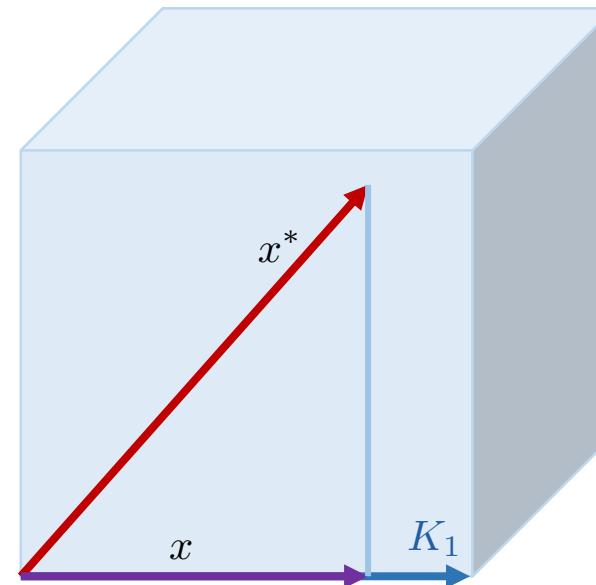


Rethinking Algorithms: Self-Healing Iterative Methods

- **Krylov iterative solvers**
- **Krylov methods** aim at approximating the solution to a linear problem in a subspace.
- Over the iterations, a nested sequence of **Krylov subspaces** is generated, adding one basis vector in each iteration.
- **Orthonormalization** ensures a orthonormal basis is formed (Classical Gram-Schmidt, Modified Gram Schmidt...).

$$K_0 \subset K_1 \subset K_2 \subset \dots$$

$$K_i(A, r) = \text{span}\{b, Ab, A^2b, \dots A^{i-1}b\}$$

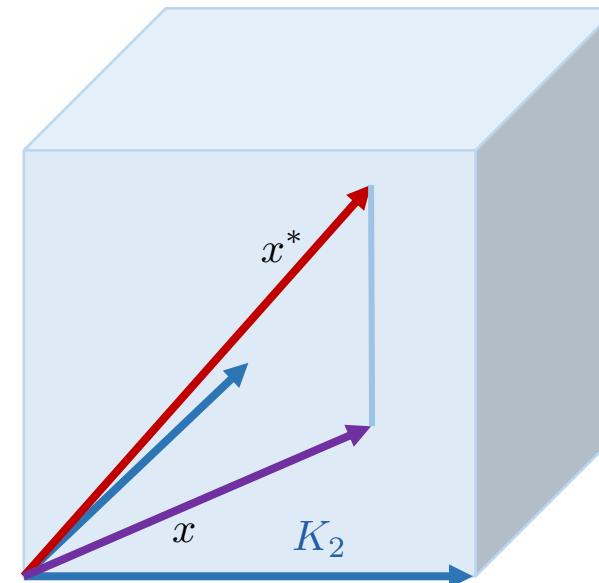


Rethinking Algorithms: Self-Healing Iterative Methods

- **Krylov iterative solvers**
- **Krylov methods** aim at approximating the solution to a linear problem in a subspace.
- Over the iterations, a nested sequence of **Krylov subspaces** is generated, adding one basis vector in each iteration.
- **Orthonormalization** ensures a orthonormal basis is formed (Classical Gram-Schmidt, Modified Gram Schmidt...).

$$K_0 \subset K_1 \subset K_2 \subset \dots$$

$$K_i(A, r) = \text{span}\{b, Ab, A^2b, \dots A^{i-1}b\}$$



Rethinking Algorithms: Self-Healing Iterative Methods

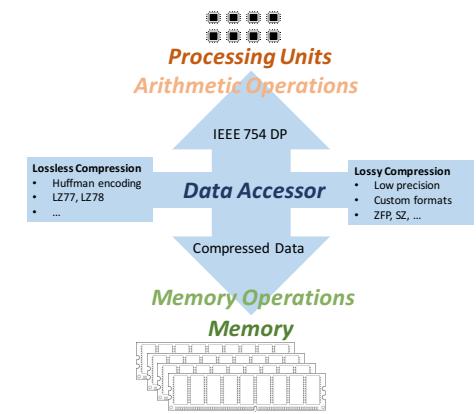
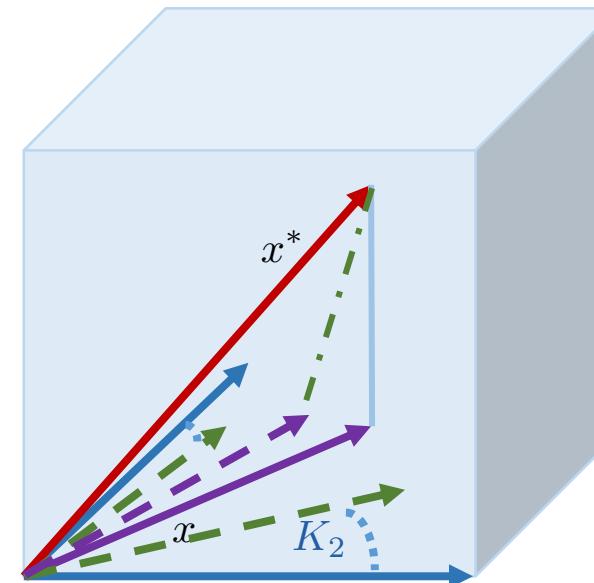
- **Krylov iterative solvers**
- **Krylov methods** aim at approximating the solution to a linear problem in a subspace.
- Over the iterations, a nested sequence of **Krylov subspaces** is generated, adding one basis vector in each iteration.
- **Orthonormalization** ensures a orthonormal basis is formed (Classical Gram-Schmidt, Modified Gram Schmidt...).

Compressed Basis (CB-) GMRES

- Use double precision in all arithmetic operations;
- Store Krylov basis vectors in lower precision;
 - Search directions are no longer DP-orthogonal;
 - Hessenberg system maps solution to “perturbed” Krylov subspace;
 - Additional iterations may be needed;
 - As long as the loss-of-orthogonality is moderate, we should see moderate convergence degradation;

$$K_0 \subset K_1 \subset K_2 \subset \dots$$

$$K_i(A, r) = \text{span}\{b, Ab, A^2b, \dots, A^{i-1}b\}$$

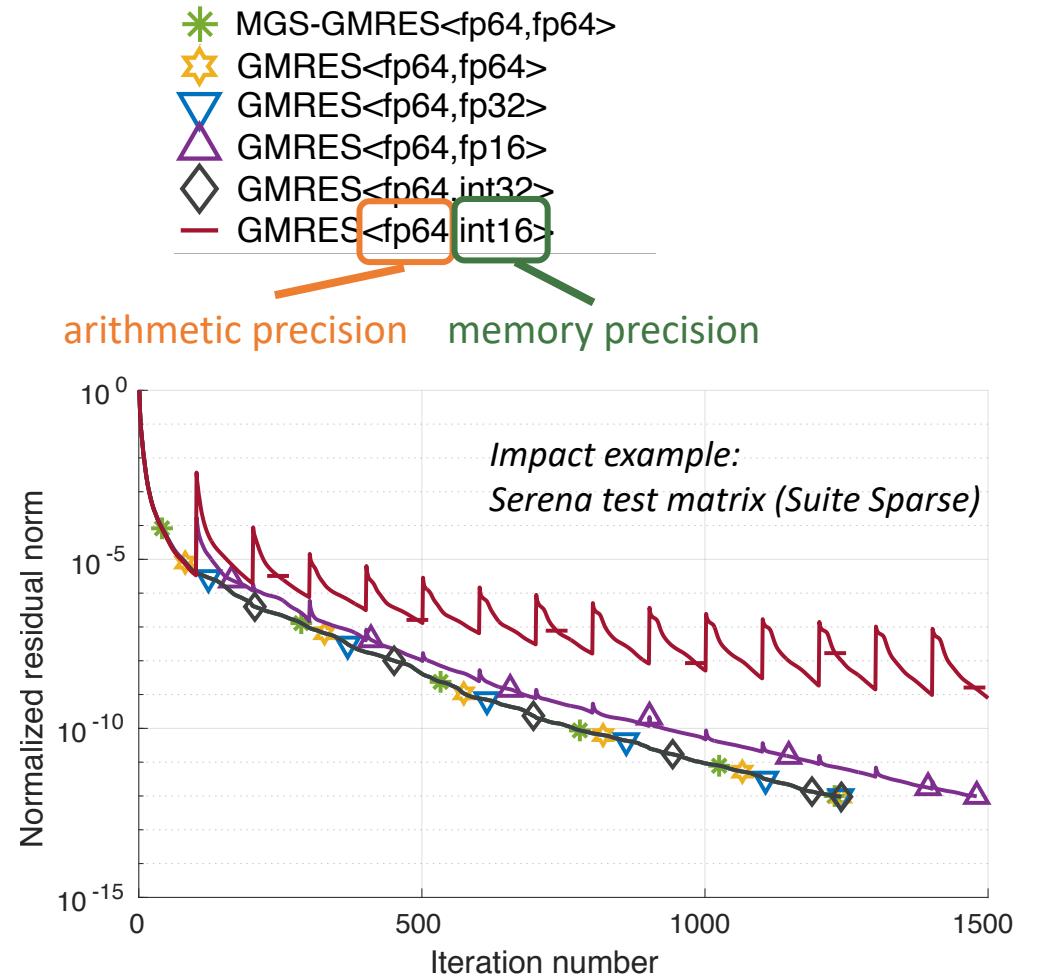


Compressed Basis (CB-) GMRES

- **Krylov iterative solvers**
- **Krylov methods** aim at approximating the solution to a linear problem in a subspace.
- Over the iterations, a nested sequence of **Krylov subspaces** is generated, adding one basis vector in each iteration.
- **Orthonormalization** ensures a orthonormal basis is formed (Classical Gram-Schmidt, Modified Gram Schmidt...).

Compressed Basis (CB-) GMRES

- Use double precision in all arithmetic operations;
- Store Krylov basis vectors in lower precision;
 - Search directions are no longer DP-orthogonal;
 - Hessenberg system maps solution to “perturbed” Krylov subspace;
 - Additional iterations may be needed;
 - As long as the loss-of-orthogonality is moderate, we should see moderate convergence degradation;



Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) NVIDIA V100 GPU

Double precision GMRES

Initial residual norm Relative residual $\sim 10^{-12}$

$\text{sqrt}(r^T r)$: 9670.36

Final residual norm

$\text{sqrt}(r^T r)$: 9.6639e-09

GMRES iteration count: 23271

GMRES execution time: 43801 ms

Single precision GMRES

Initial residual norm Relative residual $\sim 10^{-7}$

$\text{sqrt}(r^T r)$: 9670.36

Final residual norm

$\text{sqrt}(r^T r)$: 0.00175464

GMRES iteration count: 25000

GMRES execution time: 27376 ms

~2x faster!

Running iterative methods in different precision formats

Linear System $Ax=b$ with $\text{cond}(A) \approx 10^7$

(apache2 from SuiteSparse) NVIDIA V100 GPU

Double precision GMRES

Initial residual norm **Relative residual $\sim 10^{-12}$**

$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 9.6639e-09$

GMRES iteration count: 23271

GMRES execution time: 43801 ms

Single precision GMRES

Initial residual norm **Relative residual $\sim 10^{-7}$**

$\sqrt{r^T r} : 9670.36$

Final residual norm

$\sqrt{r^T r} : 0.00175464$

GMRES iteration count: 25000

GMRES execution time: 27376 ms

Compressed Basis GMRES

Initial residual norm **Relative residual $\sim 10^{-12}$**

$\sqrt{r^T r} : 9670.36$

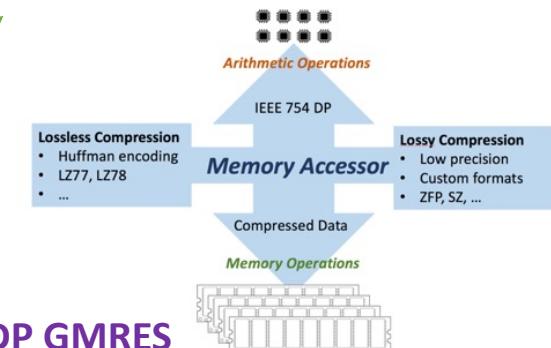
Final residual norm

$\sqrt{r^T r} : 9.6591e-09$

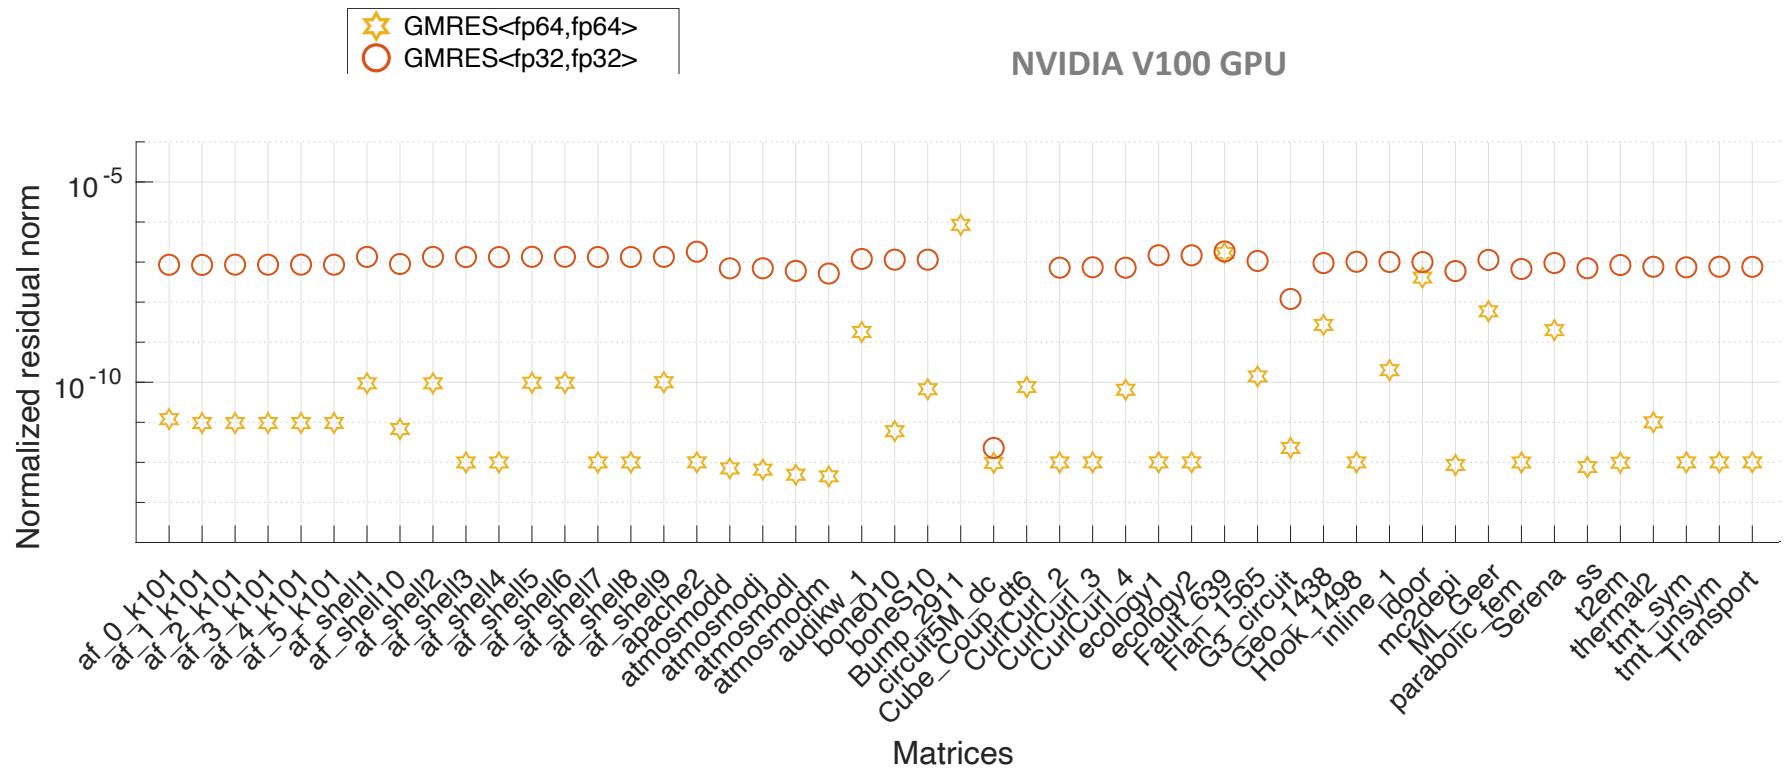
GMRES iteration count: 23271

GMRES execution time: 29369 ms

Accuracy of DP GMRES
Performance similar to SP GMRES

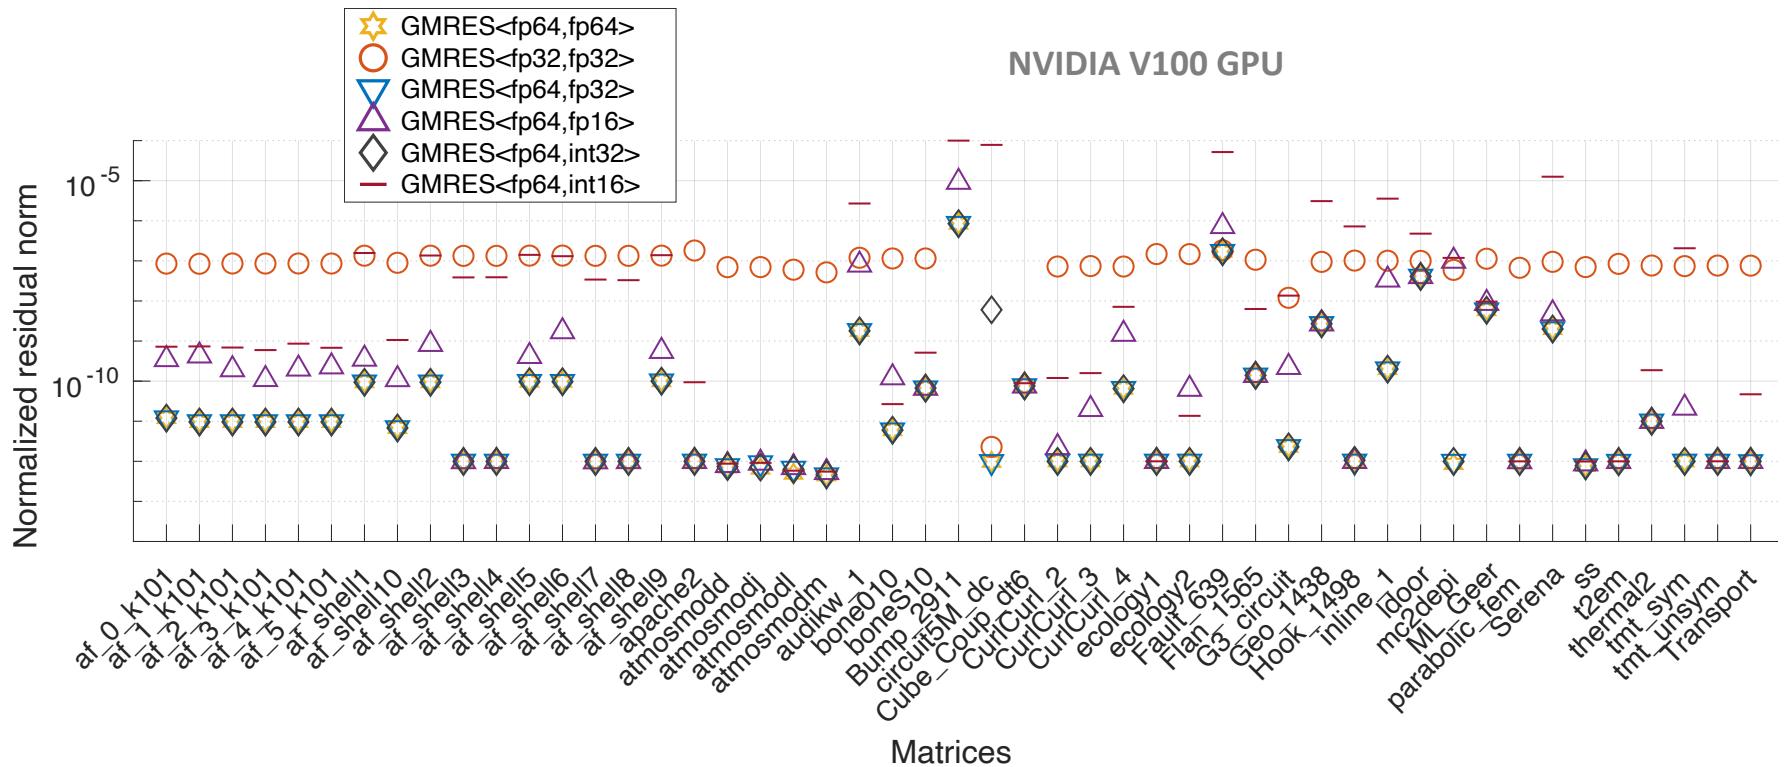


Compressed Basis GMRES



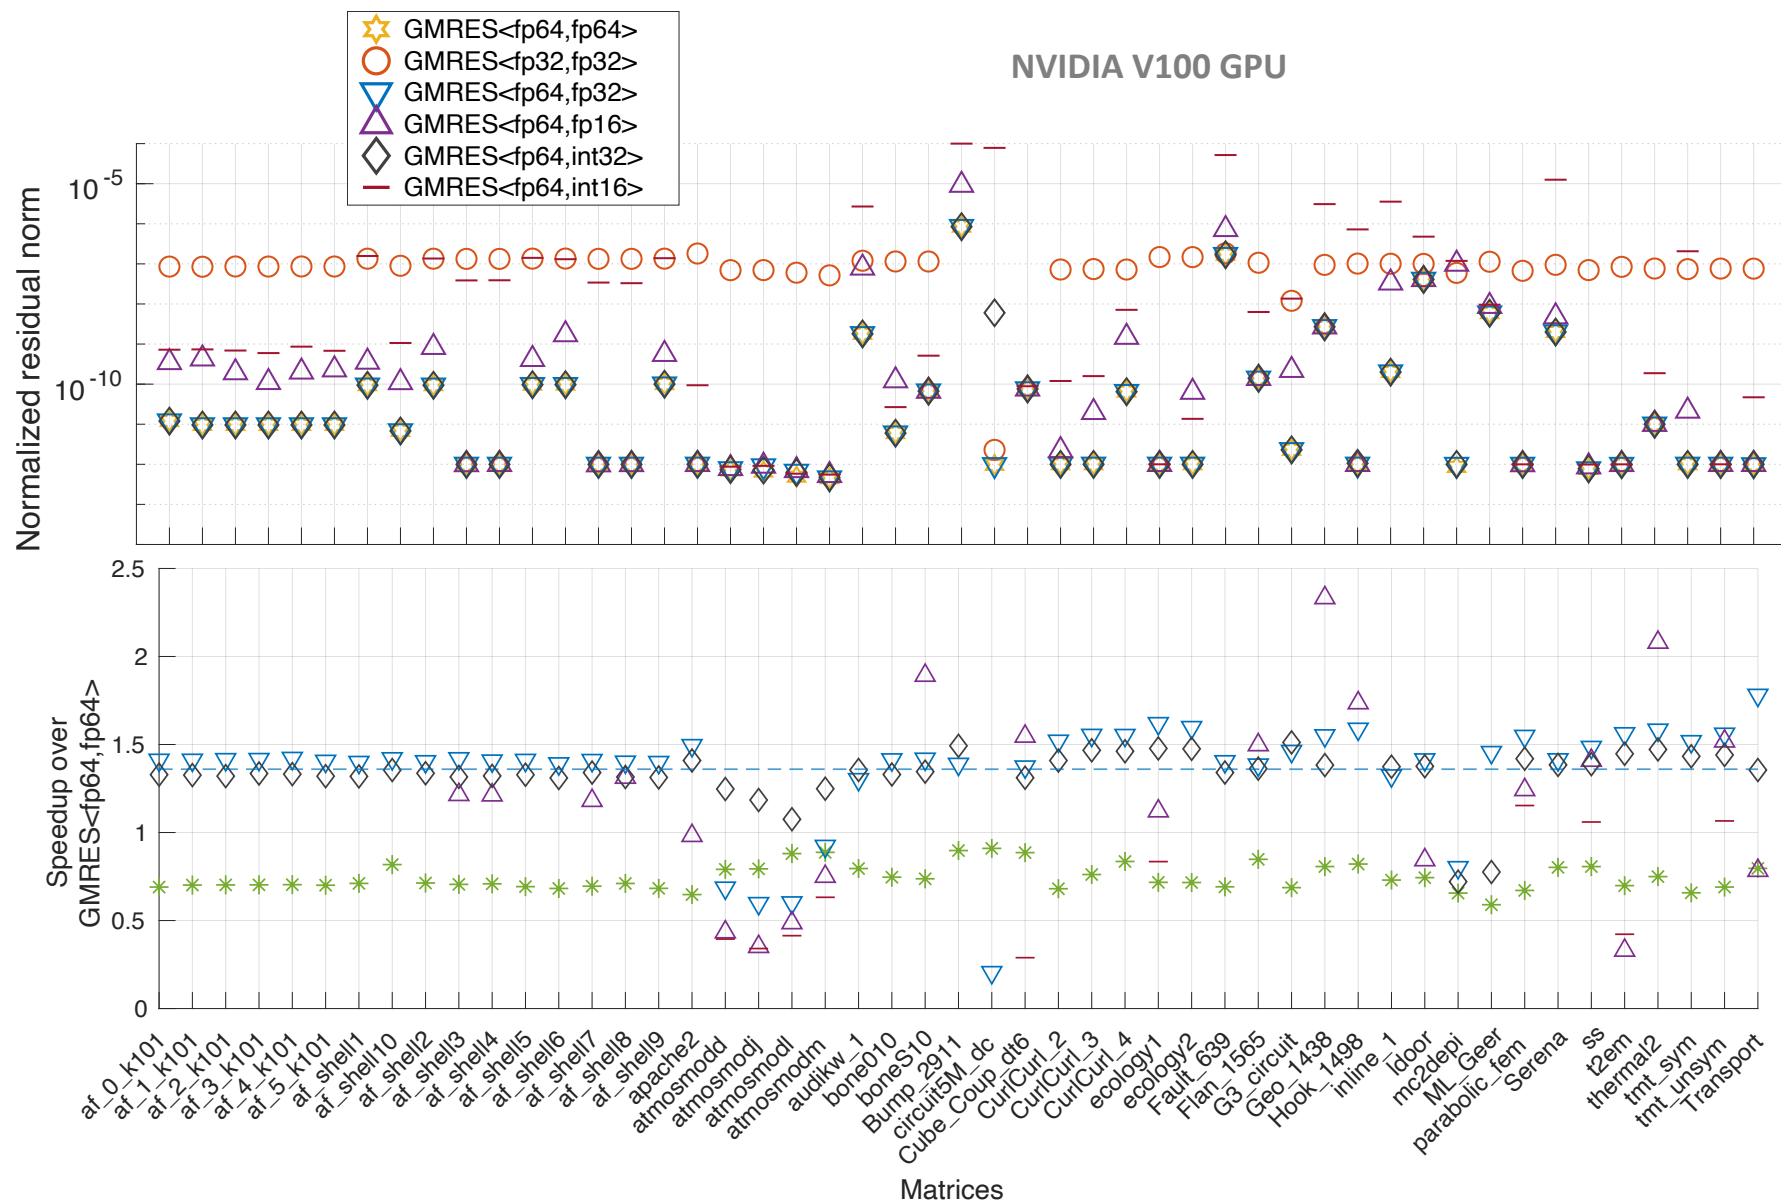
Compressed Basis GMRES

- CB-GMRES using 32-bit storage preserves DP accuracy (SP-GMRES does not)



Compressed Basis GMRES

- CB-GMRES using 32-bit storage preserves DP accuracy (SP-GMRES does not)
- Speedups problem-dependent
- Speedup $\varnothing 1.4x$ (for restart 100)
- 16-bit storage mostly inefficient



Integration into MFEM

Improve current Ginkgo-MFEM integration:

- ✓ MFEM and Ginkgo operate directly on same data without copies
- ✓ New GinkgoExecutor class automatically matches MFEM Device configuration - for CPU, CUDA, or HIP
- ✓ Ginkgo can use MFEM matrix-free operators in solvers

Add Ginkgo preconditioners to MFEM:

- ✓ Ginkgo preconditioners can be used with Ginkgo solvers, or used with MFEM solvers
- ✓ Includes Ginkgo's new ILU-ISAI/IC-ISAI preconditioners, which use the Incomplete Sparse Approximate Inverse to apply the ILU or IC factorization for improved GPU performance

Add new Ginkgo solver to MFEM:

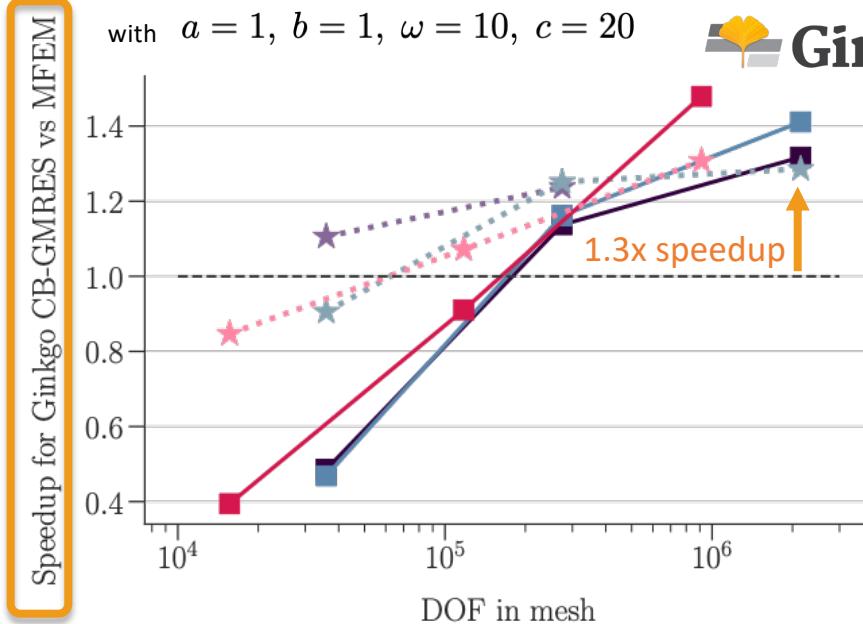
- ✓ Integration for Ginkgo's Compressed Basis GMRES solver, which uses mixed precision techniques for speedup (see example to right)

Example: Speeding up MFEM's “example 22” on NVIDIA and AMD GPUs

Example 22 solves harmonic oscillation problems, with a forced oscillation imposed at the boundary. For this test, we use variant 1:

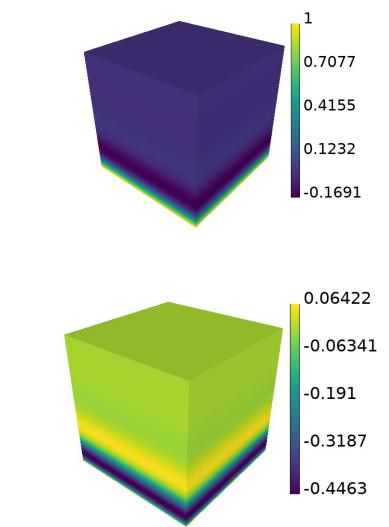
$$-\nabla \cdot (a \nabla u) - \omega^2 bu + i\omega cu = 0$$

with $a = 1, b = 1, \omega = 10, c = 20$



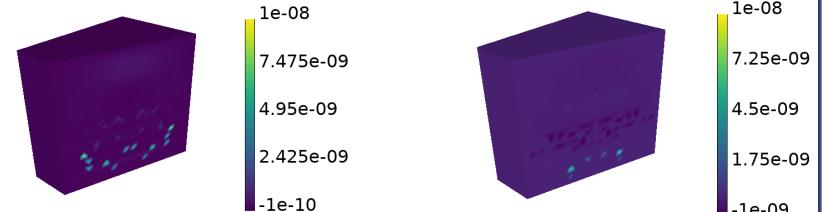
Speedup of Ginkgo's Compressed Basis-GMRES solver vs MFEM's GMRES solver for three different orders of basis functions (p) for MFEM's example 22. The tests use the “partial assembly” type of MFEM matrix-free operators.

CUDA 10.1/NVIDIA V100 and ROCm 4.0/AMD MI50. GMRES(50) used for both solvers. CB-GMRES used float/double.



From top: Real part of solution, imaginary part of solution.

Below: Slice of difference in solution output using MFEM solver versus Ginkgo CB-GMRES. Real part (left), imaginary part (right)



Natalie Beams (Univ. of Tennessee)

Using the memory accessor to boost accuracy

Instead of improving the performance of memory-bound high precision algorithms, the memory accessor can be used to increase the accuracy of memory-bound low precision algorithms – at no cost.

Design

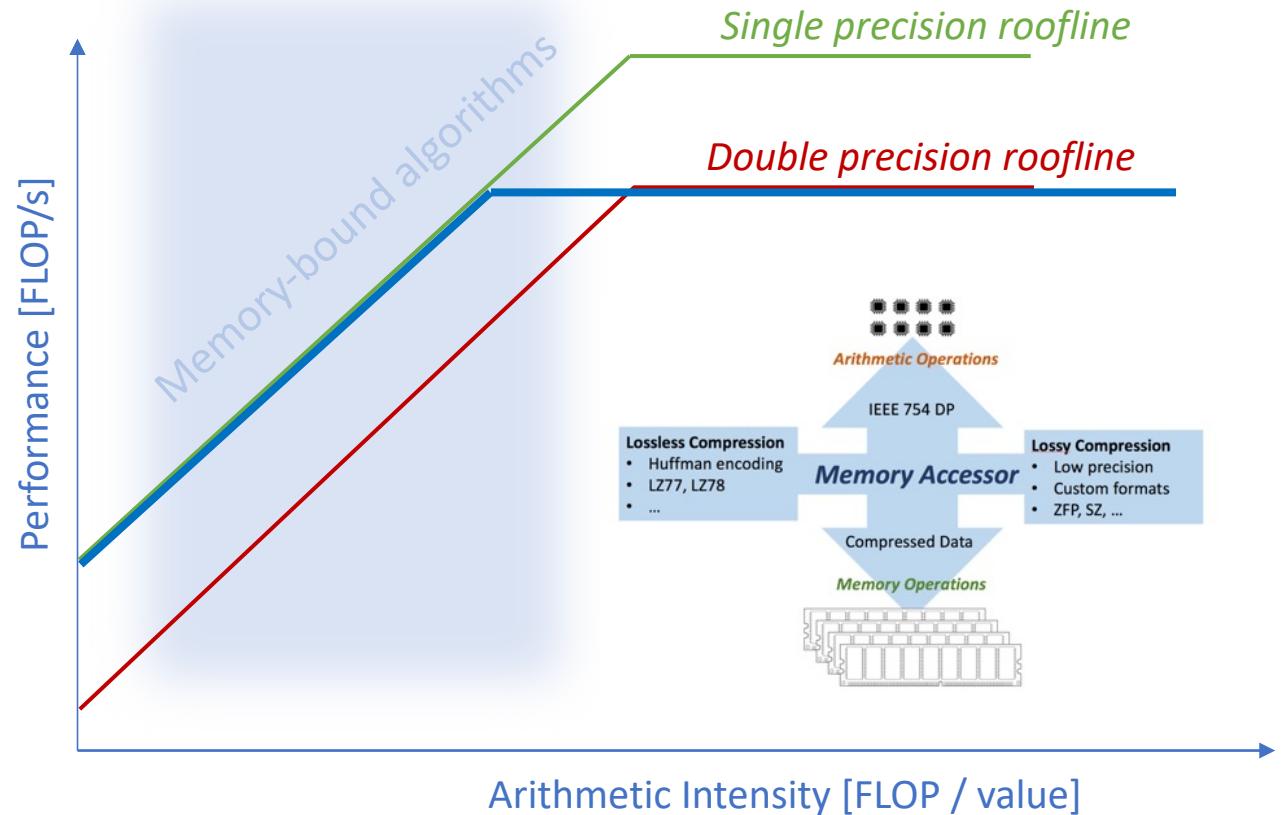
- Memory access in low precision (e.g. fp32);
- Computations in high precision (e.g. fp64);

Characteristics

- Performance of low precision BLAS;
- Higher accuracy than low precision BLAS;

Usage

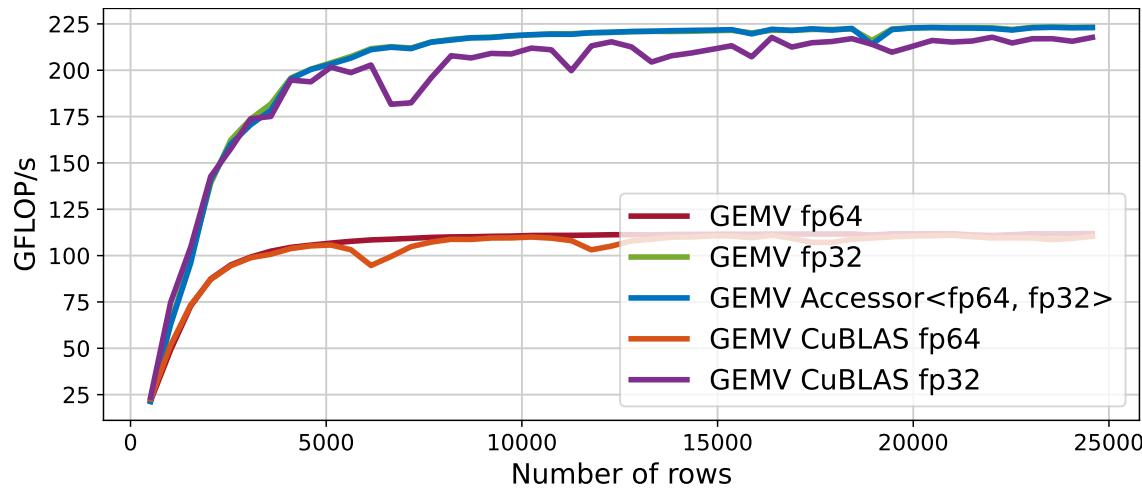
1. Can replace low precision BLAS to increase accuracy;
2. Can replace high precision BLAS if information loss is acceptable;
(without having to deal with explicit mixed precision usage)



Accessor-BLAS: Replacing LP BLAS to improve accuracy

GEMV

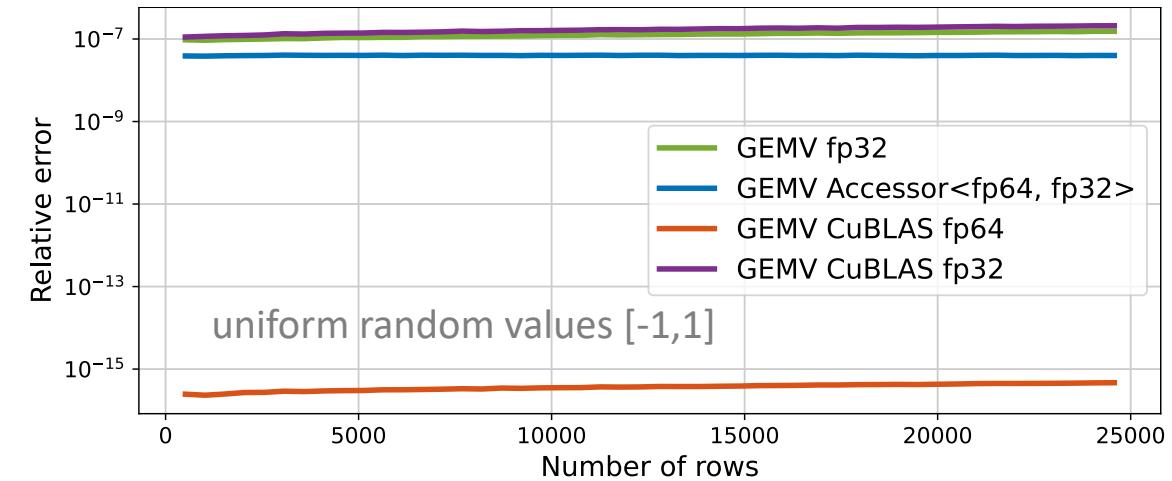
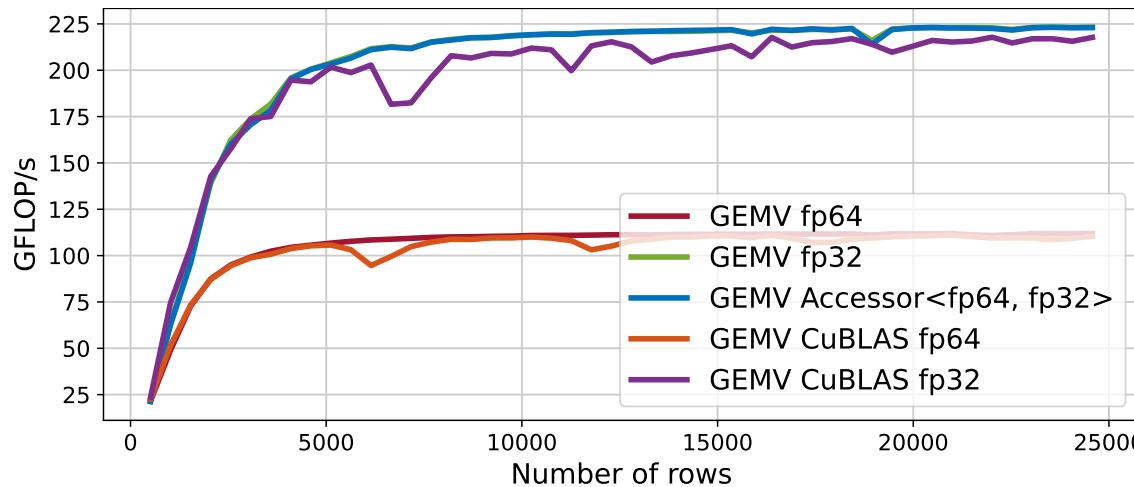
NVIDIA V100 GPU (Summit)



Accessor-BLAS: Replacing LP BLAS to improve accuracy

GEMV

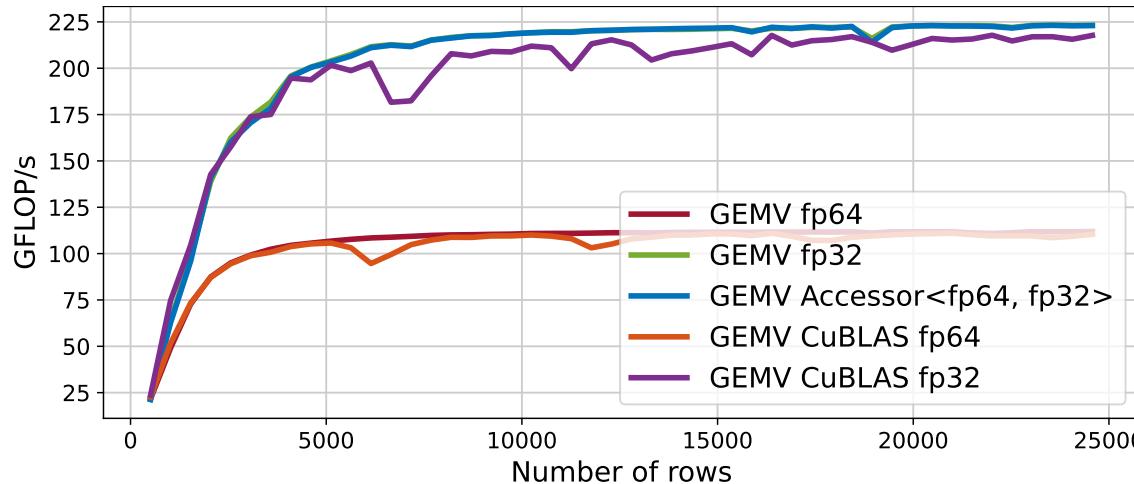
NVIDIA V100 GPU (Summit)



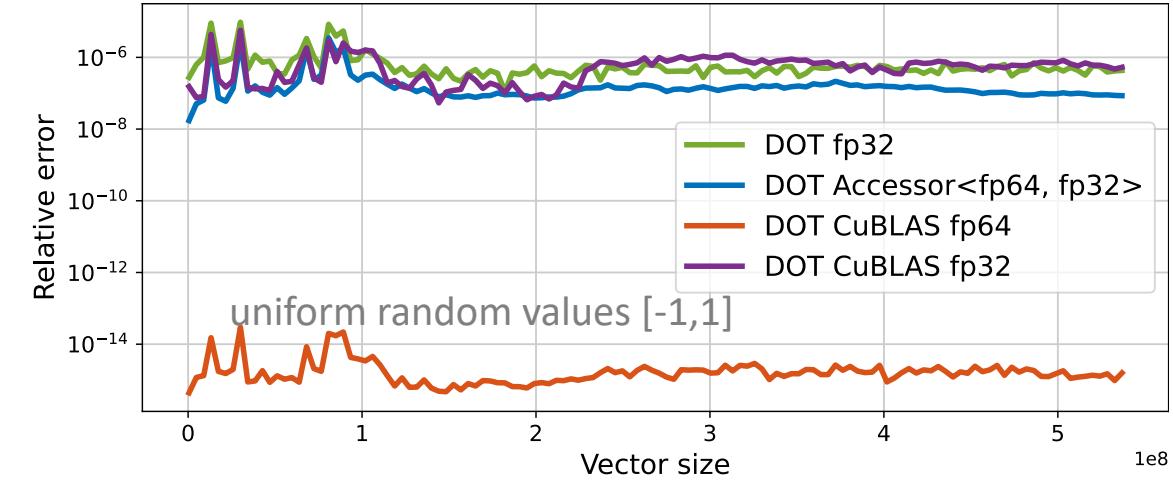
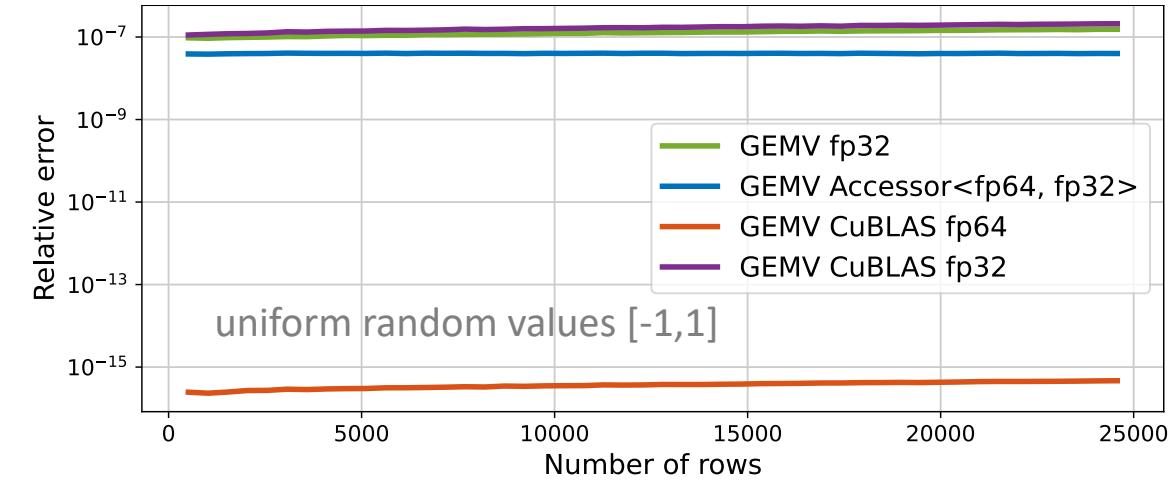
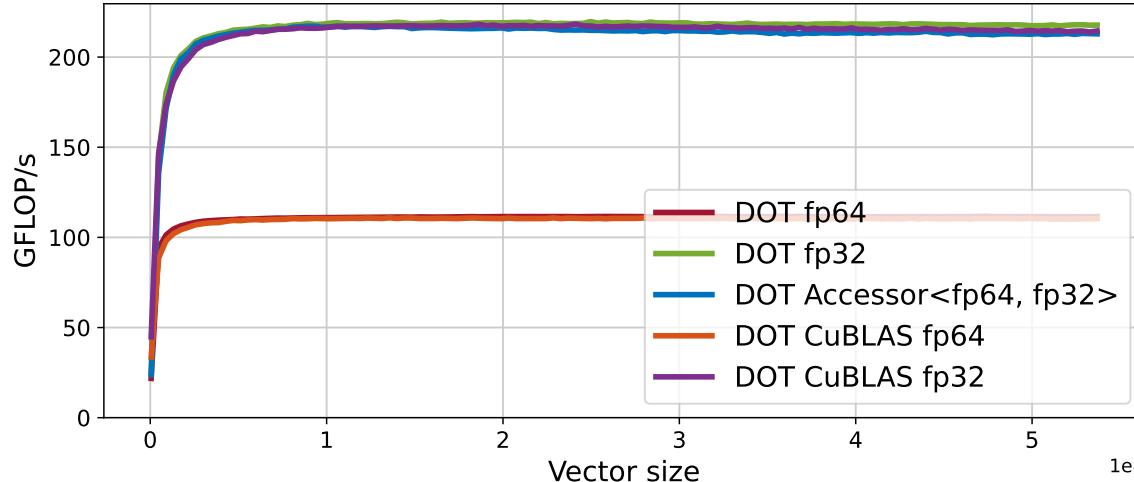
Accessor-BLAS: Replacing LP BLAS to improve accuracy

GEMV

NVIDIA V100 GPU (Summit)

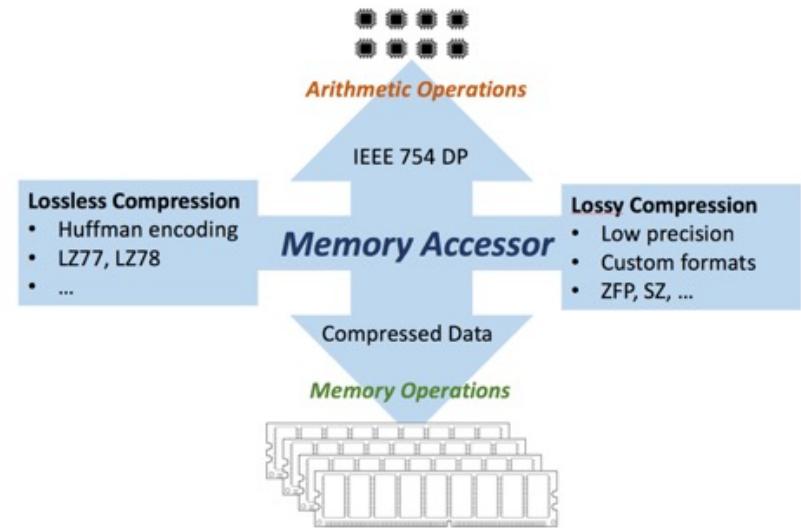


DOT



Summary and resources

- For memory-bound algorithms, mixed precision can boost performance through reduced data movement.
- **Memory accessor** allows to compress data in main memory but do all arithmetic in high (**double**) precision.
- **Approximate operators** (preconditioners, lower multigrid levels) and **self-healing iterative methods** can accept/compensate **information loss**.
- Memory-bound low precision algorithms can increase accuracy at no cost.



Accessor-based GEMV and DOT available as open-source code:
<https://github.com/ginkgo-project/Accessor-BLAS>

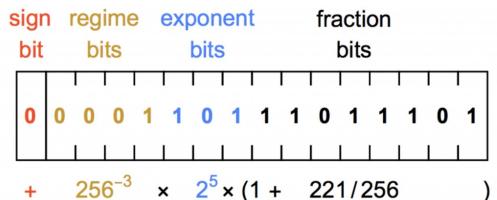
Mixed Precision block-Jacobi preconditioning:
<https://github.com/ginkgo-project/ginkgo/tree/develop/examples/adaptiveprecision-blockjacobi>

Mixed Precision Iterative Refinement:
<https://github.com/ginkgo-project/ginkgo/tree/develop/examples/mixed-precision-ir>

Compressed-Basis GMRES:
<https://github.com/ginkgo-project/ginkgo/tree/develop/examples/cb-gmres>

Let's try harder

- IEEE 754 fp64 in arithmetic operations
- More sophisticated in-register compression
 - Custom formats
 - Compression techniques (SZ, ZFP)
- Store data in compressed format



John L. Gustafson

Dynamic Range

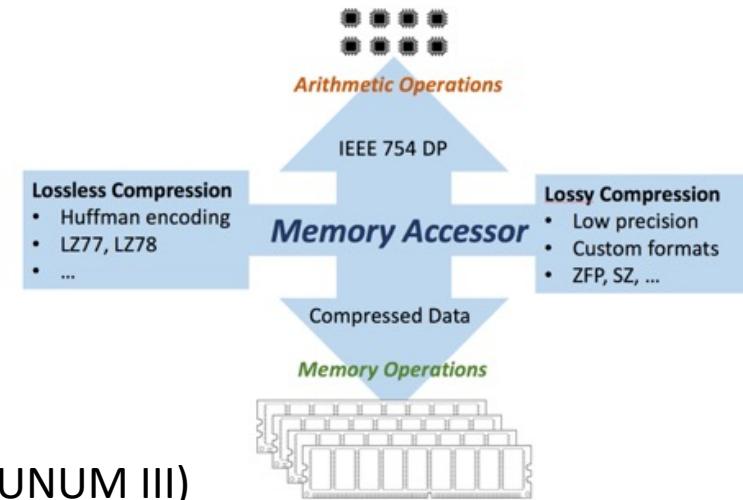
Size [bits]	IEEE exponent size [bits]	Approx. IEEE dynamic range	Approx. Posit dynamic range	Posit exp. bits
16	5	$[6 \cdot 10^{-8}, 7 \cdot 10^4]$	$[1 \cdot 10^{-17}, 7 \cdot 10^{16}]$	2
32	8	$[1 \cdot 10^{-45}, 3 \cdot 10^{38}]$	$[8 \cdot 10^{-37}, 1 \cdot 10^{36}]$	2
64	11	$[5 \cdot 10^{-324}, 2 \cdot 10^{308}]$	$[2 \cdot 10^{-75}, 5 \cdot 10^{75}]$	2

Special values

- IEEE defines ± 0 , $\pm \infty$ and NaN (quiet and signaling) as special values
- A lot of NaN representations (fp32 has $2^{24} - 1 \approx 10^7$ different NaNs)
- Posit only has 2: 0 and NaR
- NaR (Not a Real) is used as an error-value (like NaN and $\pm \infty$)

Gradual over- and underflow

- IEEE supports gradual underflow with subnormal numbers (fraction has an implicit 0.)
- No support for gradual overflow
- Posit supports both gradual over- and underflow through the regime
- The farther away from 1.0, the fewer fraction bits



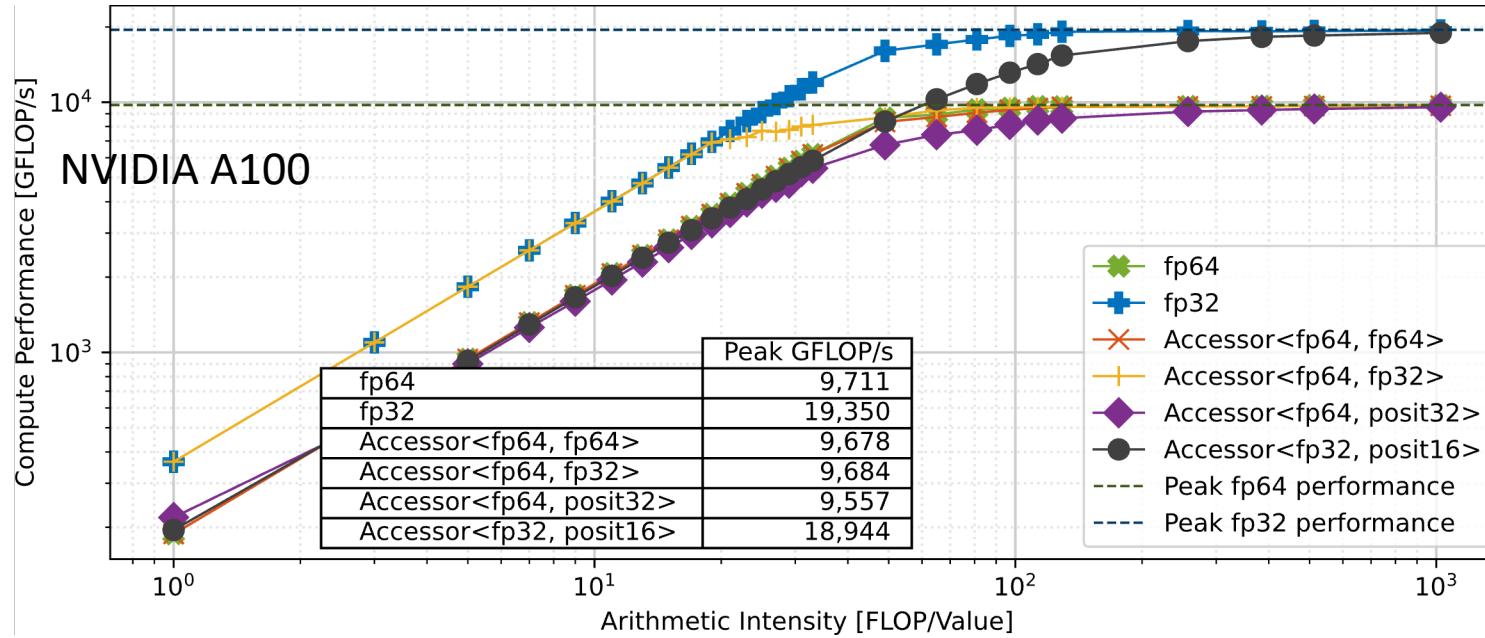
POSIT (UNUM III)

Using POSIT as memory format

sign regime exponent fraction
bit bits bits

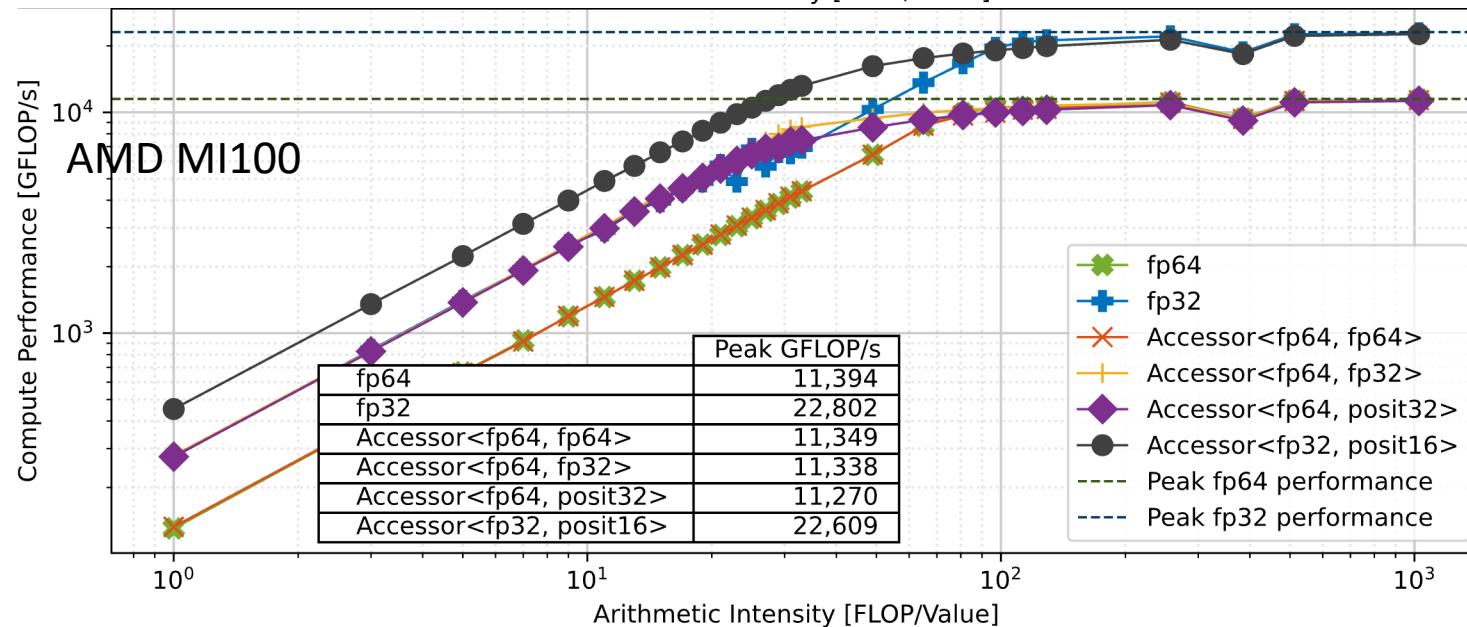
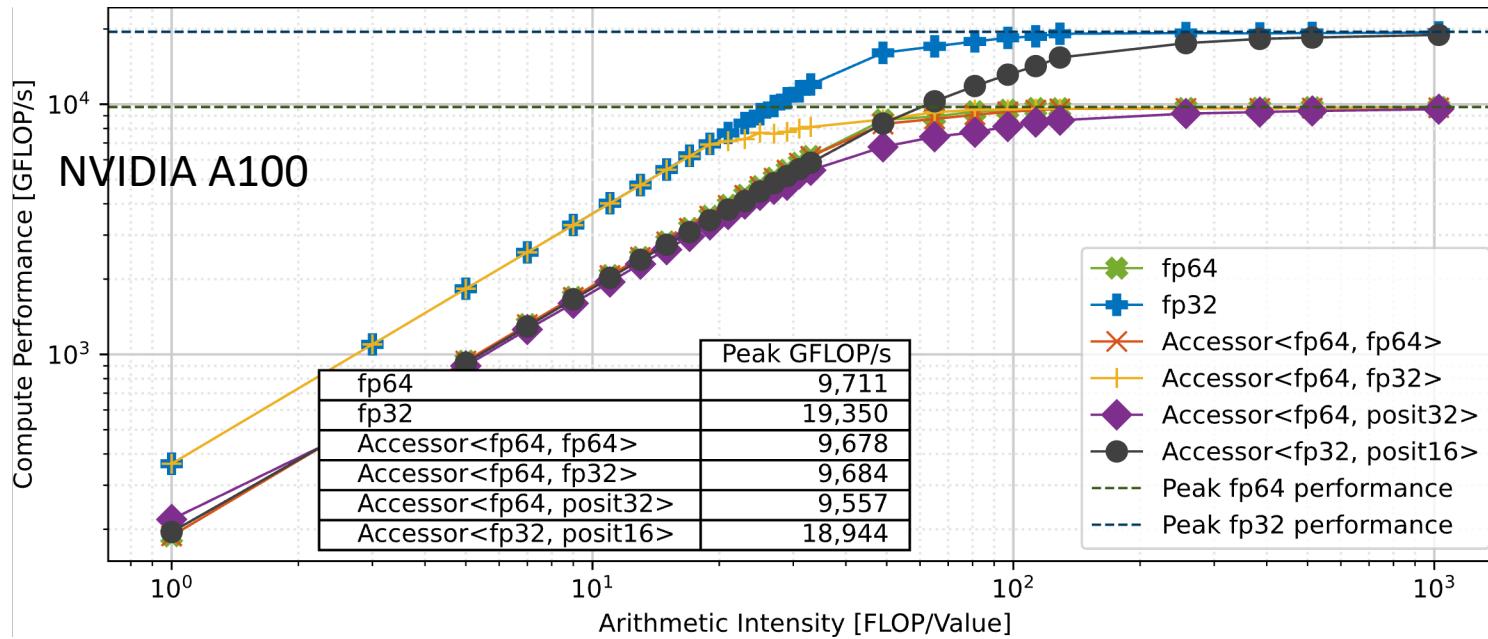
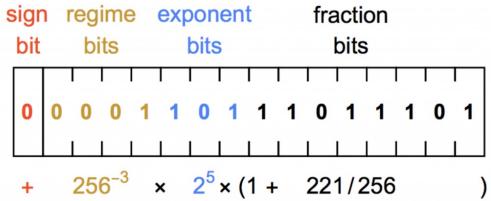
+ 256^{-3} $\times 2^5 \times (1 + 221/256)$)

Unum Type III
John L. Gustafson



T. Grützmacher

Using POSIT as memory format



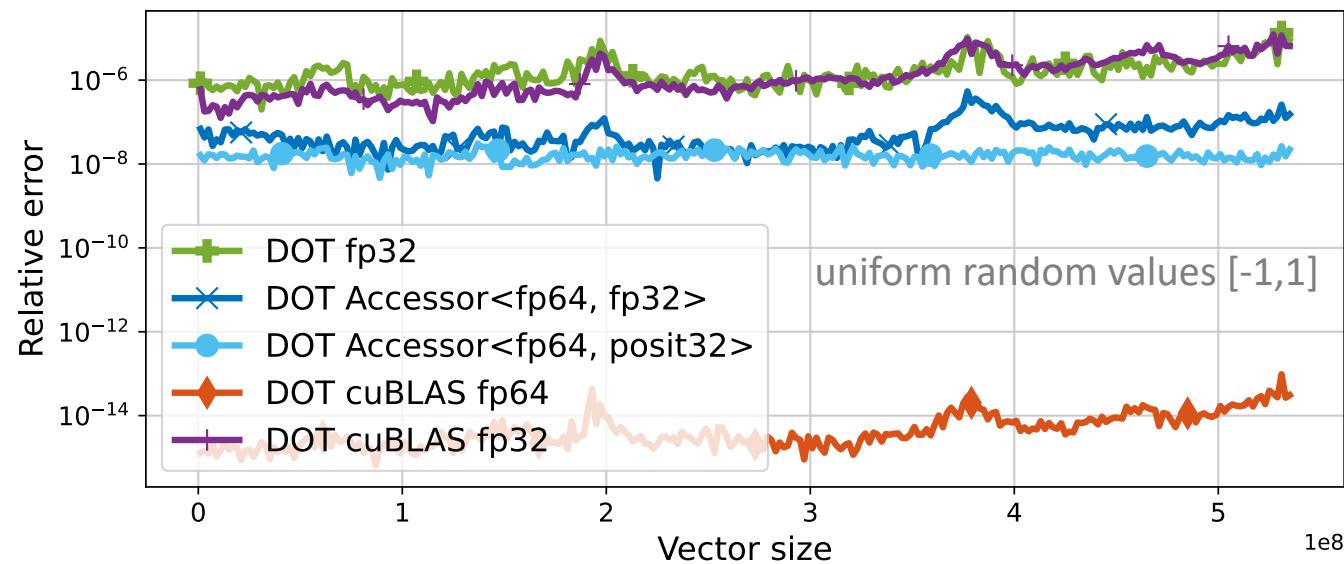
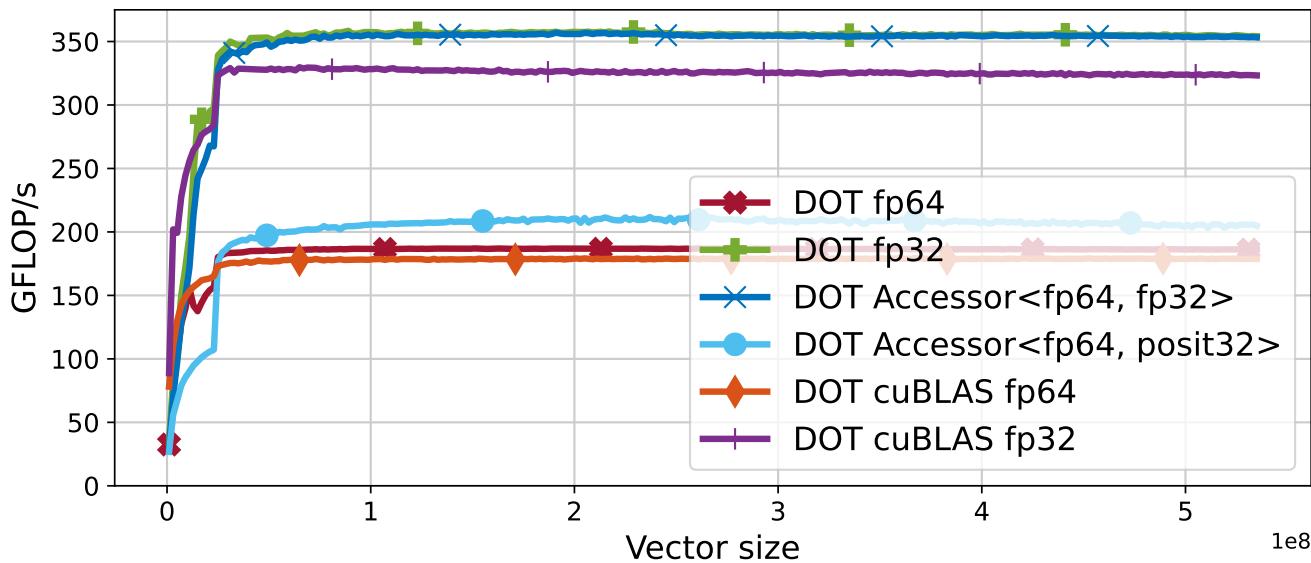
T. Grützmacher

Using POSIT as memory format

NVIDIA A100



T. Grützmacher



Using ZFP / SZ compression as memory format



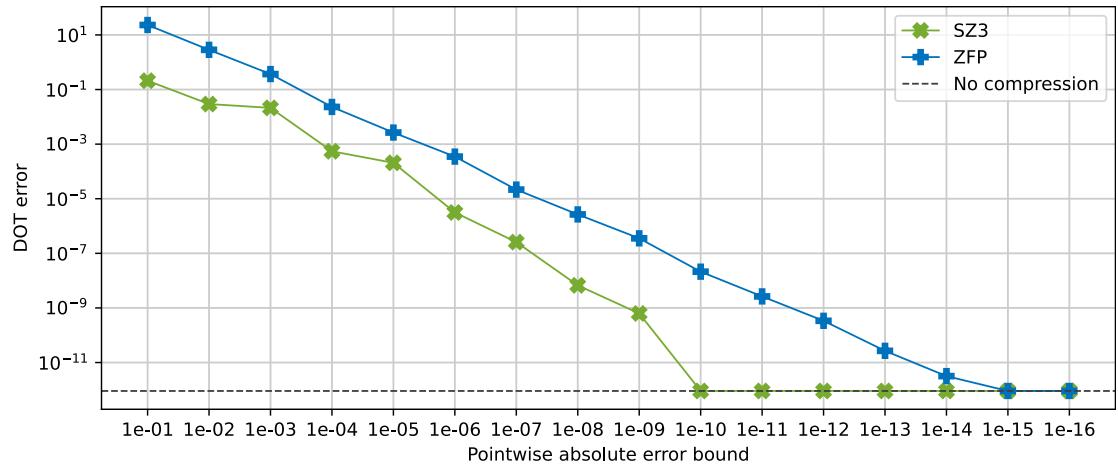
F. Capello

R. Underwood

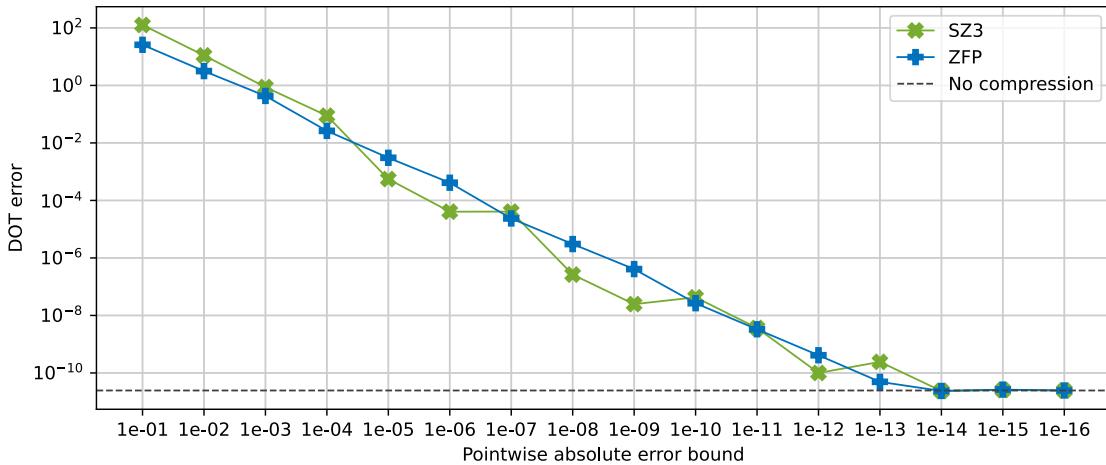
T. Grützmacher

Using ZFP / SZ compression as memory format

Random data



1D sine function



F. Capello

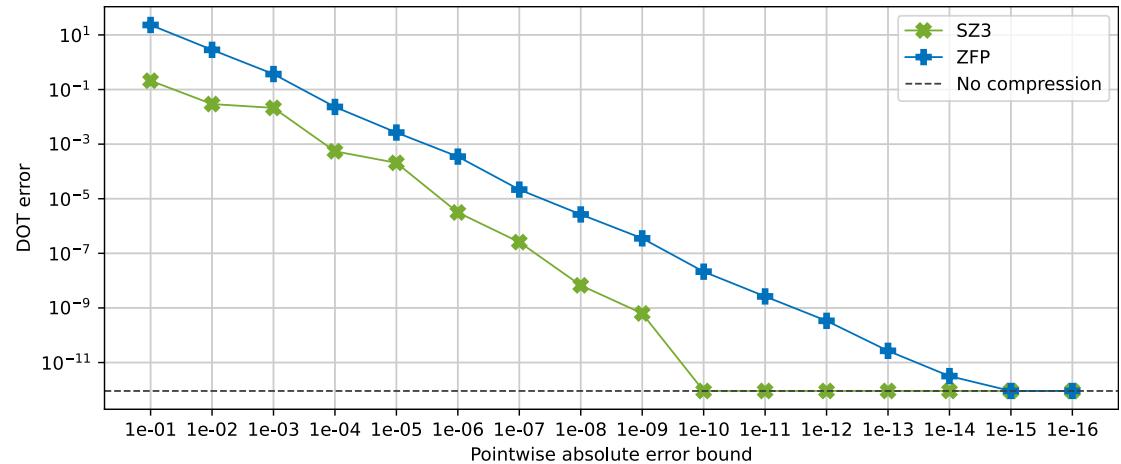
R. Underwood

T. Grütmacher

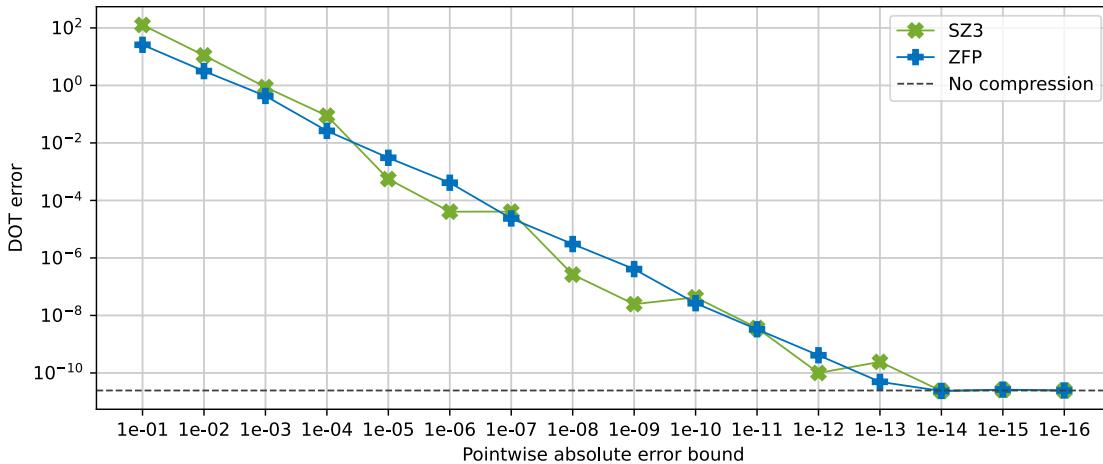
Using ZFP / SZ compression as memory format



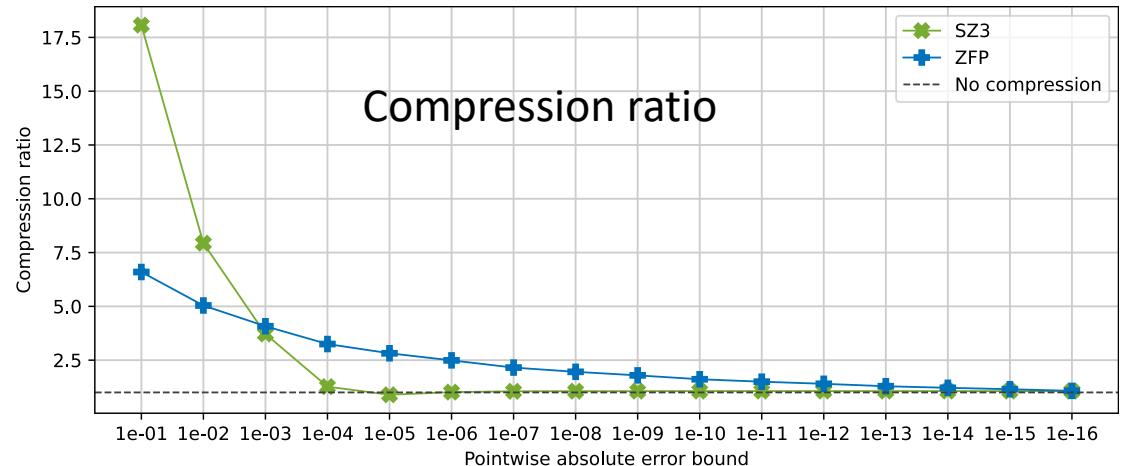
Random data



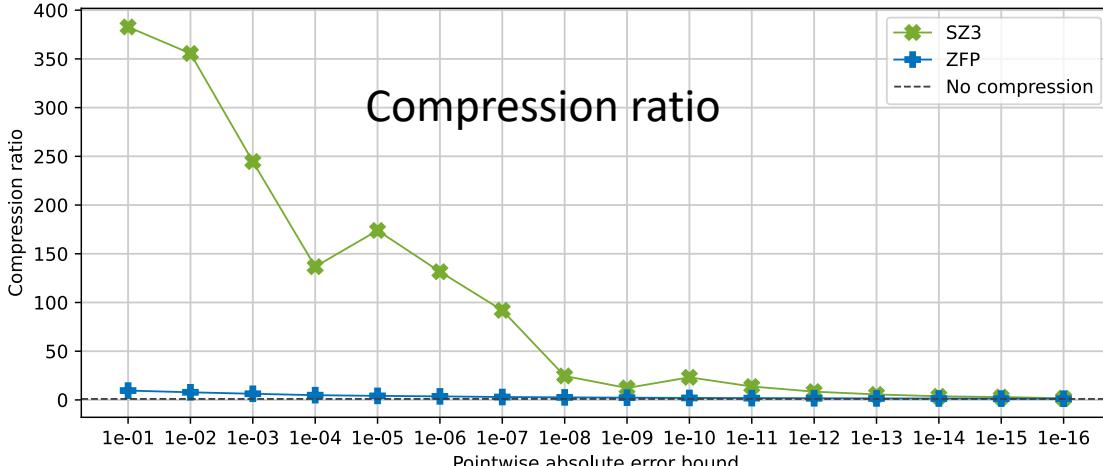
1D sine function



Compression ratio

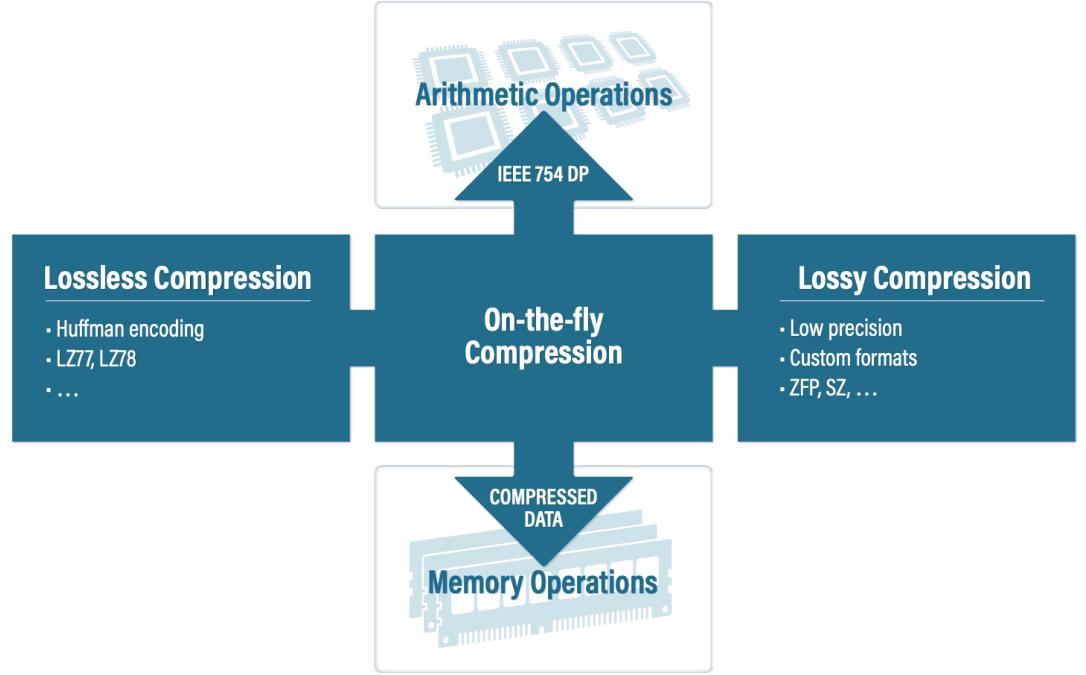


Compression ratio



Let's try harder

- IEEE 754 fp64 in arithmetic operations
- More sophisticated in-register compression
 - Custom formats
 - Compression techniques (SZ, ZFP)
- Store data in compressed format



Trade-off:

- Aggressive compression comes with larger information loss
- Element-wise compression allows only for moderate compression ratios
- Block-wise compression makes random access difficult
- Register count limits the block size (hardware specific)
- Data-dependent compression efficiency