

9th JLESC Workshop

April 15th-17th, 2019 | Knoxville, TN



ParILUT - A New Parallel Threshold ILU

Hartwig Anzt, Edmond Chow, Jack Dongarra



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE

**Joint Laboratory
for Extreme-Scale Computing**



Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $nnz(L + U) = nnz(A)$).

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $nnz(L + U) = nnz(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $\text{nnz}(L + U) = \text{nnz}(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Exact LU Factorization

- Decompose system matrix into product $A = L \cdot U$.
- Based on Gaussian elimination.
- Triangular solves to solve a system $Ax = b$:

$$Ly = b \Rightarrow y \quad \Rightarrow \quad L y = b \Rightarrow x$$

- De-Facto standard for solving dense problems.
- *What about sparse? Often significant fill-in...*

$$\left(\begin{array}{cccccc} \times & \times & \times & \times & \times & \times \\ \times & \times & \times & \times & & \\ \times & \times & \times & \times & & \\ \times & & \times & \times & \times & \\ \times & & & \times & \times & \times \\ \times & & & \times & \times & \times \\ \times & & & & \times & \times \end{array} \right)$$

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $\text{nnz}(L + U) = \text{nnz}(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Exact LU Factorization

- Decompose system matrix into product $A = L \cdot U$.
- Based on Gaussian elimination.
- Triangular solves to solve a system $Ax = b$:

$$Ly = b \Rightarrow y \quad \Rightarrow \quad L y = b \Rightarrow x$$

- De-Facto standard for solving dense problems.
- *What about sparse? Often significant fill-in...*

$$\left(\begin{array}{cccccc|cc} x & x & x & x & & x & x \\ x & x & x & & & & \\ x & x & x & x & & & \\ x & & x & x & x & & \\ x & & & x & x & x & x \\ x & & x & x & x & x & x \\ x & & & x & x & & \\ & x & x & x & & x & x \end{array} \right)$$

Incomplete LU Factorization (ILU)

- **Focused on restricting fill-in to a specific sparsity pattern \mathcal{S} .**

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $\text{nnz}(L + U) = \text{nnz}(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Exact LU Factorization

- Decompose system matrix into product $A = L \cdot U$.
- Based on Gaussian elimination.
- Triangular solves to solve a system $Ax = b$:

$$Ly = b \Rightarrow y \quad \Rightarrow \quad Ly = b \Rightarrow x$$

- De-Facto standard for solving dense problems.
- *What about sparse? Often significant fill-in...*

$$\left(\begin{array}{ccccccccc} x & x & x & x & & x & x \\ x & x & x & & & & & \\ x & x & x & x & & & & \\ x & & x & x & x & & & \\ x & & & x & x & x & & \\ x & & & x & x & x & x & \\ x & & & & x & x & x & \\ & & & & x & x & x & \\ & & & & x & x & x & \end{array} \right)$$

Incomplete LU Factorization (ILU)

- Focused on restricting fill-in to a specific sparsity pattern \mathcal{S} .
- For **ILU(0)**, \mathcal{S} is the sparsity pattern of A .
 - Works well for many problems.
 - *Is this the best we can get for nonzero count?*

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $\text{nnz}(L + U) = \text{nnz}(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Exact LU Factorization

- Decompose system matrix into product $A = L \cdot U$.
- Based on Gaussian elimination.
- Triangular solves to solve a system $Ax = b$:

$$Ly = b \Rightarrow y \quad \Rightarrow \quad Ly = b \Rightarrow x$$

- De-Facto standard for solving dense problems.
- *What about sparse? Often significant fill-in...*

$$\left(\begin{array}{cccccc|cc} x & x & x & x & & x & x \\ x & x & x & & & & \\ x & x & x & x & & & \\ x & & x & x & x & & \\ x & & x & x & x & x & x \\ x & & x & x & x & x & x \\ x & & & x & x & x & x \\ & x & x & x & & x & x \end{array} \right)$$

Incomplete LU Factorization (ILU)

- **Focused on restricting fill-in** to a specific sparsity pattern \mathcal{S} .
- For **ILU(0)**, \mathcal{S} is the sparsity pattern of A .
 - Works well for many problems.
 - *Is this the best we can get for nonzero count?*
- Fill-in threshold ILU (**ILUT**) bases \mathcal{S} on the significance of elements (e.g. magnitude).
 - Often **better preconditioners** than level-based ILU.
 - Difficult to parallelize.

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $nnz(L + U) = nnz(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Rethink the overall strategy!

- Use a parallel iterative process to generate factors.

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $nnz(L + U) = nnz(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Rethink the overall strategy!

- Use a parallel iterative process to generate factors.
- The preconditioner should have a moderate number of nonzero elements,
but we don't care too much about intermediate data.

Motivation

We are looking for a factorization-based preconditioner such that $A \approx L \cdot U$.
is a good approximation with moderate nonzero count (e.g. $\text{nnz}(L + U) = \text{nnz}(A)$).

- *Where should these nonzero elements be located?*
- *How can we compute the preconditioner in a highly parallel fashion?*

Rethink the overall strategy!

- Use a parallel iterative process to generate factors.
- The preconditioner should have a moderate number of nonzero elements,
but we don't care too much about intermediate data.

1. *Select a set of nonzero locations.*
2. *Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.*
3. *Maybe change some locations in favor of locations that result in a better preconditioner.*
4. *Repeat until the preconditioner quality stagnates.*

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & * \\ * & & * & & * \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & & & * & \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

ILU residual $R = A - L \times U$

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & \\ * & & * & & \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & & & * & \\ * & * & & * & \\ * & * & & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

\leftarrow

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & \\ * & & & * & \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
 2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
 3. Maybe change some locations in favor of locations that result in a better preconditioner.
 4. Repeat until the preconditioner quality stagnates.
- This is an optimization problem...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & \\ * & & * & & \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix} \times \begin{pmatrix} * & & & \\ * & * & & \\ * & * & * & \\ * & & & * \\ * & * & & \\ * & * & & \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & & & \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix}$$
$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & & & \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
 2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
 3. Maybe change some locations in favor of locations that result in a better preconditioner.
 4. Repeat until the preconditioner quality stagnates.
- This is an optimization problem...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & * \\ * & & * & & * \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & & & * & \\ * & * & & * & \\ * & * & * & * & \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & * & * & * \end{pmatrix}$$
$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
 2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
 3. Maybe change some locations in favor of locations that result in a better preconditioner.
 4. Repeat until the preconditioner quality stagnates.
- This is an optimization problem...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & * \\ * & & * & & * \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & & & * & \\ * & * & & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & * & * & * \end{pmatrix}$$
$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & * \\ * & & * & & * \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem...

ILU residual
matrix pattern

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & & * & * & \\ * & * & & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & * & * & * \end{pmatrix}$$
$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem with $\text{nnz}(A - L \cdot U)$ equations and $\text{nnz}(L + U)$ variables.

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & & * \\ * & & * & & * \\ * & * & & * & * \\ * & * & & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$
$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem with $\text{nnz}(A - L \cdot U)$ equations and $\text{nnz}(L + U)$ variables.
- We may want to compute the values in L, U such that $R = A - L \cdot U = 0|_{\mathcal{S}}$, the approximation being exact in the locations included in \mathcal{S} , but not outside!

$\text{nnz}(L + U)$ equations
 $\text{nnz}(L + U)$ variables

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem with $\text{nnz}(A - L \cdot U)$ equations and $\text{nnz}(L + U)$ variables.
- We may want to compute the values in L, U such that $R = A - L \cdot U = 0|_{\mathcal{S}}$, the approximation being exact in the locations included in \mathcal{S} , but not outside!
- This is the underlying idea of Edmond Chow’s parallel ILU algorithm¹:

$$F(L, U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i > j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i \leq j \end{cases}$$

- Converges in the asymptotic sense towards incomplete factors L, U such that $R = A - L \cdot U = 0|_{\mathcal{S}}$

¹Chow and Patel. “Fine-grained Parallel Incomplete LU Factorization”. In: SIAM J. on Sci. Comp. (2015).

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

- This is an optimization problem with $\text{nnz}(A - L \cdot U)$ equations and $\text{nnz}(L + U)$ variables.
- We may want to compute the values in L, U such that $R = A - L \cdot U = 0|_{\mathcal{S}}$, the approximation being exact in the locations included in \mathcal{S} , but not outside!
- This is the underlying idea of Edmond Chow’s parallel ILU algorithm¹:

$$F(L, U) = \begin{cases} \frac{1}{u_{jj}} \left(a_{ij} - \sum_{k=1}^{j-1} l_{ik} u_{kj} \right), & i > j \\ a_{ij} - \sum_{k=1}^{i-1} l_{ik} u_{kj}, & i \leq j \end{cases}$$

- We may not need high accuracy here, because we may change the pattern again...
- One single fixed-point sweep.

Fixed-point sweep approximates incomplete factors.

¹Chow and Patel. “Fine-grained Parallel Incomplete LU Factorization”. In: SIAM J. on Sci. Comp. (2015).

Considerations

1. *Select a set of nonzero locations.*
2. *Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.*
3. *Maybe change some locations in favor of locations that result in a better preconditioner.*
4. *Repeat until the preconditioner quality stagnates.*

Fixed-point sweep
approximates
incomplete factors.

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. Maybe change some locations in favor of locations that result in a better preconditioner.
4. Repeat until the preconditioner quality stagnates.

Compute ILU residual & check convergence.

- Comparing sparsity patterns extremely difficult.
- Maybe use the ILU residual as convergence check.

Fixed-point sweep approximates incomplete factors.

Considerations

1. *Select a set of nonzero locations.*
2. *Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.*
3. *Maybe change some locations in favor of locations that result in a better preconditioner.*
4. *Repeat until the preconditioner quality stagnates.*

- The sparsity pattern of A might be a **good initial start** for nonzero locations.

Compute ILU residual & check convergence.

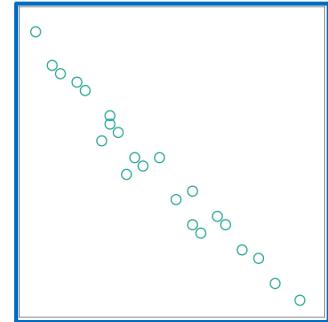
Fixed-point sweep approximates incomplete factors.

Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. **Maybe change some locations in favor of locations that result in a better preconditioner.**
4. Repeat until the preconditioner quality stagnates.

Identify locations with nonzero ILU residual.

Compute ILU residual & check convergence.



- The sparsity pattern of A might be a good initial start for nonzero locations.
- Then, the approximation will be exact for all locations $\mathcal{S}_0 = \mathcal{S}(L_0 + U_0)$ and nonzero in locations $\mathcal{S}_1 = (\mathcal{S}(A) \cup \mathcal{S}(L_0 \cdot U_0)) \setminus \mathcal{S}(L_0 + U_0)$ ¹.

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

Fixed-point sweep approximates incomplete factors.

Considerations

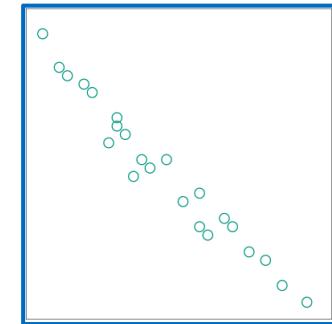
1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. **Maybe change some locations in favor of locations that result in a better preconditioner.**
4. Repeat until the preconditioner quality stagnates.

- The sparsity pattern of A might be a **good initial start** for nonzero locations.
- Then, the approximation will be exact for all locations $\mathcal{S}_0 = \mathcal{S}(L_0 + U_0)$ and nonzero in locations $\mathcal{S}_1 = (\mathcal{S}(A) \cup \mathcal{S}(L_0 \cdot U_0)) \setminus \mathcal{S}(L_0 + U_0)$ ¹.
- Adding all these locations (**level-fill!**) might be good idea...

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix}$$

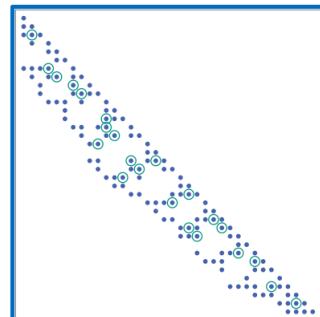
Identify locations with nonzero ILU residual.

Compute ILU residual & check convergence.



Add locations to sparsity pattern of incomplete factors.

Fixed-point sweep approximates incomplete factors.



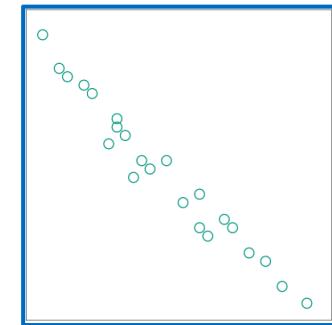
Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. **Maybe change some locations in favor of locations that result in a better preconditioner.**
4. Repeat until the preconditioner quality stagnates.

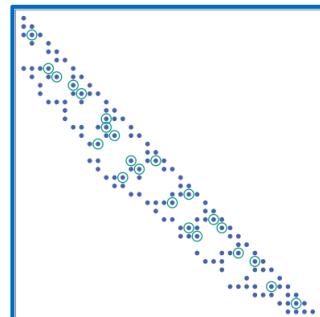
- The sparsity pattern of A might be a **good initial start** for nonzero locations.
- Then, the approximation will be exact for all locations $\mathcal{S}_0 = \mathcal{S}(L_0 + U_0)$ and nonzero in locations $\mathcal{S}_1 = (\mathcal{S}(A) \cup \mathcal{S}(L_0 \cdot U_0)) \setminus \mathcal{S}(L_0 + U_0)$ ¹.
- Adding all these locations (**level-fill!**) might be good idea, **but adding these will again generate new nonzero residuals** $\mathcal{S}_2 = (\mathcal{S}(A) \cup \mathcal{S}(L_1 \cdot U_1)) \setminus \mathcal{S}(L_1 + U_1)$

$$\begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} - \begin{pmatrix} * & & & & \\ * & * & & & \\ * & * & * & & \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} \times \begin{pmatrix} * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \\ * & * & * & * & * \end{pmatrix} = \begin{pmatrix} * & * & * & * & * & * \\ * & * & * & * & * & * \\ * & * & * & * & * & * \\ * & * & * & * & * & * \\ * & * & * & * & * & * \\ * & * & * & * & * & * \end{pmatrix}$$

Fixed-point sweep approximates incomplete factors.



Add locations to sparsity pattern of incomplete factors.



Identify locations with nonzero ILU residual.

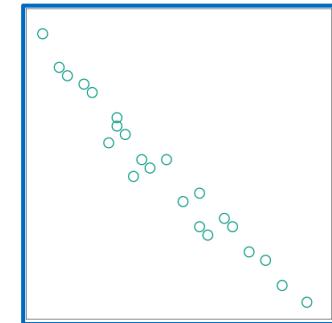
Compute ILU residual & check convergence.

Considerations

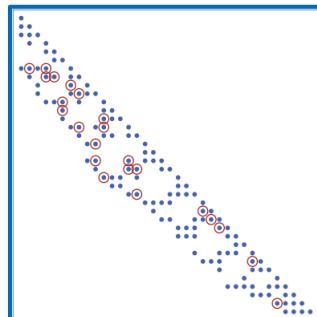
1. Select a set of nonzero locations.
 2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
 3. **Maybe change some locations in favor of locations that result in a better preconditioner.**
 4. Repeat until the preconditioner quality stagnates.
- At some point we should remove some locations again, e.g. the smallest elements, and start over looking at locations $R = A - L_k \cdot U_k \dots$

Identify locations with nonzero ILU residual.

Compute ILU residual & check convergence.



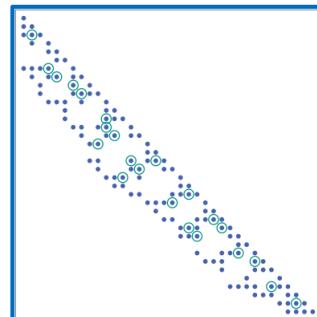
Remove smallest elements from incomplete factors.



Select a threshold separating smallest elements.

Fixed-point sweep approximates incomplete factors.

Add locations to sparsity pattern of incomplete factors.



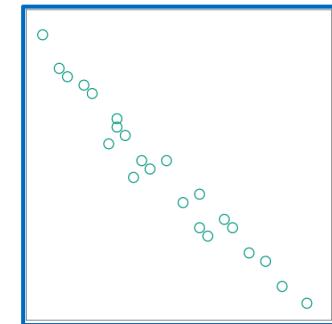
Considerations

1. Select a set of nonzero locations.
2. Compute values in those locations such that $A \approx L \cdot U$ is a “good” approximation.
3. **Maybe change some locations in favor of locations that result in a better preconditioner.**
4. Repeat until the preconditioner quality stagnates.

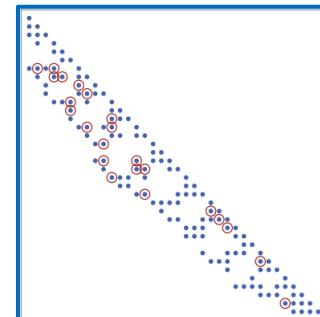
- At some point we should remove some locations again, e.g. the smallest elements, and start over looking at locations $R = A - L_k \cdot U_k \dots$
- We need another sweep, then...

Identify locations with nonzero ILU residual.

Compute ILU residual & check convergence.



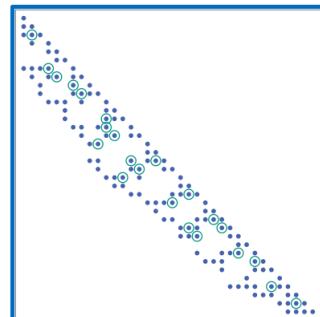
Remove smallest elements from incomplete factors.



Select a threshold separating smallest elements.

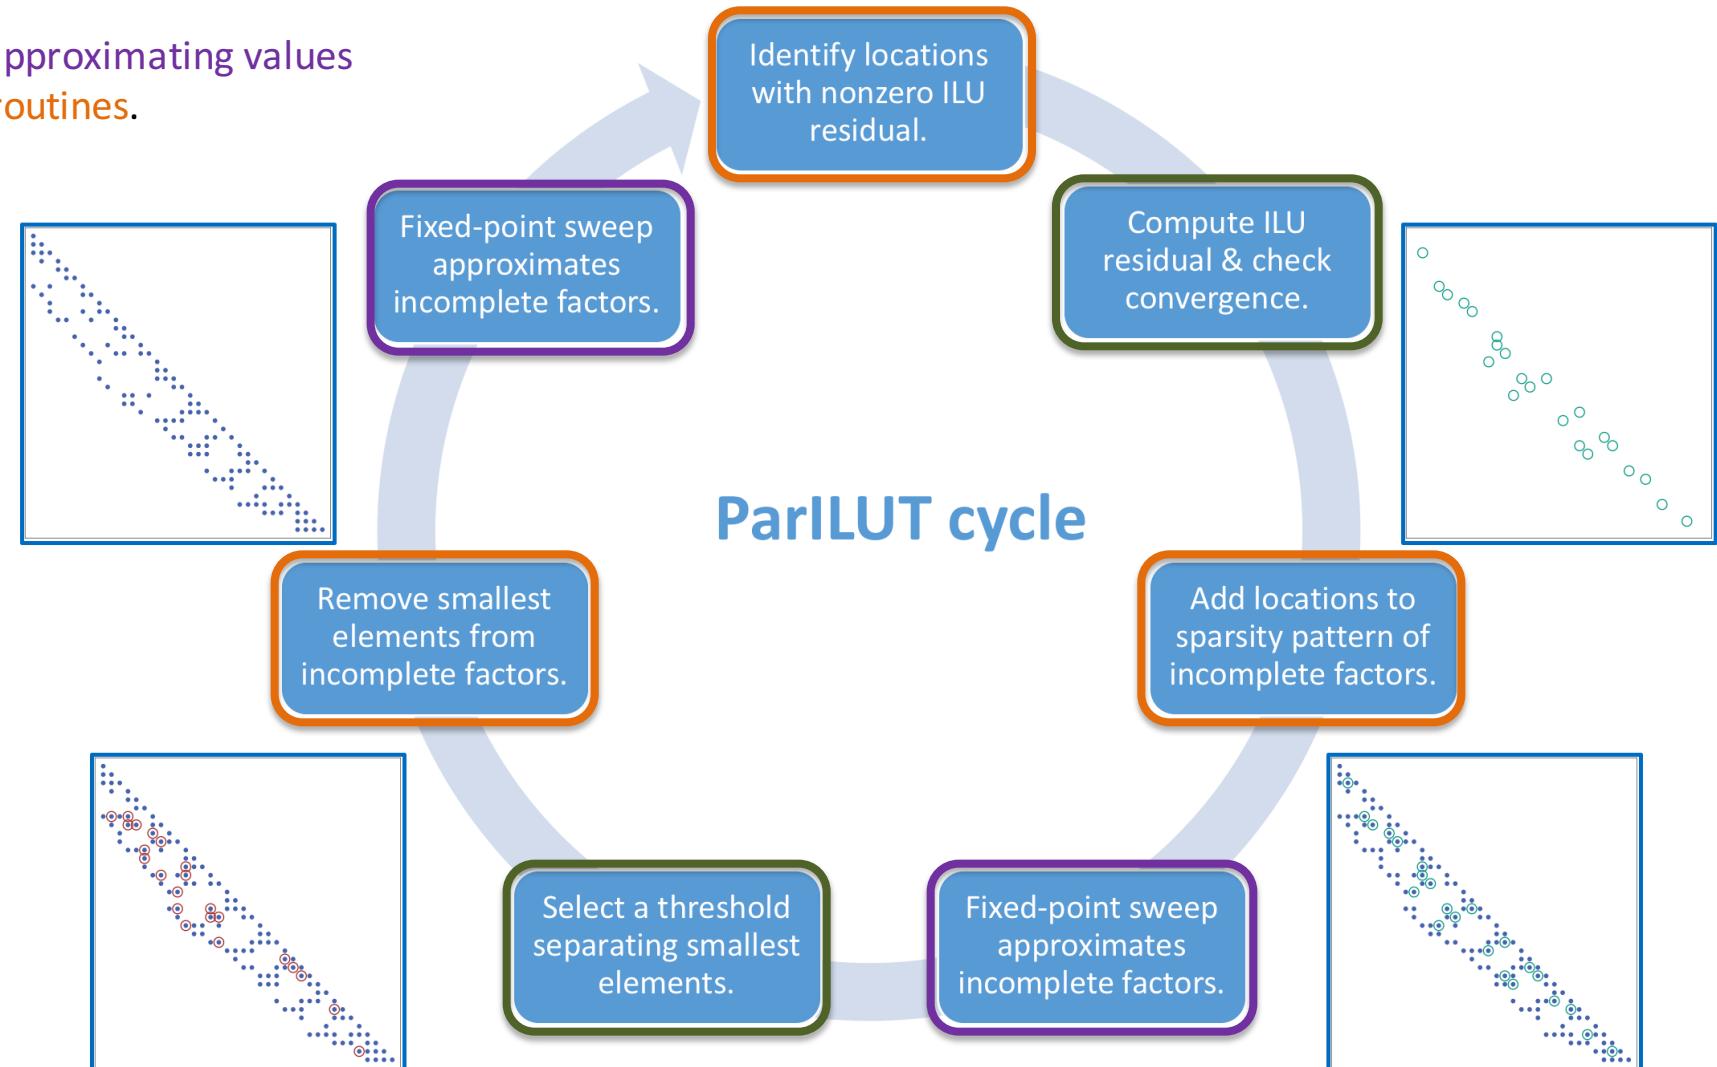
Fixed-point sweep approximates incomplete factors.

Add locations to sparsity pattern of incomplete factors.



ParILUT

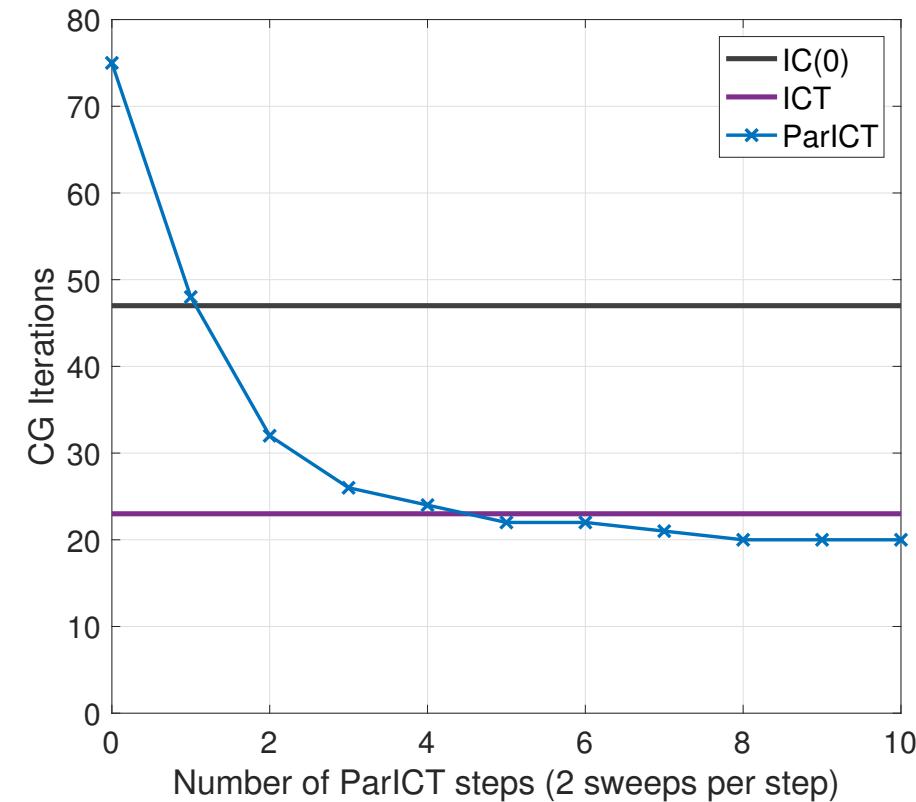
Interleaving fixed-point sweeps approximating values with pattern-changing symbolic routines.



¹Anzt et al. “ParILUT – A new parallel threshold ILU”. In: SIAM J. on Sci. Comp. (2018).

ParILUT quality

Anisotropic fluid flow problem
n: 741, nz: 4,951

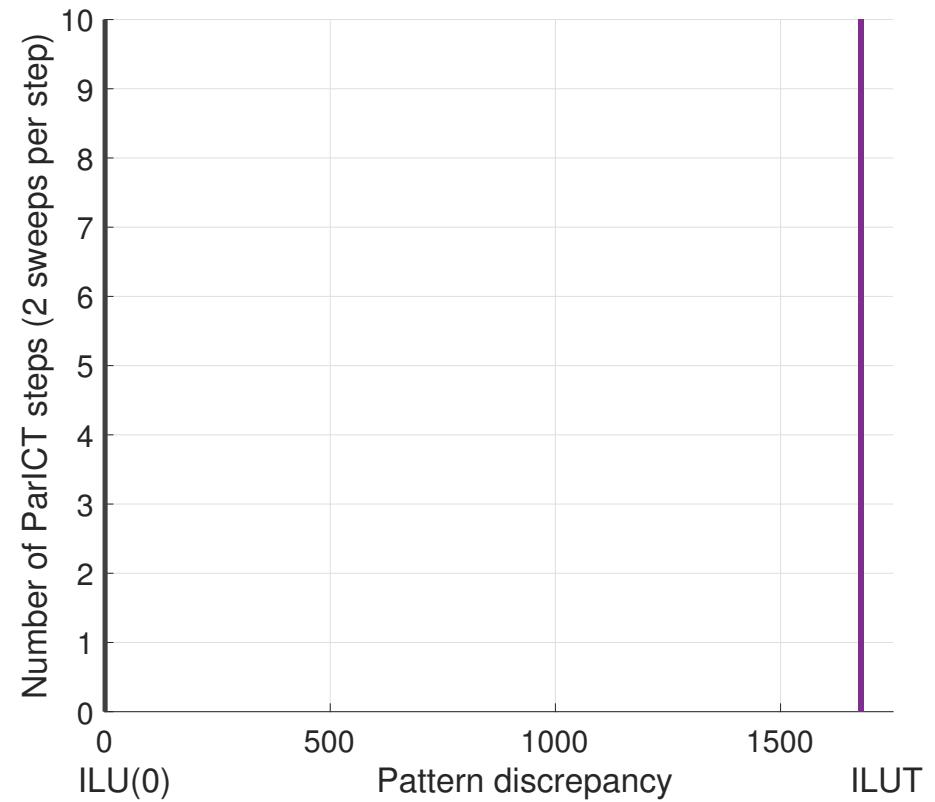
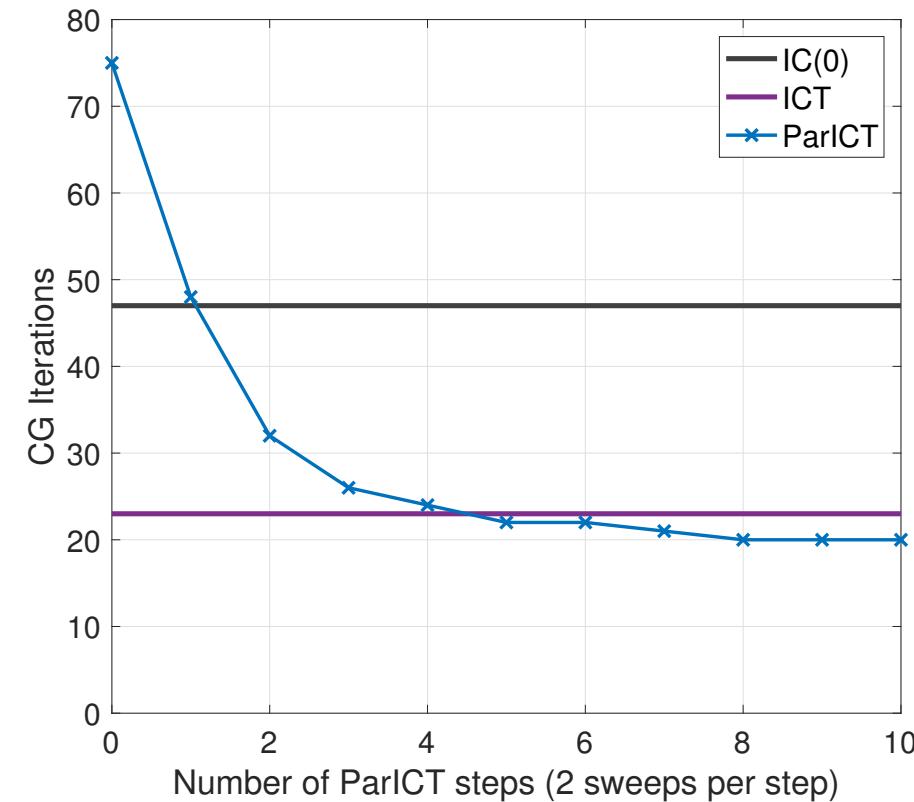


- Top-level solver iterations as quality metric.
- Few sweeps give a “better” preconditioner than ILU(0).
- Better than ILUT?

¹Anzt et al. “ParILUT – A new parallel threshold ILU”. In: SIAM J. on Sci. Comp. (2018).

ParILUT quality

Anisotropic fluid flow problem
n: 741, nz: 4,951

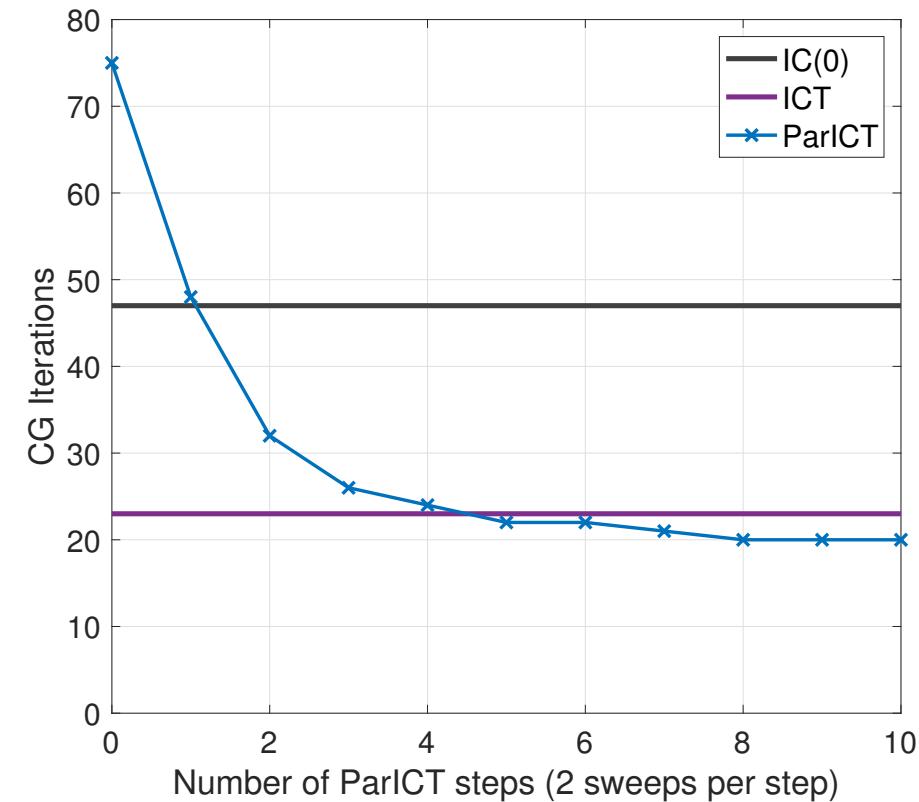


- Top-level solver iterations as quality metric.
- Few sweeps give a “better” preconditioner than ILU(0).
- Better than ILUT?

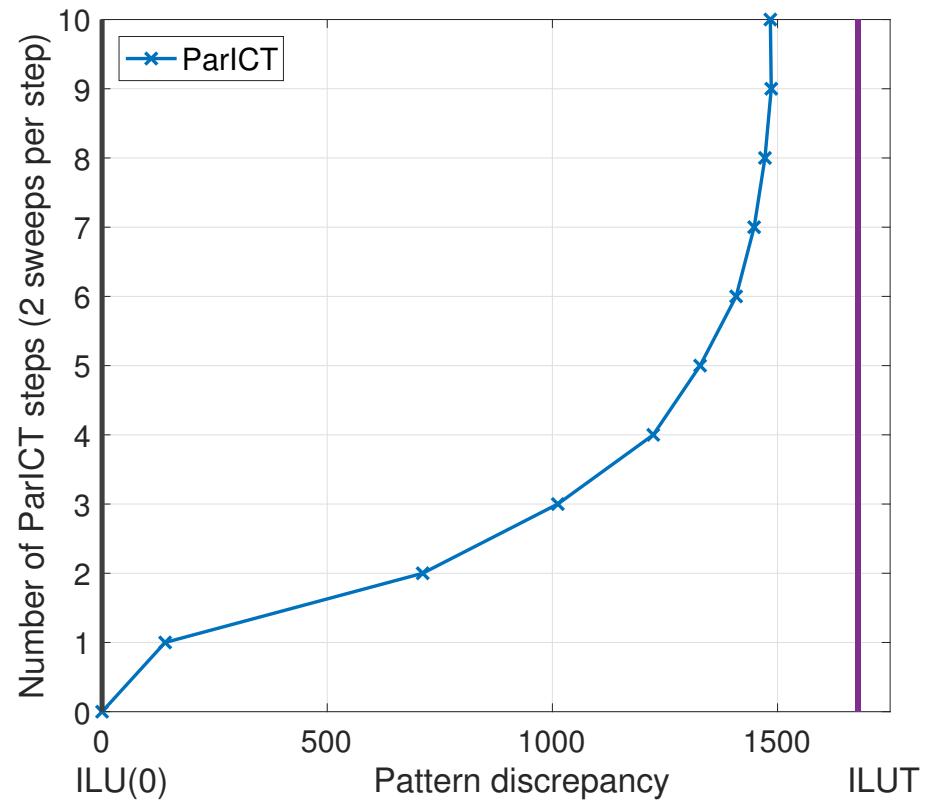
¹Anzt et al. “ParILUT – A new parallel threshold ILU”. In: SIAM J. on Sci. Comp. (2018).

ParILUT quality

Anisotropic fluid flow problem
n: 741, nz: 4,951



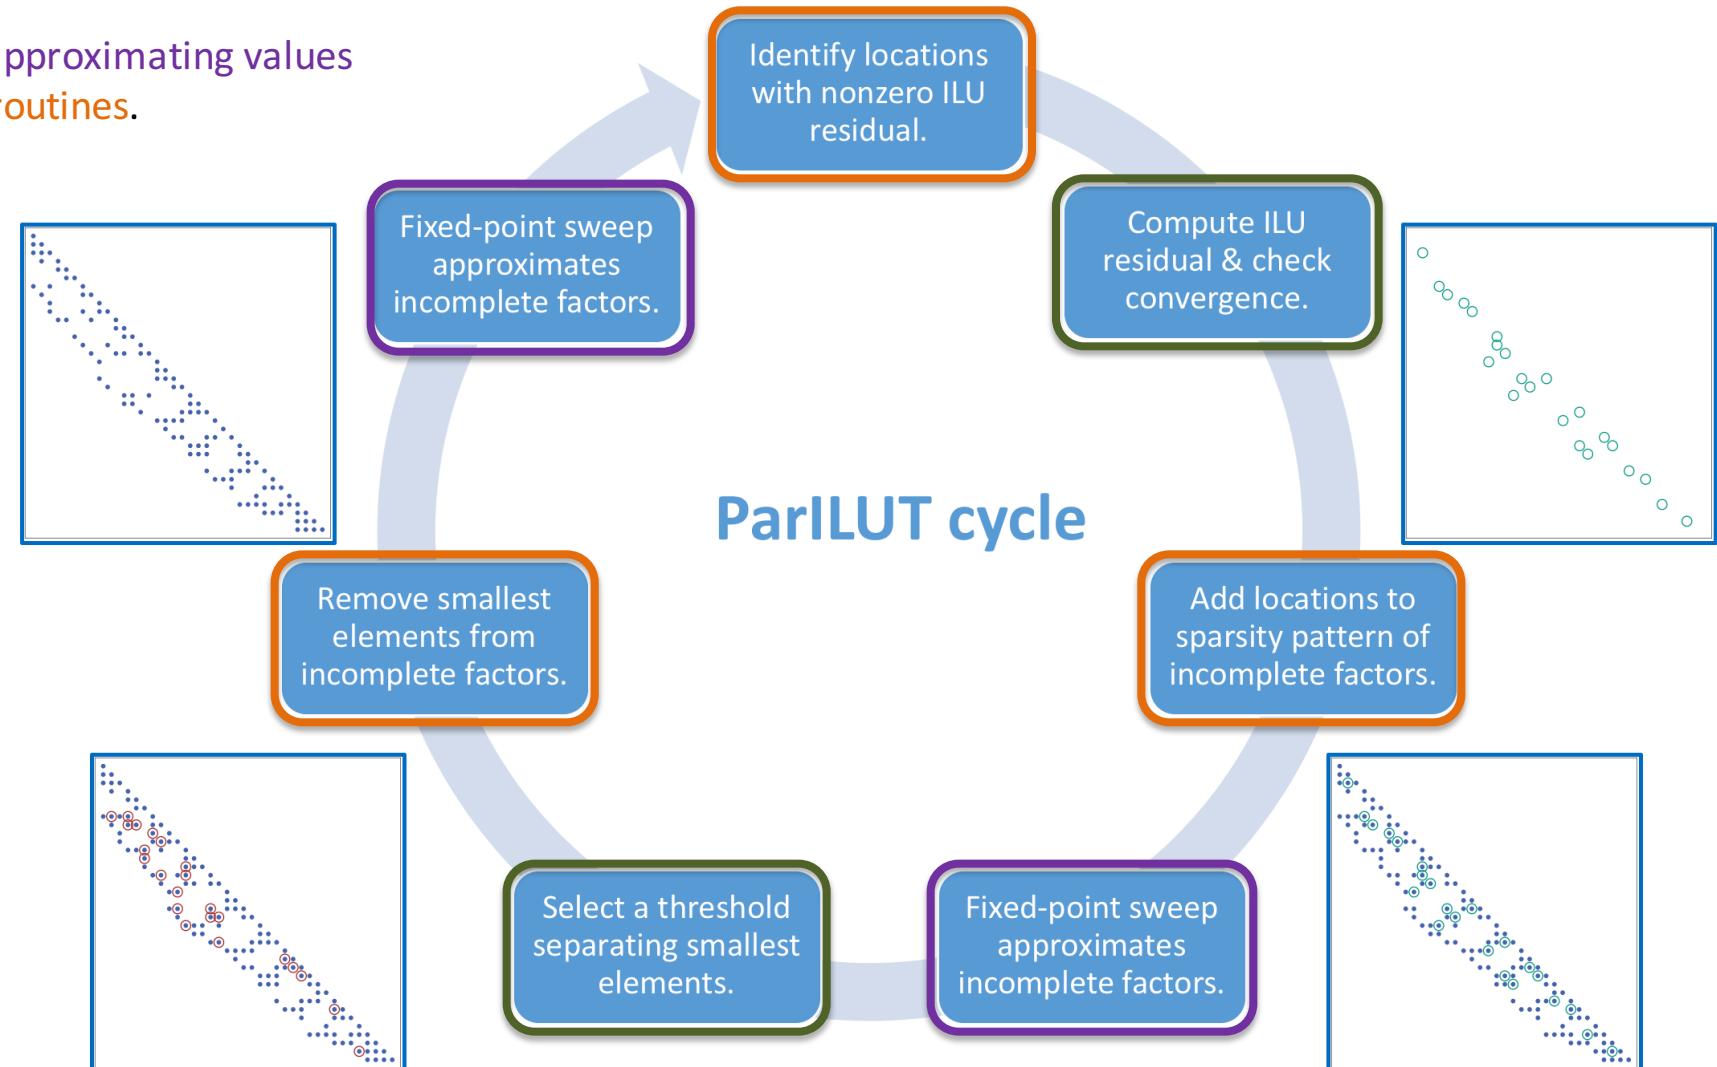
- Top-level solver iterations as quality metric.
- Few sweeps give a “better” preconditioner than ILU(0).
- Better than ILUT?



- Pattern stagnates after few sweeps.
- Pattern “more like” ILUT than ILU(0).

ParILUT

Interleaving fixed-point sweeps approximating values with pattern-changing symbolic routines.

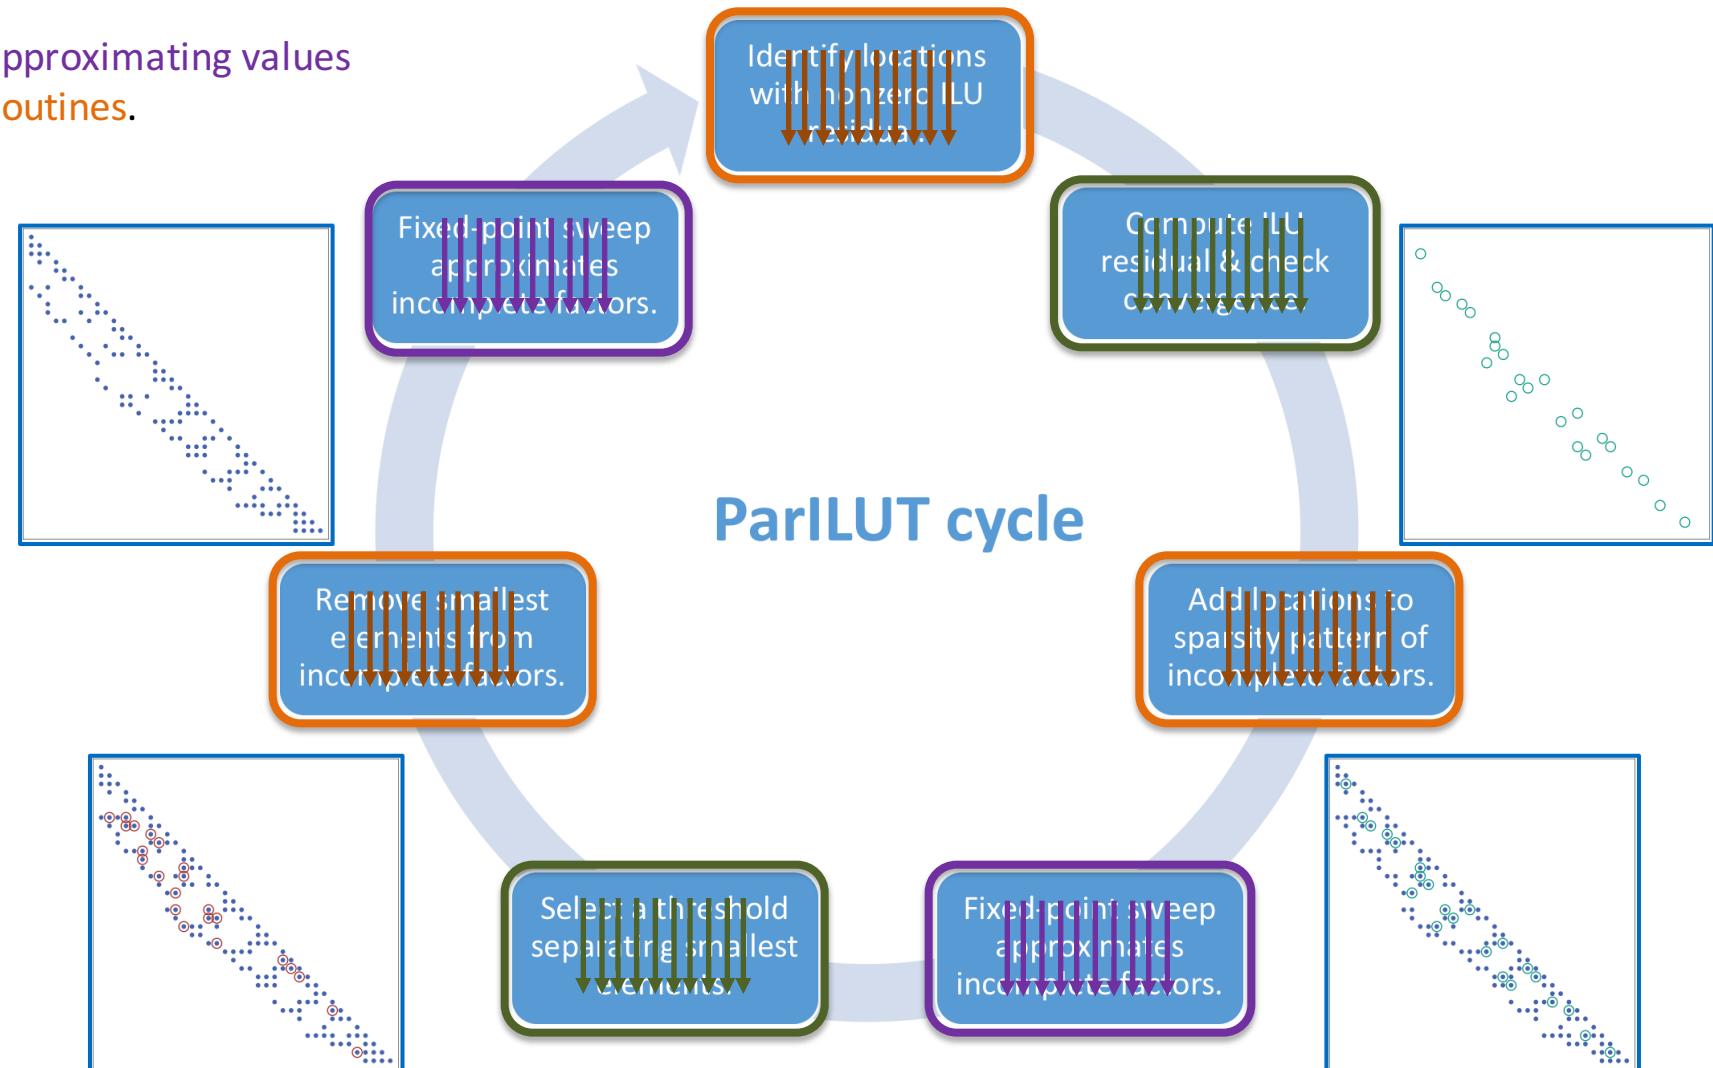


¹Anzt et al. "ParILUT – A new parallel threshold ILU". In: SIAM J. on Sci. Comp. (2018).

ParILUT – a parallel threshold ILU

Parallelism inside the building blocks.

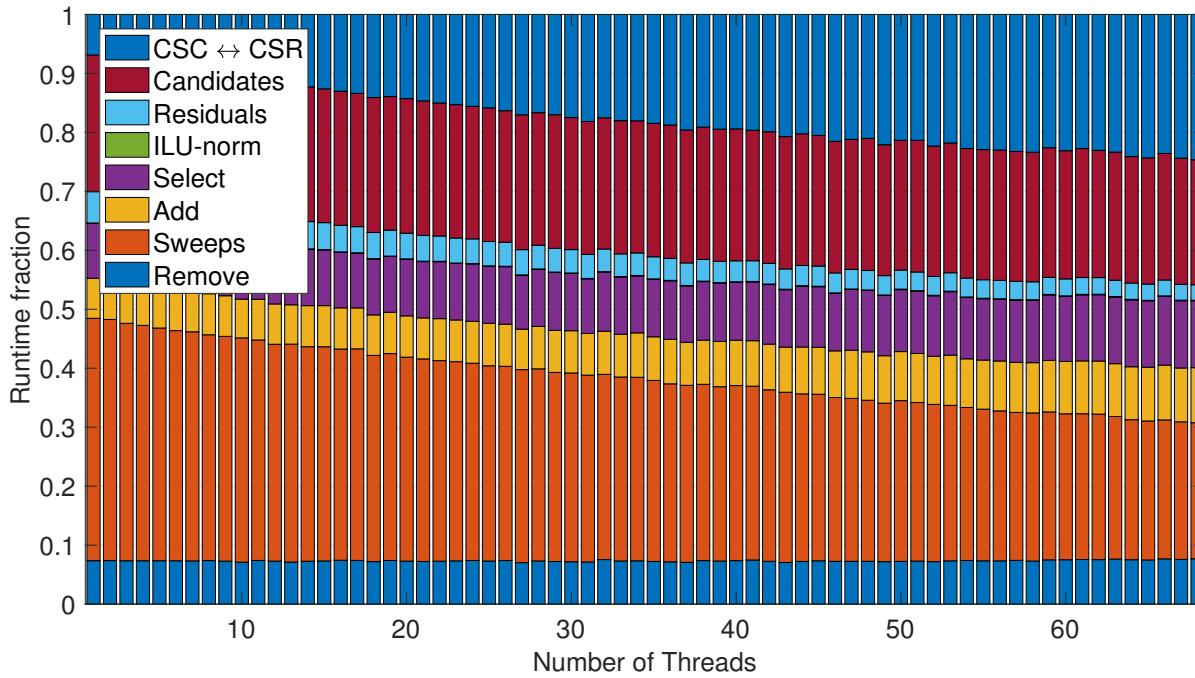
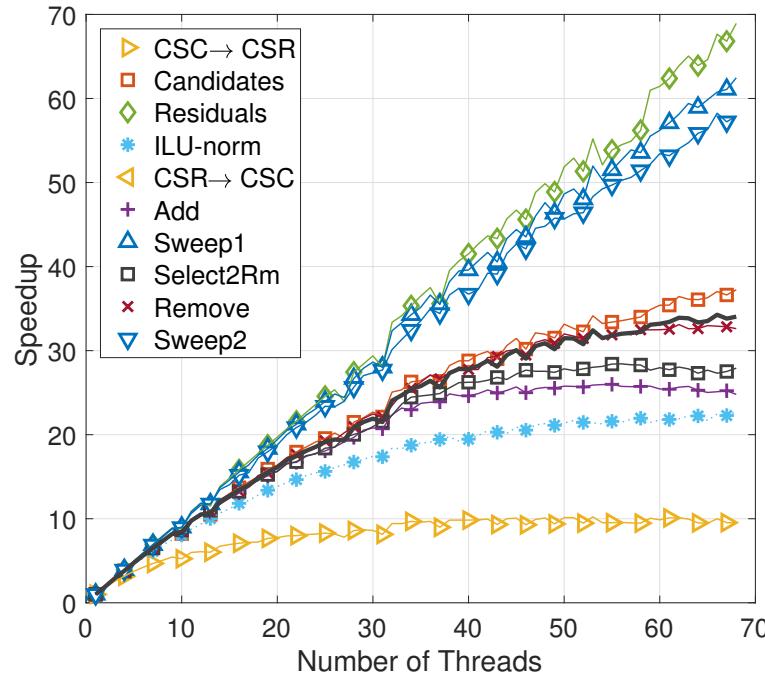
Interleaving fixed-point sweeps approximating values with pattern-changing symbolic routines.



Scalability

Intel Xeon Phi 7250 "Knights Landing"
68 cores @1.40 GHz,
16GB MCDRAM @490 GB/s

thermal2 matrix from SuiteSparse, RCM ordering, 8 el/row.

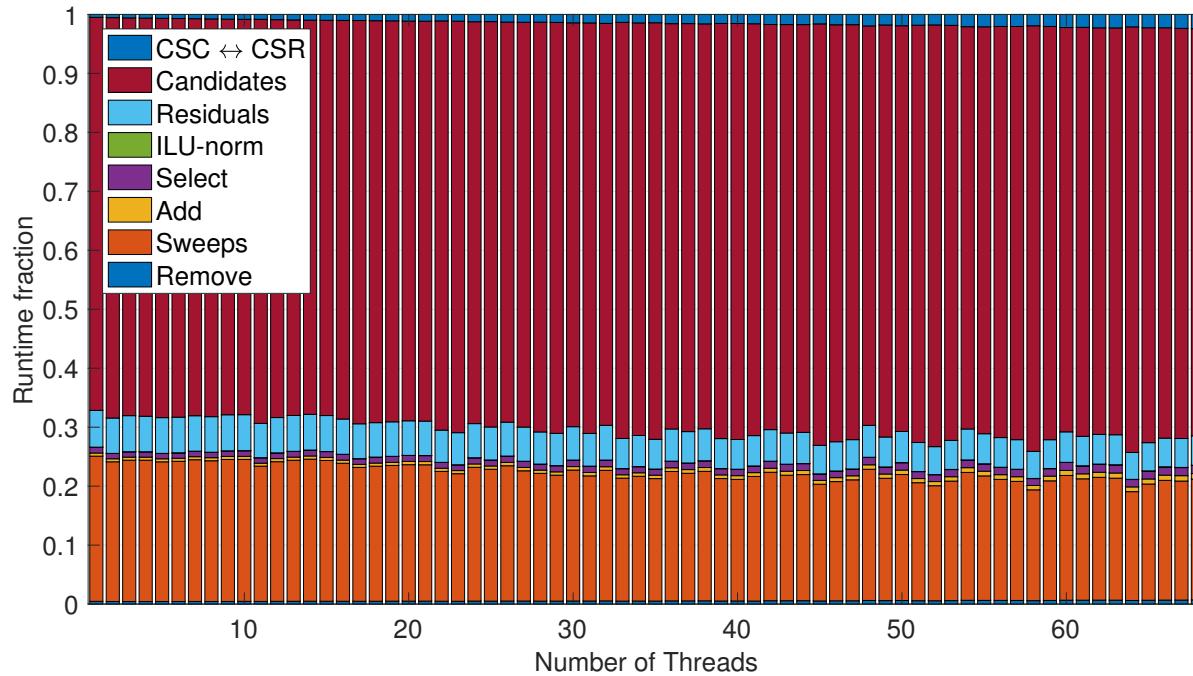
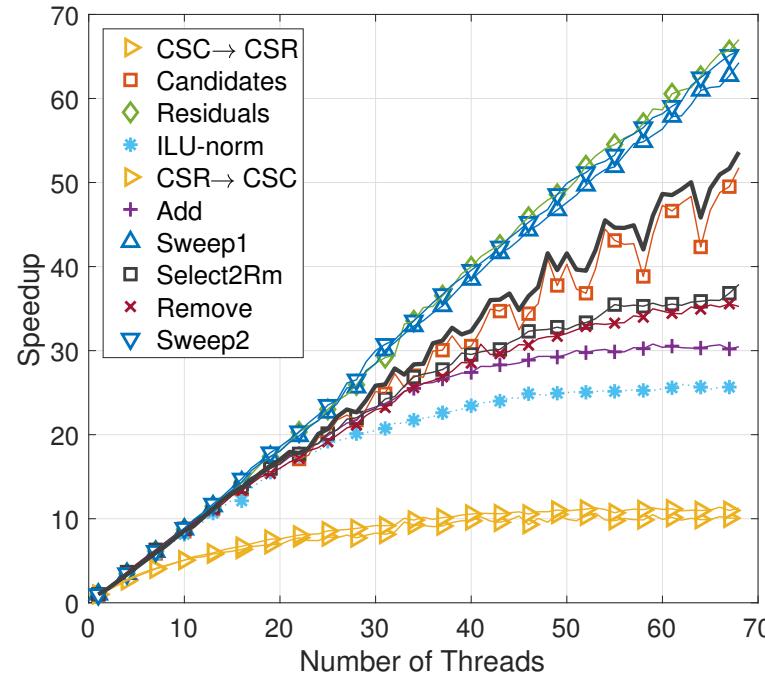


- Building blocks scale with 15% - 100% parallel efficiency.
- Transposition and sort are the bottlenecks.
- Overall speedup ~35x when using 68 KNL cores.

Scalability

Intel Xeon Phi 7250 "Knights Landing"
68 cores @1.40 GHz,
16GB MCDRAM @490 GB/s

topopt120 matrix from topology optimization, 67 el/row.



- Building blocks scale with 15% - 100% parallel efficiency.
- Dominated by candidate search.
- Overall speedup ~52x when using 68 KNL cores.

¹Anzt et al. "ParILUT – A new parallel threshold ILU". In: SIAM J. on Sci. Comp. (2018).

Performance

Intel Xeon Phi 7250 "Knights Landing"
68 cores @1.40 GHz,
16GB MCDRAM @490 GB/s

Runtime of 5 ParILUT / ParICT steps and speedup over SuperLU ILUT*.

| Matrix | Origin | Rows | Nonzeros | Ratio | | | SuperLU | ParILUT | | ParICT | |
|--------------|------------------------------|-----------|-----------|-------|--|--|----------|---------|--------|--------|--------|
| ani7 | 2D Anisotropic Diffusion | 203,841 | 1,407,811 | 6.91 | | | 10.48 s | 0.45 s | 23.34 | 0.30 s | 35.16 |
| apache2 | Suite Sparse Matrix Collect. | 715,176 | 4,817,870 | 6.74 | | | 62.27 s | 1.24 s | 50.22 | 0.65 s | 95.37 |
| cage11 | Suite Sparse Matrix Collect. | 39,082 | 559,722 | 14.32 | | | 60.89 s | 0.54 s | 112.56 | -- | -- |
| jacobianMat9 | Fun3D Fluid Flow Problem | 90,708 | 5,047,042 | 55.64 | | | 153.84 s | 7.26 s | 21.19 | -- | -- |
| thermal2 | Thermal Problem (Suite Sp.) | 1,228,045 | 8,580,313 | 6.99 | | | 91.83 s | 1.23 s | 74.66 | 0.68 s | 134.25 |
| tmt_sym | Suite Sparse Matrix Collect. | 726,713 | 5,080,961 | 6.97 | | | 53.42 s | 0.70 s | 76.21 | 0.41 s | 131.25 |
| topopt120 | Geometry Optimization | 132,300 | 8,802,544 | 66.53 | | | 44.22 s | 14.40 s | 3.07 | 8.24 s | 5.37 |
| torso2 | Suite Sparse Matrix Collect. | 115,967 | 1,033,473 | 8.91 | | | 10.78 s | 0.27 s | 39.92 | -- | -- |
| venkat01 | Suite Sparse Matrix Collect. | 62,424 | 1,717,792 | 27.52 | | | 8.53 s | 0.74 s | 11.54 | -- | -- |

*We thank Sherry Li and Meiyue Shao for technical help in generating the performance numbers.

¹Anzt et al. "ParILUT – A new parallel threshold ILU". In: SIAM J. on Sci. Comp. (2018).

How about GPUs?

ParILUT - A Parallel Threshold ILU for GPUs

Hartwig Anzt*,†, Tobias Ribizel*, Goran Flegar‡, Edmond Chow§, Jack Dongarra†¶||

*Steinbuch Centre for Computing, Karlsruhe Institute of Technology, Germany

†Innovative Computing Lab (ICL), University of Tennessee, Knoxville, USA

‡Departamento de Ingeniería y Ciencia de Computadores, Universidad Jaume I Castellón, Spain

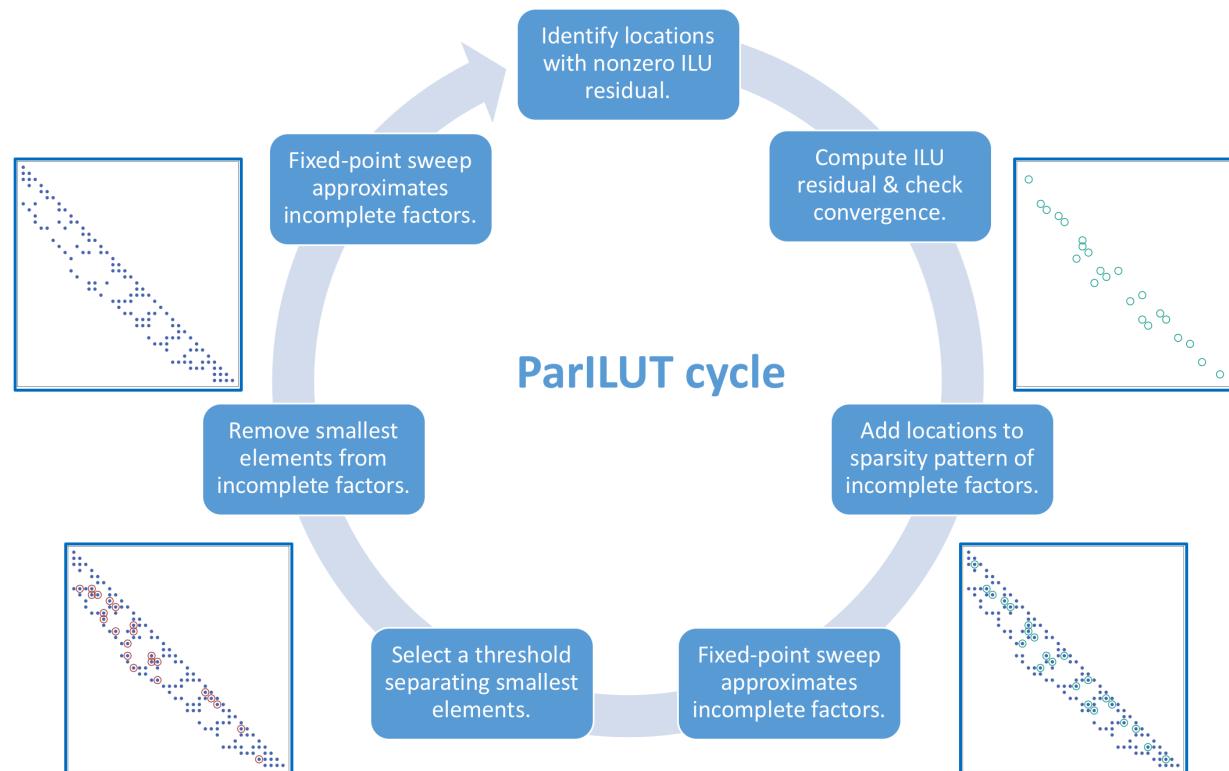
§School of Computational Science and Engineering, Georgia Institute of Technology, USA

¶University of Manchester, Manchester, UK

||Oak Ridge National Lab (ORNL), Oak Ridge, USA

hartwig.anzt@kit.edu, tobias.ribizel@student.kit.edu, flegar@uji.es, echow@cc.gatech.edu, dongarra@icl.utk.edu

Accepted for IPDPS 2019



How about GPUs?

- Fine-grained parallelism
- High bandwidth for coalescent reads
- No deep cache hierarchy
- We need to oversubscribe cores for hiding latency

Part of the ParILUT algorithm requires selecting the smallest k values for removal.

Selection and Sorting algorithms very inefficient on GPUs...

ParILUT - A Parallel Threshold ILU for GPUs

Hartwig Anzt^{*†}, Tobias Ribizel^{*}, Goran Flegar[‡], Edmond Chow[§], Jack Dongarra^{¶||}

^{*}Steinbuch Centre for Computing, Karlsruhe Institute of Technology, Germany

[†]Innovative Computing Lab (ICL), University of Tennessee, Knoxville, USA

[‡]Departamento de Ingeniería y Ciencia de Computadores, Universidad Jaume I Castellón, Spain

[§]School of Computational Science and Engineering, Georgia Institute of Technology, USA

[¶]University of Manchester, Manchester, UK

^{||}Oak Ridge National Lab (ORNL), Oak Ridge, USA

hartwig.anzt@kit.edu, tobias.ribizel@student.kit.edu, flegar@uji.es, echow@cc.gatech.edu, dongarra@icl.uk.ac.uk

Accepted for IPDPS 2019

How about GPUs?

- Fine-grained parallelism
- High bandwidth for coalescent reads
- No deep cache hierarchy
- We need to oversubscribe cores for hiding latency

ParILUT - A Parallel Threshold ILU for GPUs

Hartwig Anzt^{*†}, Tobias Ribizel^{*}, Goran Flegar[‡], Edmond Chow[§], Jack Dongarra^{¶||}

^{*}Steinbuch Centre for Computing, Karlsruhe Institute of Technology, Germany

[†]Innovative Computing Lab (ICL), University of Tennessee, Knoxville, USA

[‡]Departamento de Ingeniería y Ciencia de Computadores, Universidad Jaume I Castellón, Spain

[§]School of Computational Science and Engineering, Georgia Institute of Technology, USA

[¶]University of Manchester, Manchester, UK

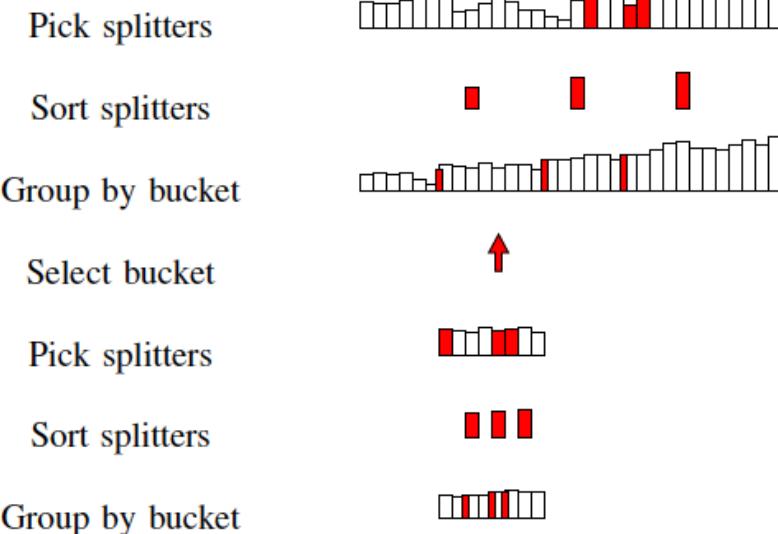
^{||}Oak Ridge National Lab (ORNL), Oak Ridge, USA

hartwig.anzt@kit.edu, tobias.ribizel@student.kit.edu, flegar@uji.es, echow@cc.gatech.edu, dongarra@icl.uk.ac.uk

Accepted for IPDPS 2019

Part of the ParILUT algorithm requires selecting the smallest k values for removal.

SampleSelect:



How about GPUs?

- Fine-grained parallelism
- High bandwidth for coalescent reads
- No deep cache hierarchy
- We need to oversubscribe cores for hiding latency

ParILUT - A Parallel Threshold ILU for GPUs

Hartwig Anzt^{*†}, Tobias Ribizel^{*}, Goran Flegar[‡], Edmond Chow[§], Jack Dongarra^{¶||}

^{*}Steinbuch Centre for Computing, Karlsruhe Institute of Technology, Germany

[†]Innovative Computing Lab (ICL), University of Tennessee, Knoxville, USA

[‡]Departamento de Ingeniería y Ciencia de Computadores, Universidad Jaume I Castellón, Spain

[§]School of Computational Science and Engineering, Georgia Institute of Technology, USA

[¶]University of Manchester, Manchester, UK

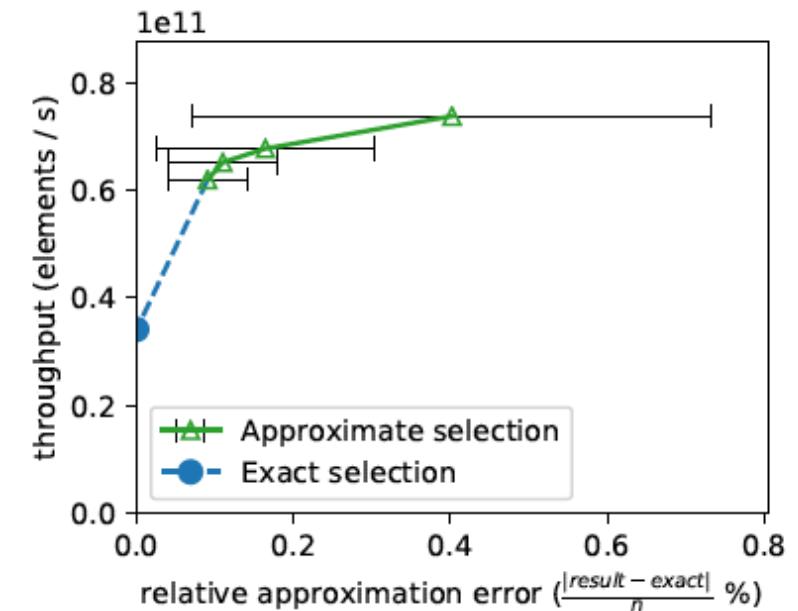
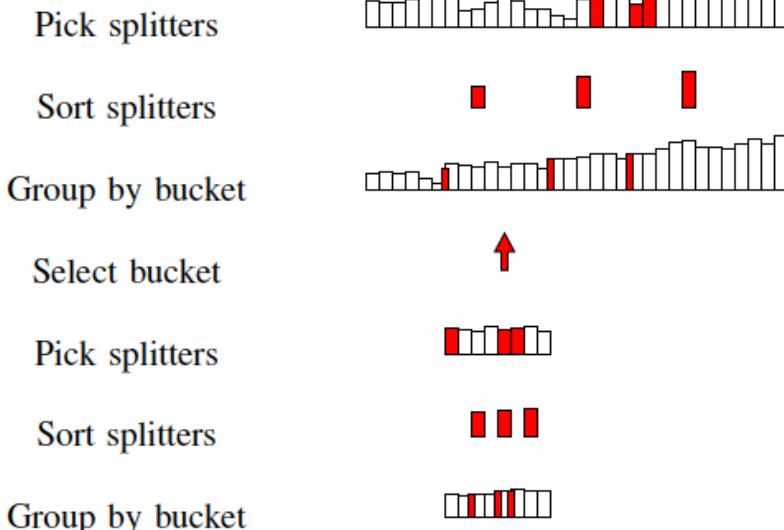
^{||}Oak Ridge National Lab (ORNL), Oak Ridge, USA

hartwig.anzt@kit.edu, tobias.ribizel@student.kit.edu, flegar@uji.es, echow@cc.gatech.edu, dongarra@icl.uk.ac.uk

Accepted for IPDPS 2019

Part of the ParILUT algorithm requires selecting the smallest k values for removal.

SampleSelect:

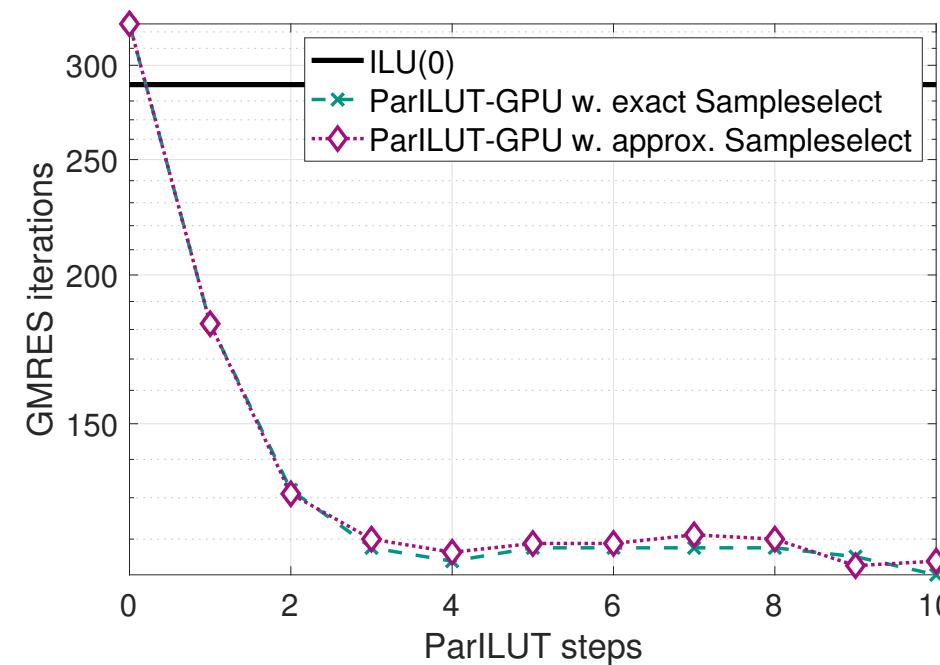


¹Ribizel and Anzt. "Approximate and Exact Selection on GPUs". In: ASHES workshop 2019, accepted.

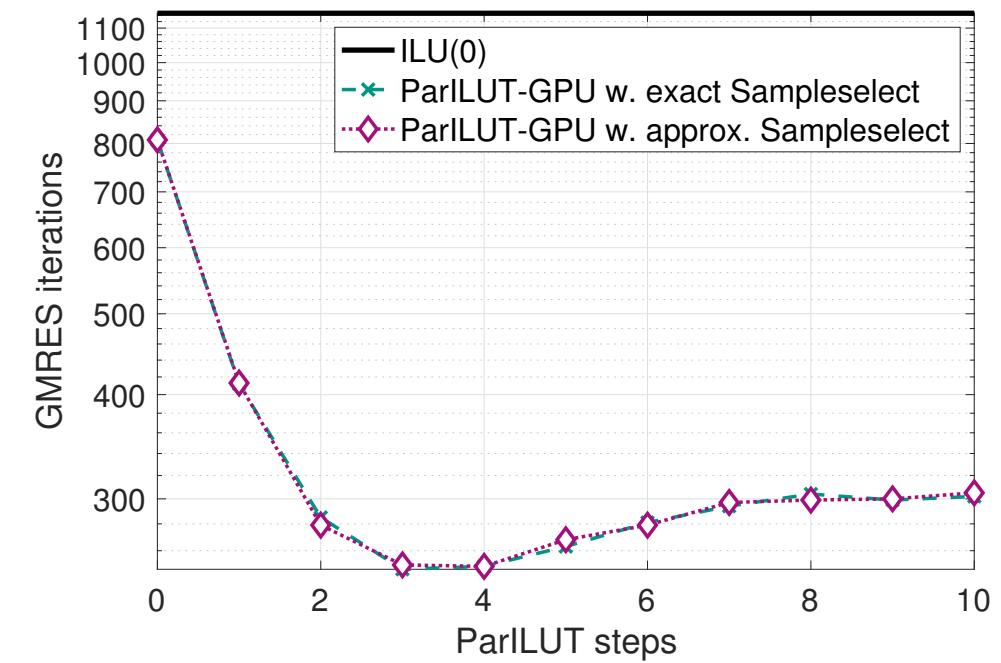
ParILUT - A Parallel Threshold ILU for GPUs

Impact of exact/approximate SampleSelect on ParILUT preconditioner quality

AN15

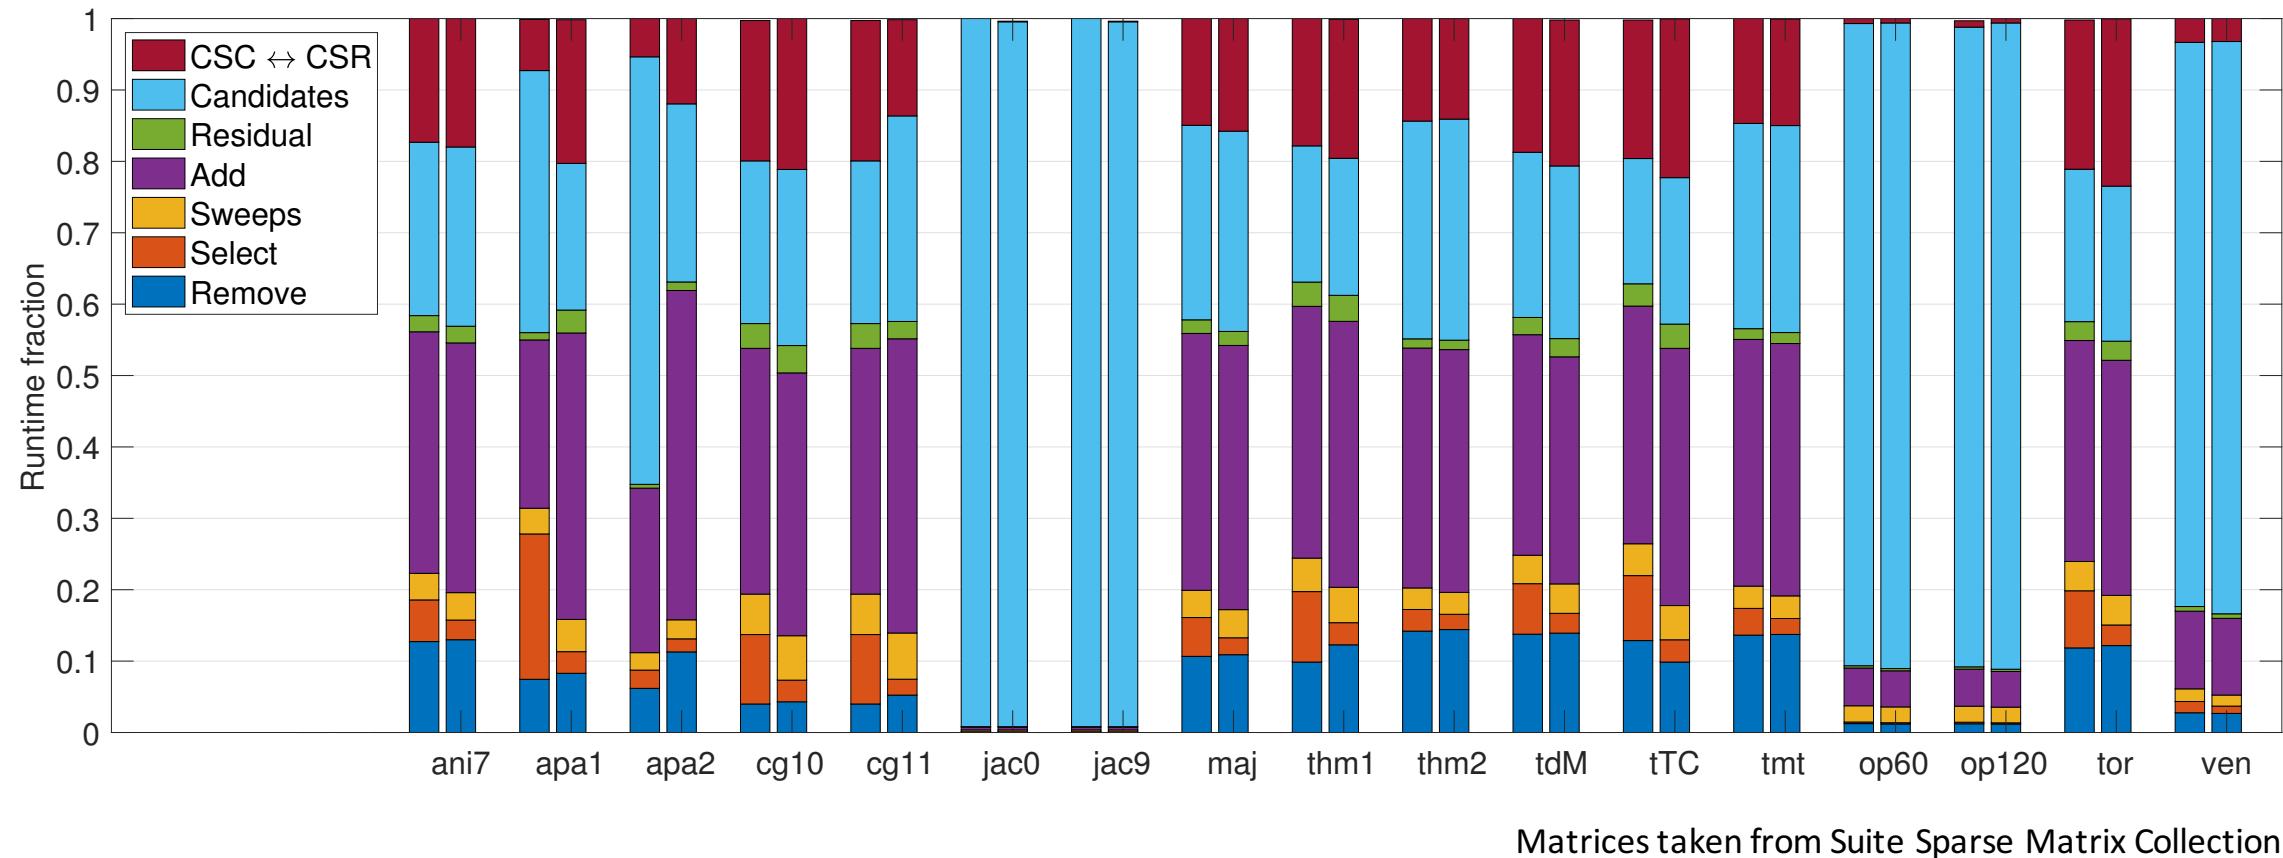


AN16



ParILUT - A Parallel Threshold ILU for GPUs

Impact of exact/approximate SampleSelect on ParILUT runtime breakdown

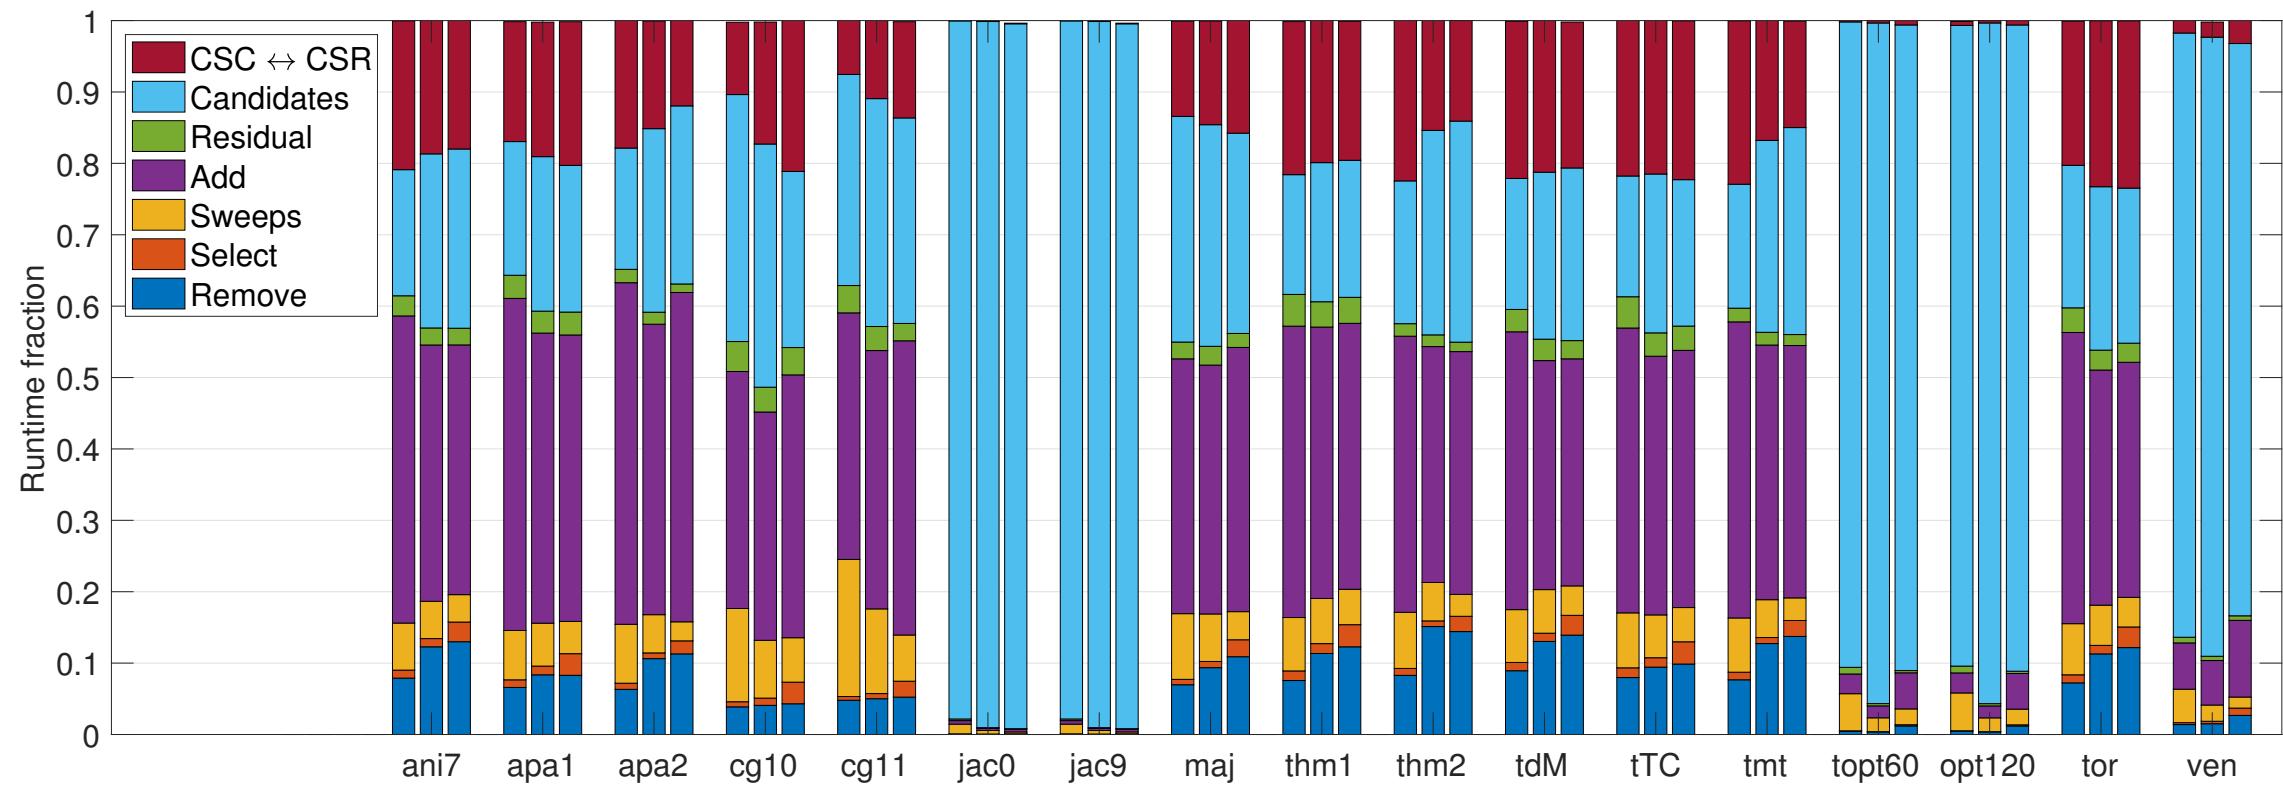


Matrices taken from Suite Sparse Matrix Collection.

¹Anzt et al. "ParILUT – A parallel threshold ILU for GPUs". In: IPDPS 2019, accepted.

ParILUT - A Parallel Threshold ILU for GPUs

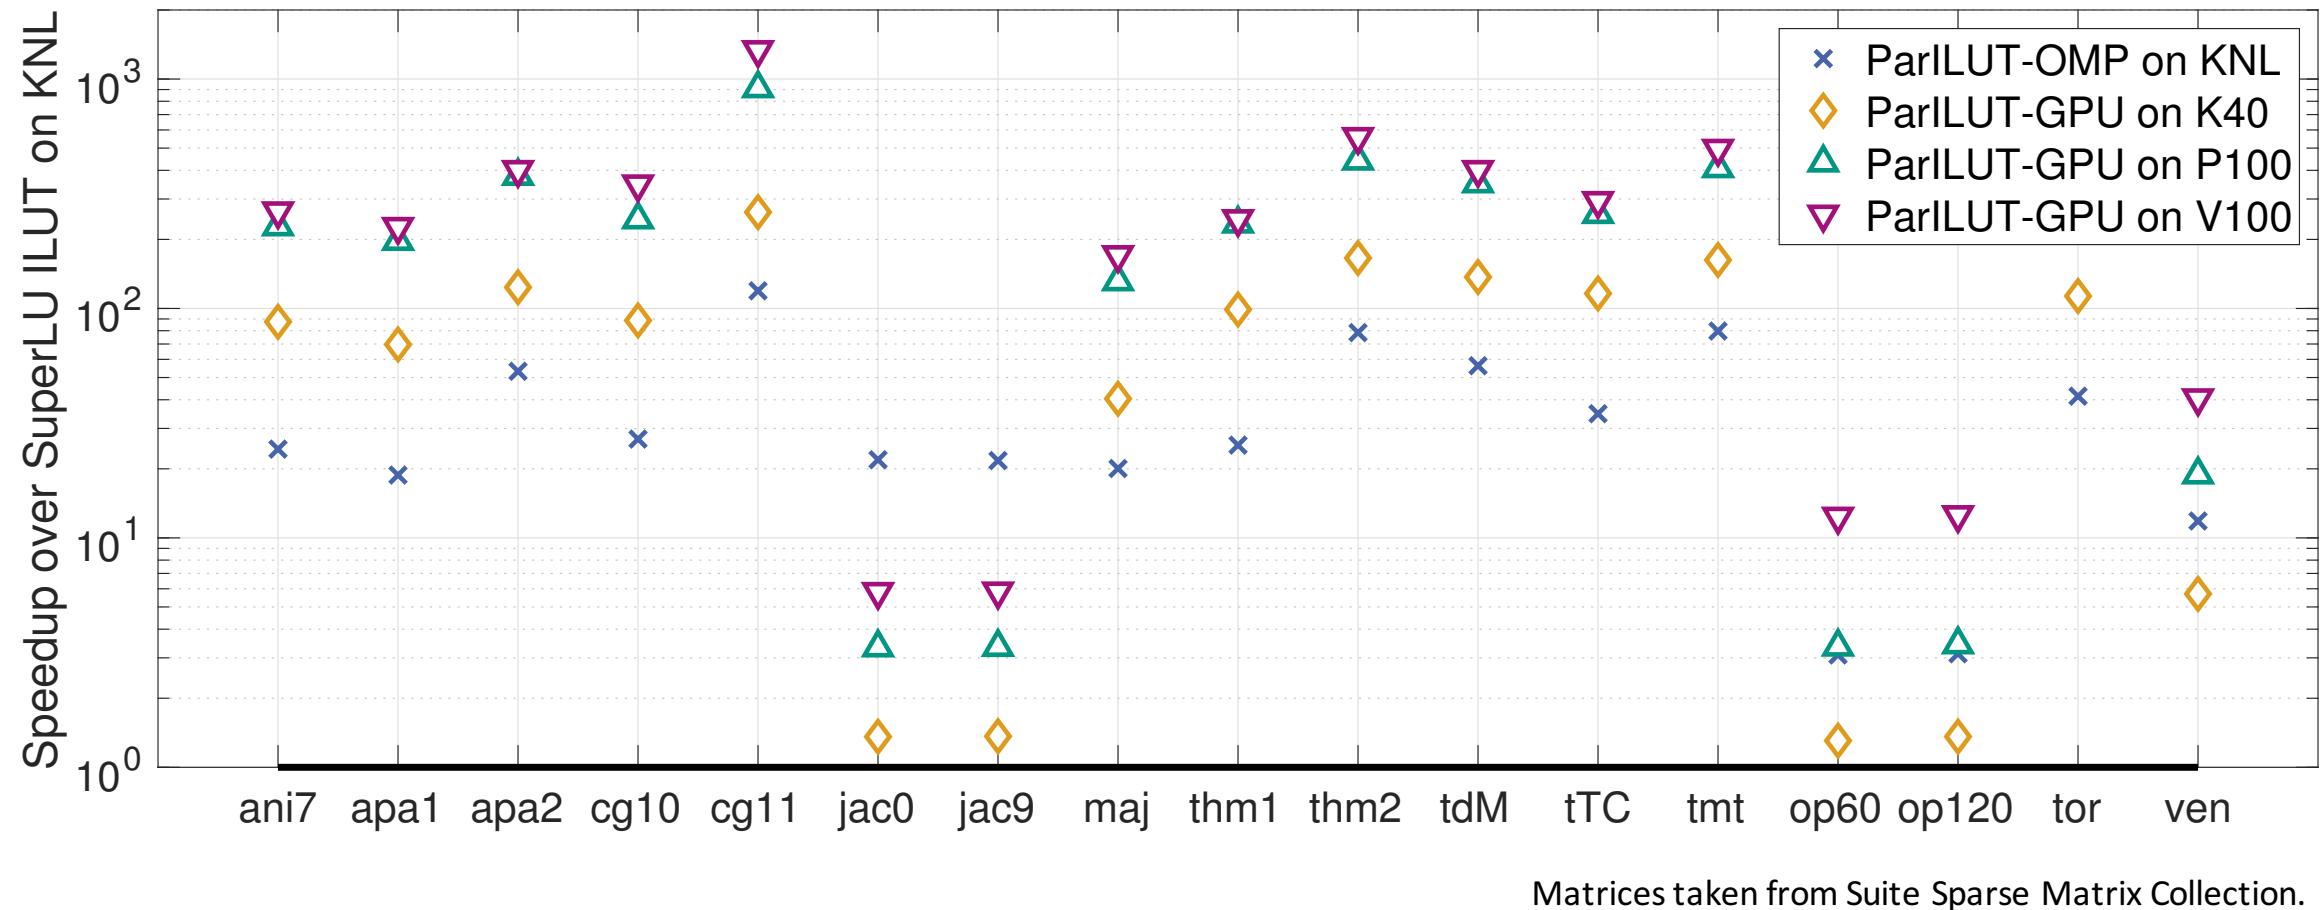
ParILUT performance across different GPU generations:
1st bar: NVIDIA K40
2nd bar: NVIDIA P100
3rd bar: NVIDIA V100



Matrices taken from Suite Sparse Matrix Collection.

¹Anzt et al. "ParILUT – A parallel threshold ILU for GPUs". In: IPDPS 2019, accepted.

ParILUT Performance across architectures



Matrices taken from Suite Sparse Matrix Collection.

¹Anzt et al. "ParILUT – A parallel threshold ILU for GPUs". In: IPDPS 2019, accepted.

Is this a future-oriented algorithm?

- **Hybrid ParILUT** version utilizing GPU and CPU, overlapping communication & computation.
- **Asynchronous** version relaxing dependencies.
- Use a **different sparsity-pattern generator**:
 - Randomized?
 - Machine learning techniques?
- **Increasing fill-in** towards “full” factorization.
- ParILUT routines available in MAGMA-sparse – they will be in Ginkgo.



THE UNIVERSITY OF
TENNESSEE
KNOXVILLE



This research was sponsored by:



The Exascale Computing Project

A Collaborative effort of the U.S. Department of Energy Office of Science And the National Nuclear Security Administration



U.S. DEPARTMENT OF
ENERGY
Office of Science

U.S. Department of Energy
ASCR Award Number DE-SC0016513

HELMHOLTZ
RESEARCH FOR GRAND CHALLENGES

Helmholtz Impuls und Vernetzungsfond
VH-NG-1241

Test matrices

| Matrix | Origin | SPD | Num. Rows | Nz | Nz/Row |
|---------------|----------------------------|-----|-----------|-----------|--------|
| ANI5 | 2D anisotropic diffusion | yes | 12,561 | 86,227 | 6.86 |
| ANI6 | 2D anisotropic diffusion | yes | 50,721 | 349,603 | 6.89 |
| ANI7 | 2D anisotropic diffusion | yes | 203,841 | 1,407,811 | 6.91 |
| APACHE1 | Suite Sparse [10] | yes | 80,800 | 542,184 | 6.71 |
| APACHE2 | Suite Sparse | yes | 715,176 | 4,817,870 | 6.74 |
| CAGE10 | Suite Sparse | no | 11,397 | 150,645 | 13.22 |
| CAGE11 | Suite Sparse | no | 39,082 | 559,722 | 14.32 |
| JACOBIANMAT0 | Fun3D fluid flow [20] | no | 90,708 | 5,047,017 | 55.64 |
| JACOBIANMAT9 | Fun3D fluid flow | no | 90,708 | 5,047,042 | 55.64 |
| MAJORBASIS | Suite Sparse | no | 160,000 | 1,750,416 | 10.94 |
| TOPOPT010 | Geometry optimization [24] | yes | 132,300 | 8,802,544 | 66.53 |
| TOPOPT060 | Geometry optimization | yes | 132,300 | 7,824,817 | 59.14 |
| TOPOPT120 | Geometry optimization | yes | 132,300 | 7,834,644 | 59.22 |
| THERMAL1 | Suite Sparse | yes | 82,654 | 574,458 | 6.95 |
| THERMAL2 | Suite Sparse | yes | 1,228,045 | 8,580,313 | 6.99 |
| THERMOMECH_TC | Suite Sparse | yes | 102,158 | 711,558 | 6.97 |
| THERMOMECH_DM | Suite Sparse | yes | 204,316 | 1,423,116 | 6.97 |
| TMT_SYM | Suite Sparse | yes | 726,713 | 5,080,961 | 6.99 |
| TORSO2 | Suite Sparse | no | 115,967 | 1,033,473 | 8.91 |
| VENKAT01 | Suite Sparse | no | 62,424 | 1,717,792 | 27.52 |

Convergence: GMRES iterations

| Matrix | no prec. | ILU(0) | ILUT | ParILUT | | | | | |
|--------------|----------|--------|------|---------|-----|-----|-----|-----|-----|
| | | | | 0 | 1 | 2 | 3 | 4 | 5 |
| ANI5 | 882 | 172 | 78 | 278 | 161 | 105 | 84 | 74 | 66 |
| ANI6 | 1,751 | 391 | 127 | 547 | 315 | 211 | 168 | 143 | 131 |
| ANI7 | 3,499 | 828 | 290 | 1,083 | 641 | 459 | 370 | 318 | 289 |
| CAGE10 | 20 | 8 | 8 | 9 | 7 | 8 | 8 | 8 | 8 |
| CAGE11 | 21 | 9 | 8 | 9 | 7 | 7 | 7 | 7 | 7 |
| JACOBIANMAT0 | 315 | 40 | 34 | 63 | 36 | 33 | 33 | 33 | 33 |
| JACOBIANMAT9 | 539 | 66 | 65 | 110 | 60 | 55 | 54 | 53 | 53 |
| MAJORBASIS | 95 | 15 | 9 | 26 | 12 | 11 | 11 | 11 | 11 |
| TOPOPT010 | 2,399 | 565 | 303 | 835 | 492 | 375 | 348 | 340 | 339 |
| TOPOPT060 | 2,852 | 666 | 397 | 963 | 584 | 445 | 417 | 412 | 410 |
| TOPOPT120 | 2,765 | 668 | 396 | 959 | 584 | 445 | 416 | 408 | 408 |
| TORSO2 | 46 | 10 | 7 | 18 | 8 | 6 | 7 | 7 | 7 |
| VENKAT01 | 195 | 22 | 17 | 42 | 18 | 17 | 17 | 17 | 17 |

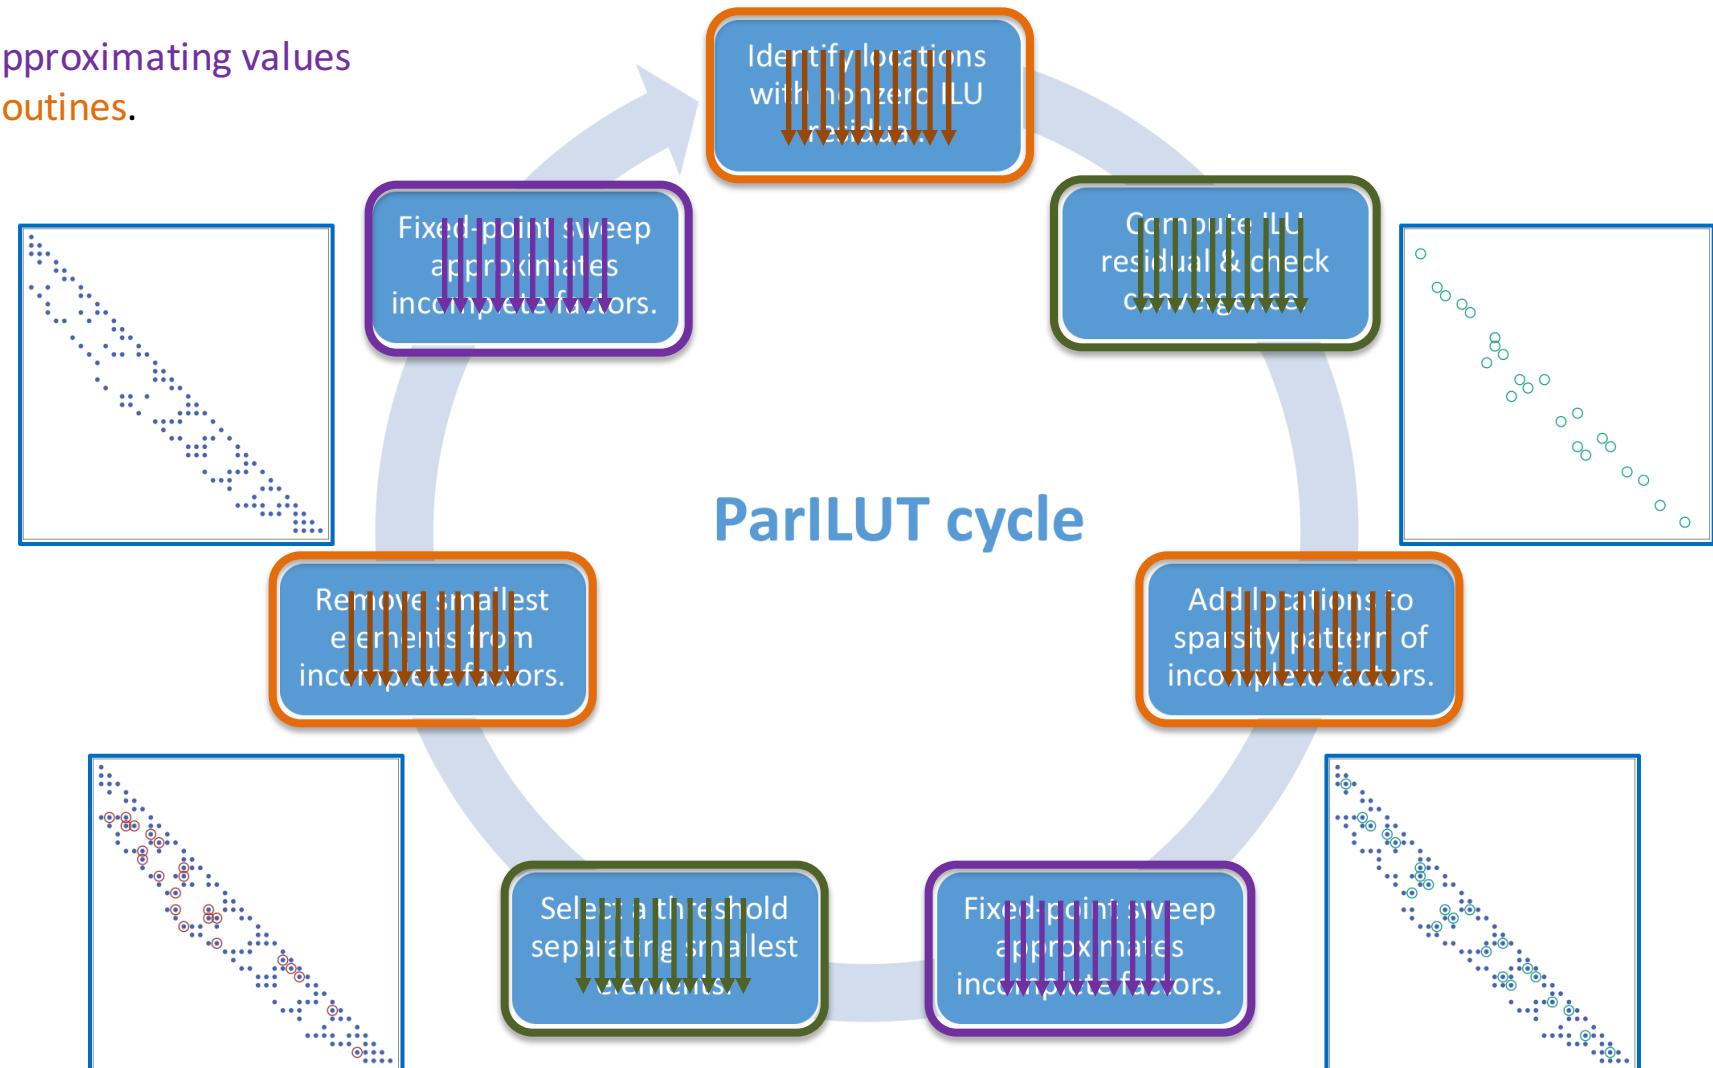
Convergence: CG iterations

| Matrix | no prec. | IC(0) | ICT | ParICT | | | | | |
|---------------|----------|-------|-------|--------|-------|-------|-------|-------|-------|
| | | | | 0 | 1 | 2 | 3 | 4 | 5 |
| ANI5 | 951 | 226 | — | 297 | 184 | 136 | 108 | 93 | 86 |
| ANI6 | 1,926 | 621 | — | 595 | 374 | 275 | 219 | 181 | 172 |
| ANI7 | 3,895 | 1,469 | — | 1,199 | 753 | 559 | 455 | 405 | 377 |
| APACHE1 | 3,727 | 368 | 331 | 1,480 | 933 | 517 | 321 | 323 | 323 |
| APACHE2 | 4,574 | 1,150 | 785 | 1,890 | 1,197 | 799 | 766 | 760 | 754 |
| THERMAL1 | 1,640 | 453 | 412 | 626 | 447 | 409 | 389 | 385 | 383 |
| THERMAL2 | 6,253 | 1,729 | 1,604 | 2,372 | 1,674 | 1,503 | 1,457 | 1,472 | 1,433 |
| THERMOMECH_DM | 21 | 8 | 8 | 8 | 7 | 7 | 7 | 7 | 7 |
| THERMOMECH_TC | 21 | 8 | 7 | 8 | 7 | 7 | 7 | 7 | 7 |
| TMT_SYM | 5,481 | 1,453 | 1,185 | 1,963 | 1,234 | 1,071 | 1,012 | 992 | 1,004 |
| TOPOPT010 | 2,613 | 692 | 331 | 845 | 551 | 402 | 342 | 316 | 313 |
| TOPOPT060 | 3,123 | 871 | — | 988 | 749 | 693 | 1,116 | — | — |
| TOPOPT120 | 3,062 | 886 | — | 991 | 837 | 784 | 2,185 | — | — |

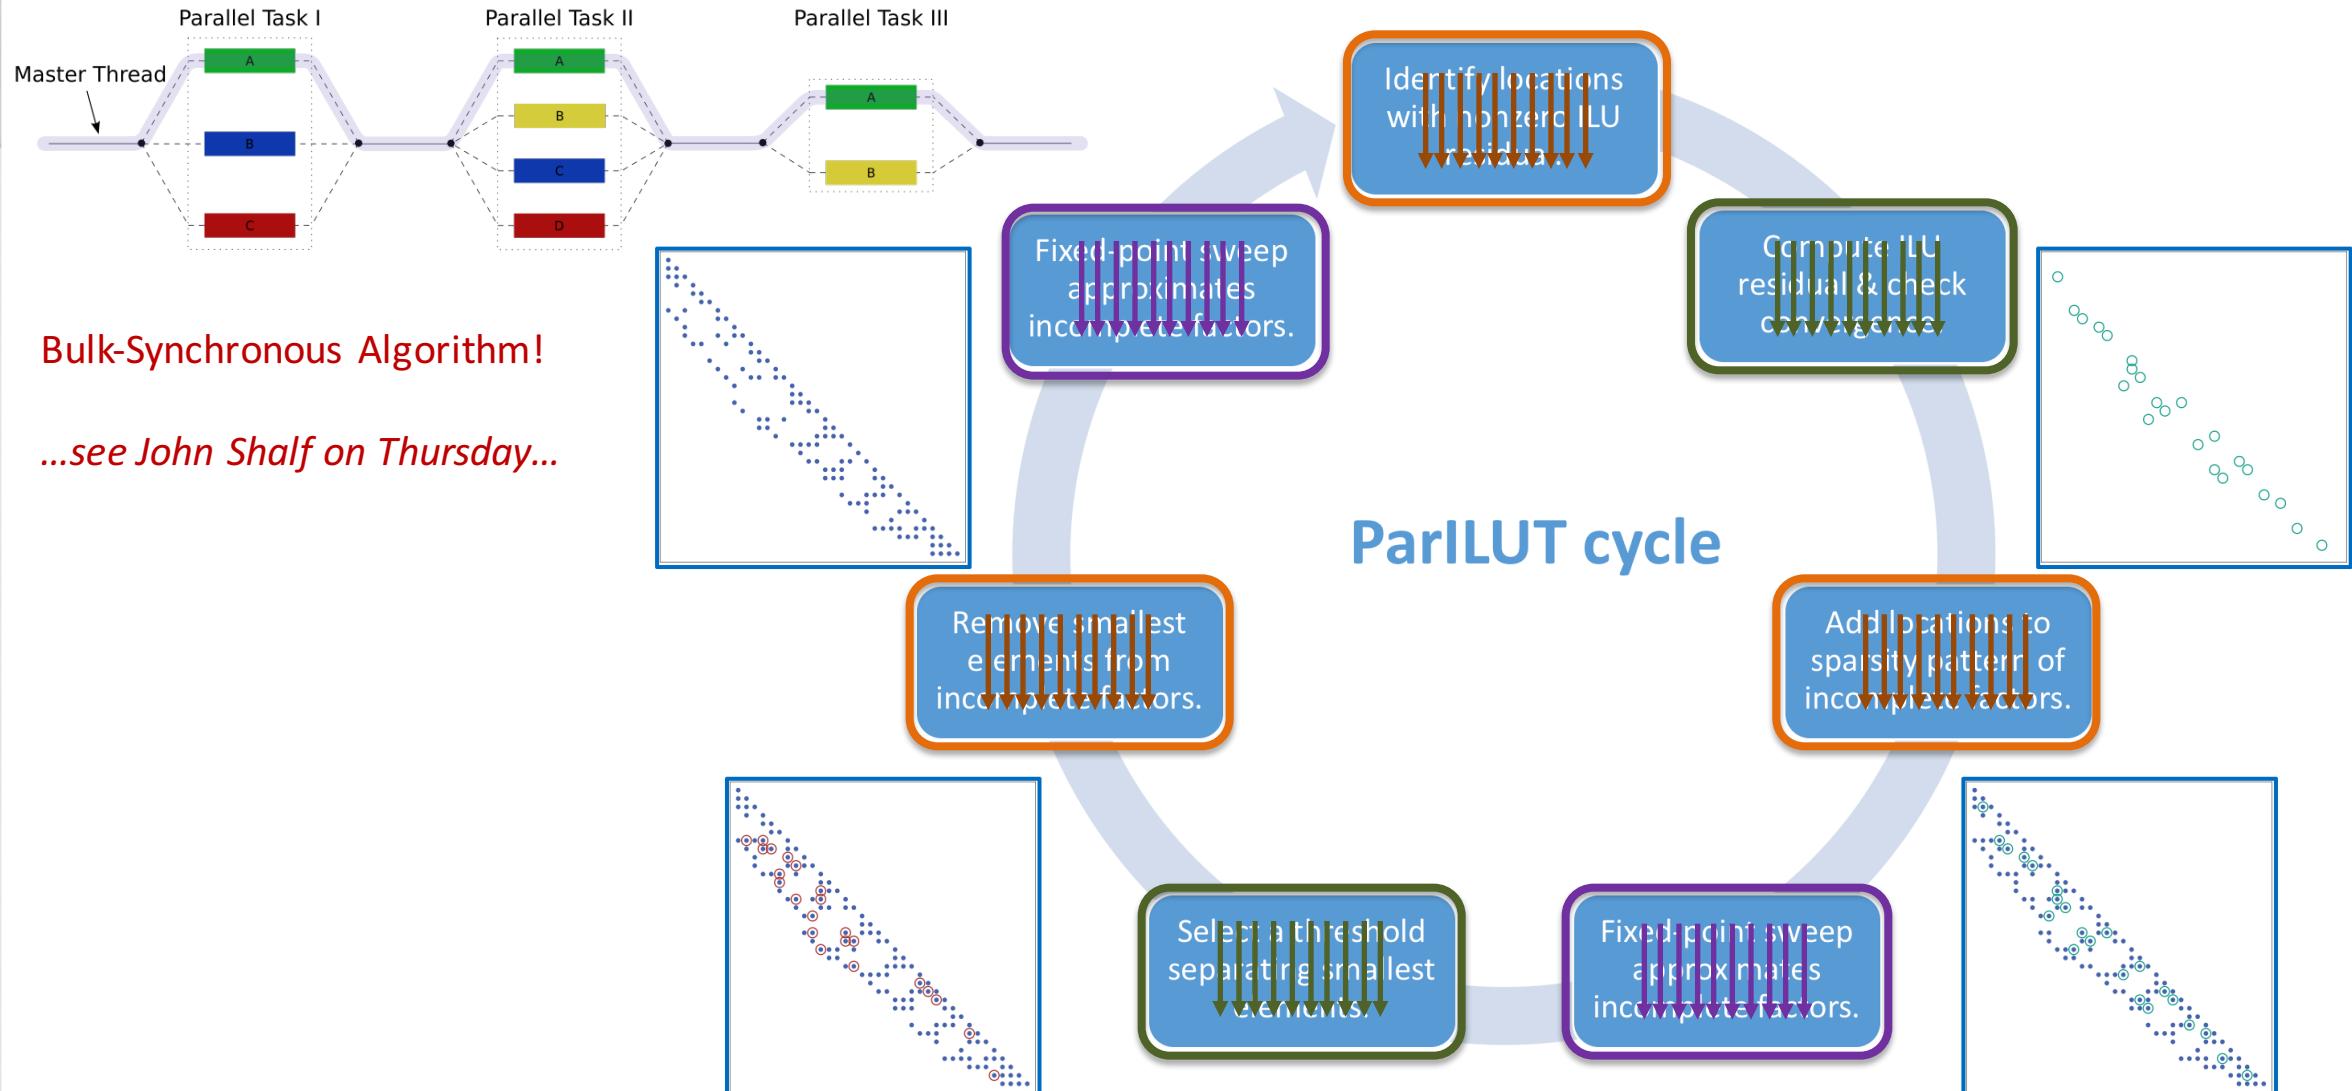
Is this a future-oriented algorithm?

Parallelism inside the building blocks.

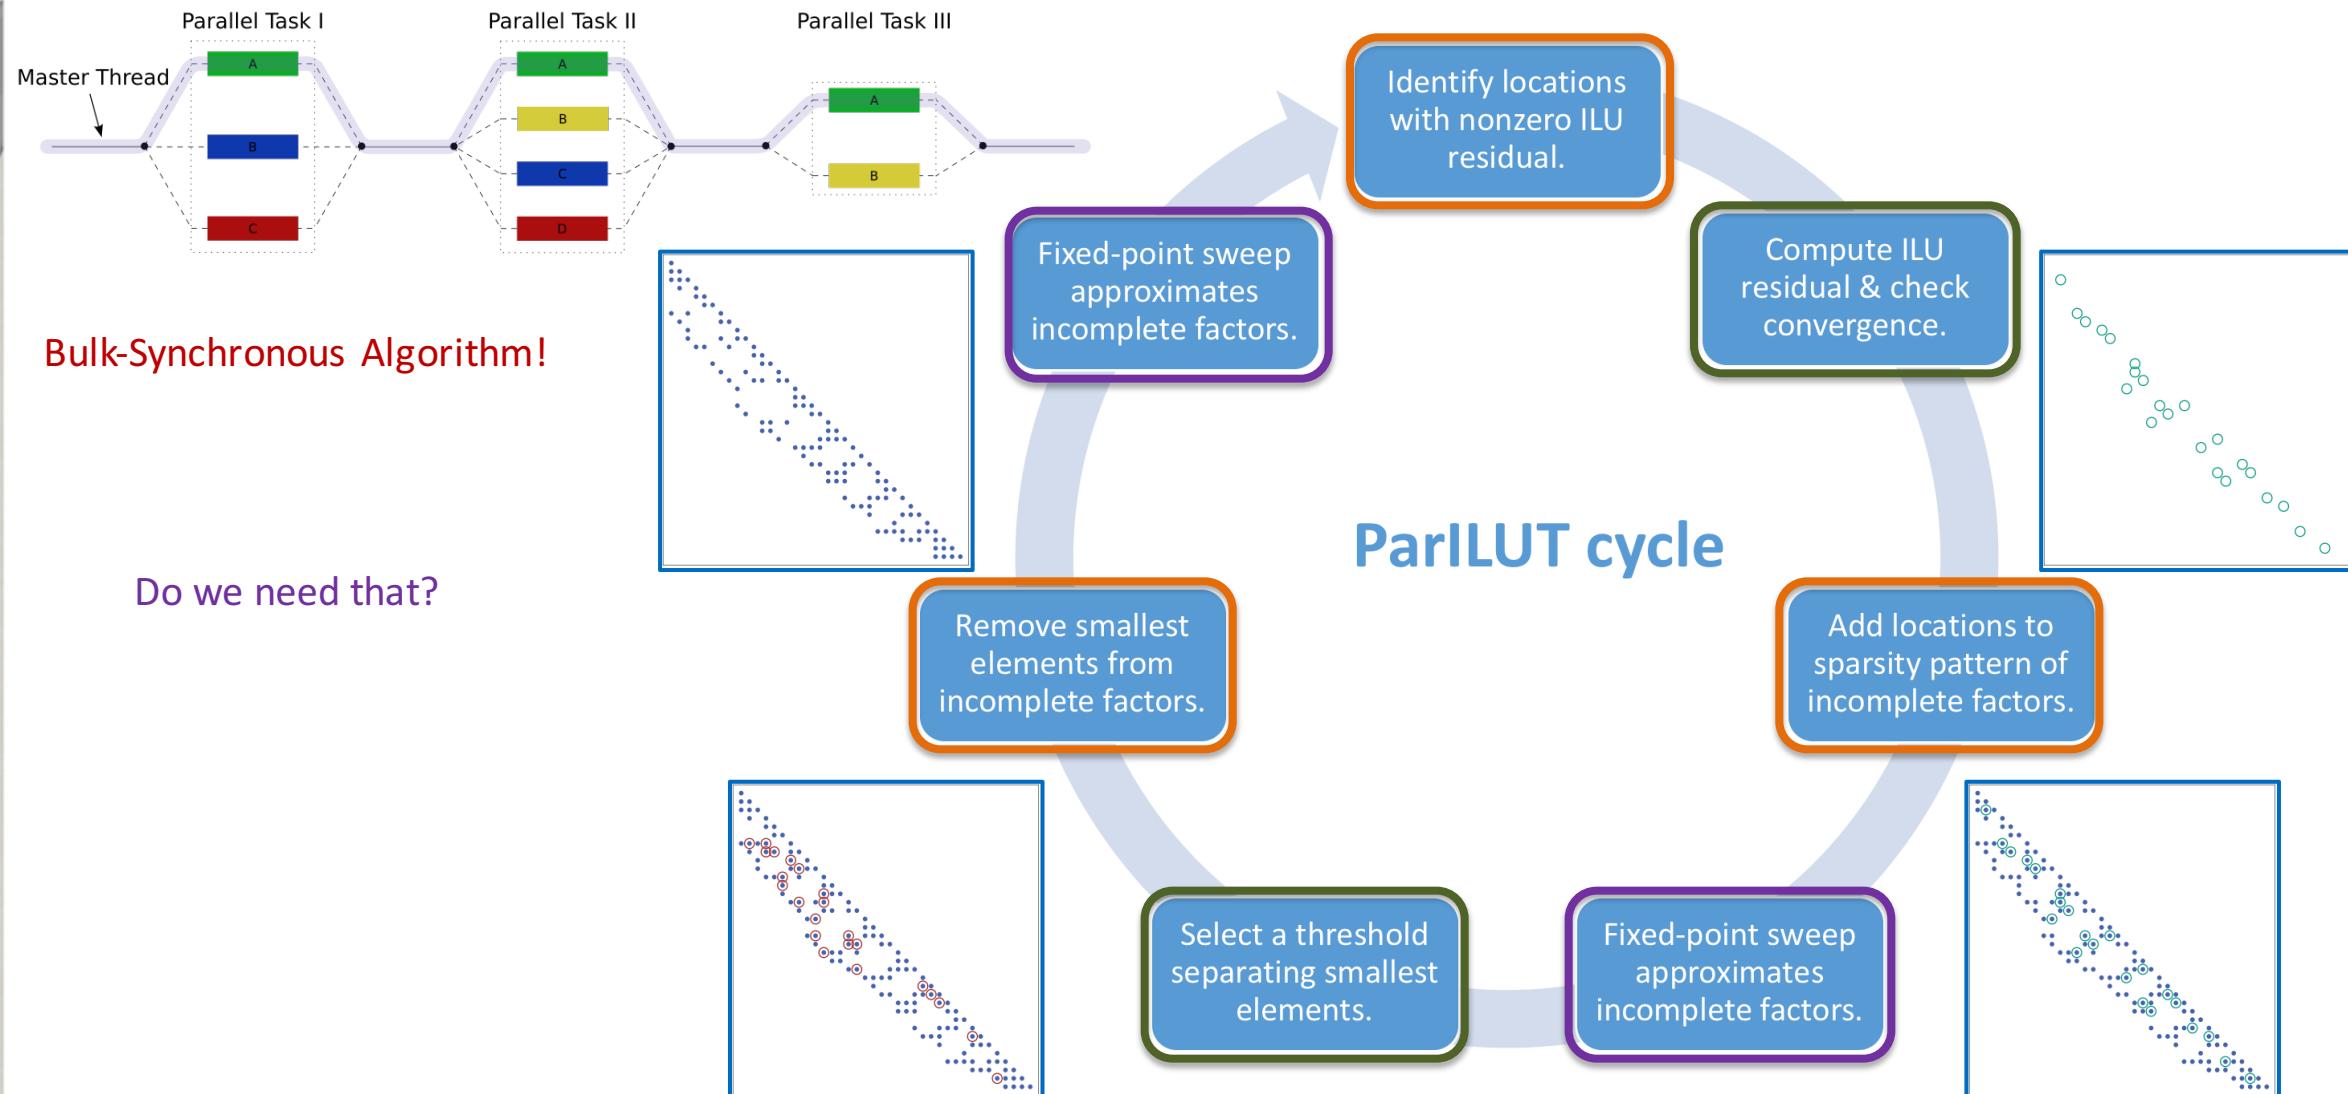
Interleaving fixed-point sweeps approximating values with pattern-changing symbolic routines.



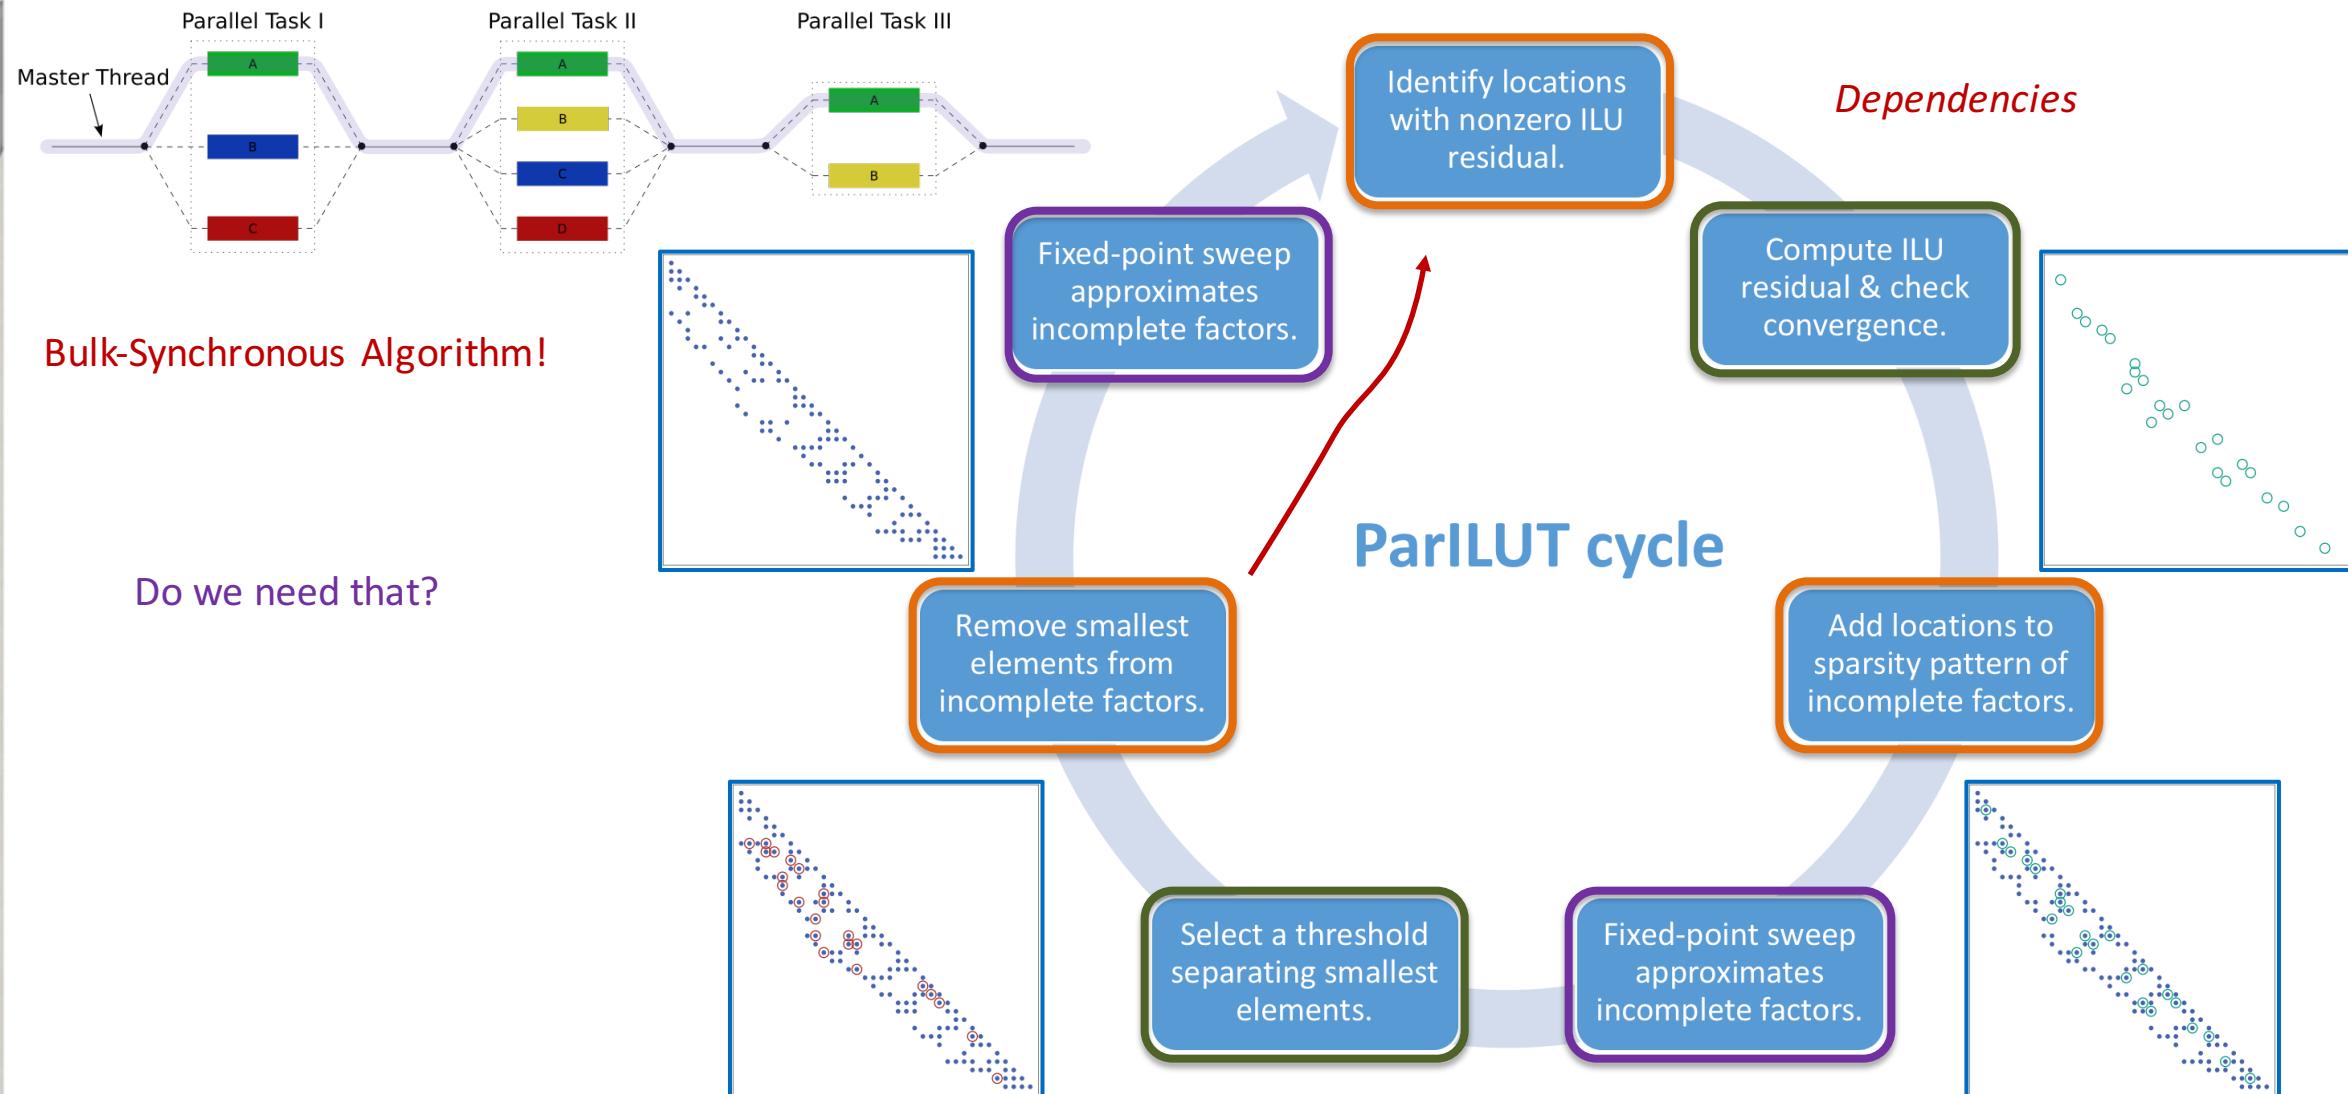
Is this a future-oriented algorithm?



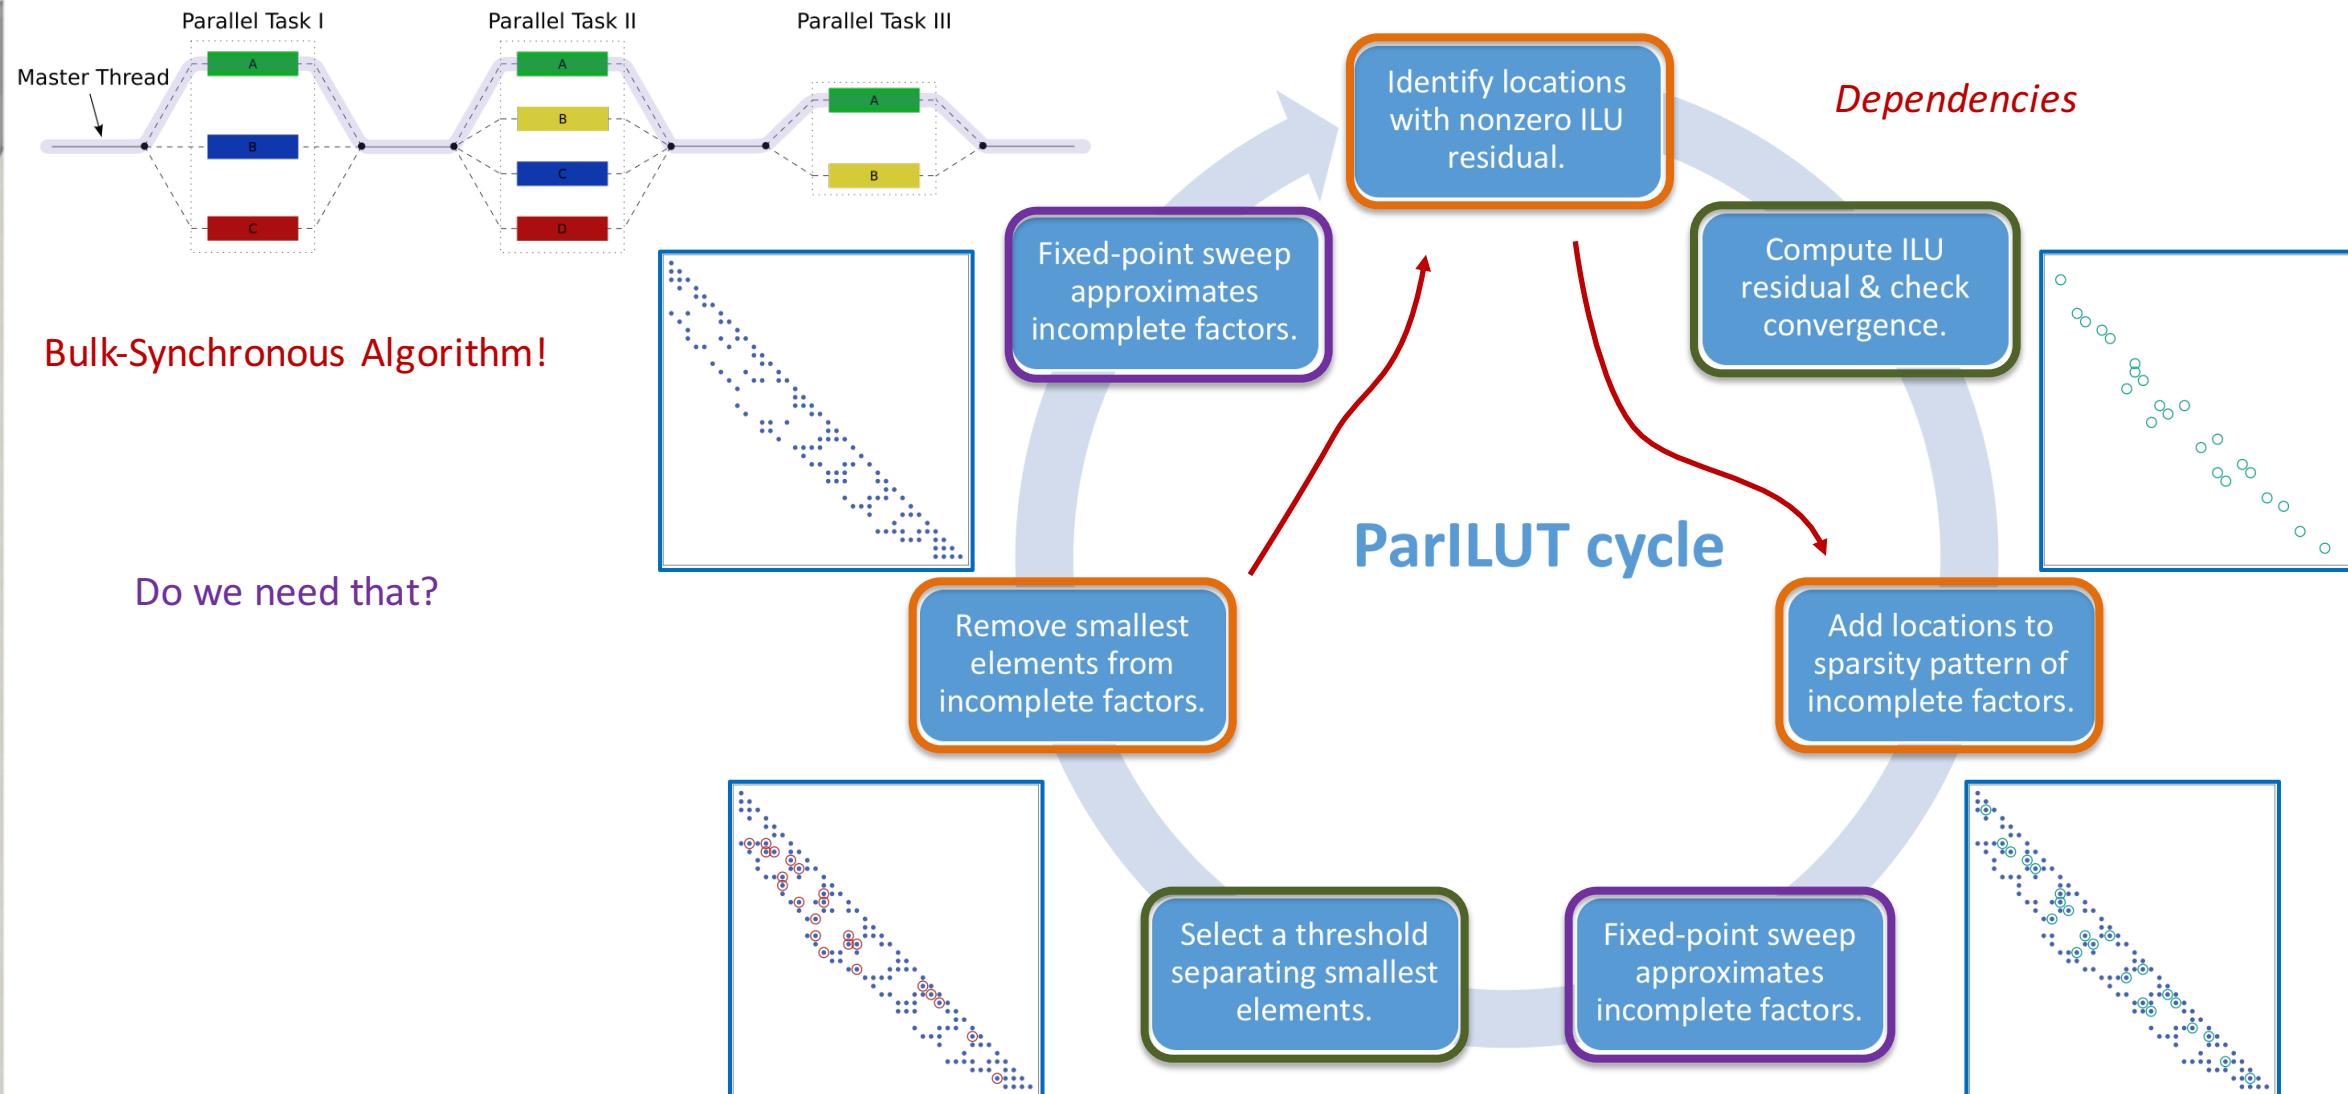
Is this a future-oriented algorithm?



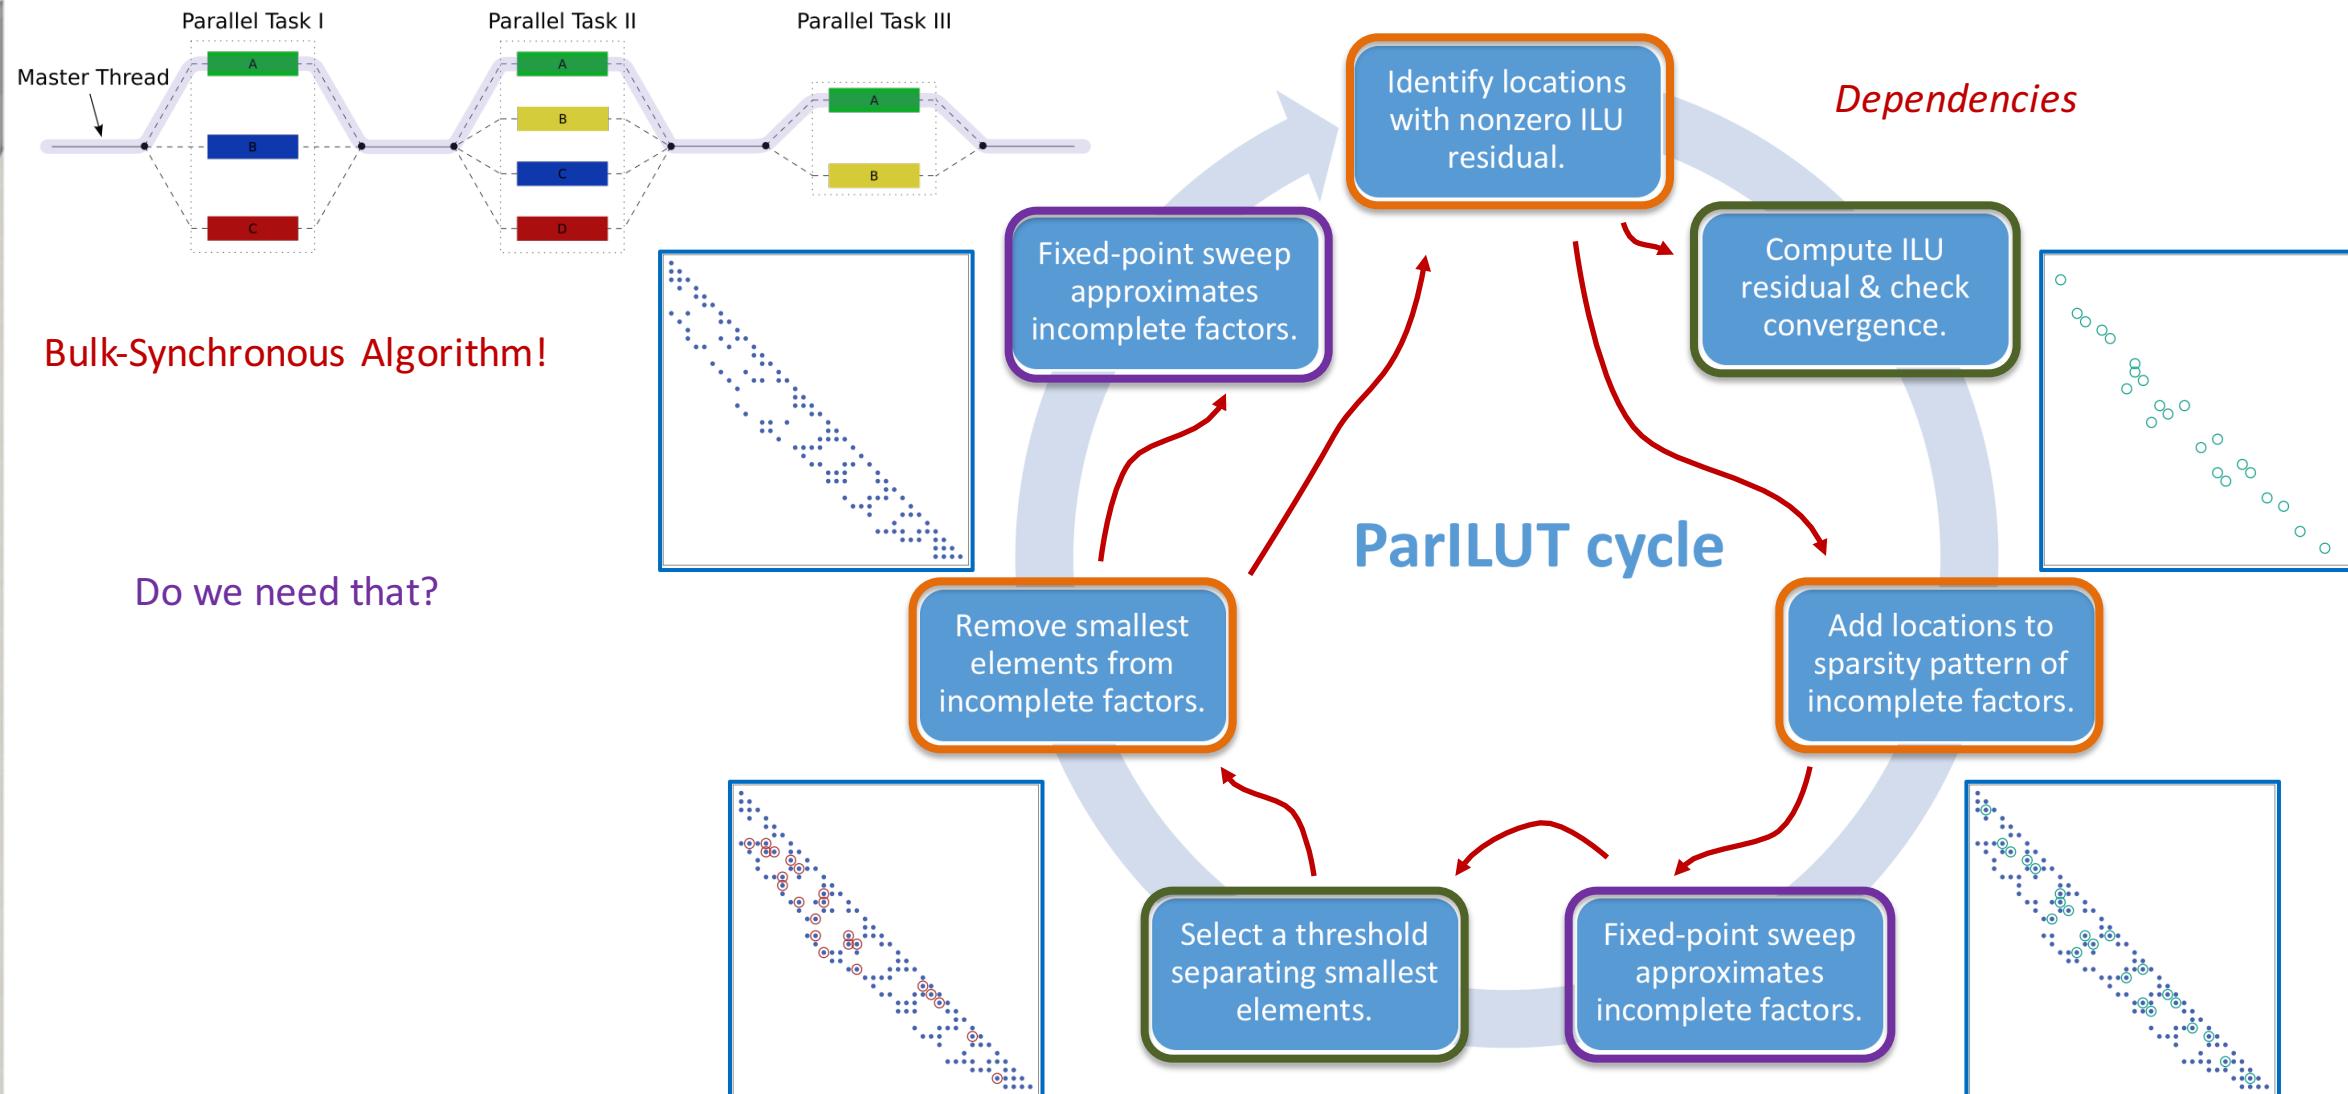
Is this a future-oriented algorithm?



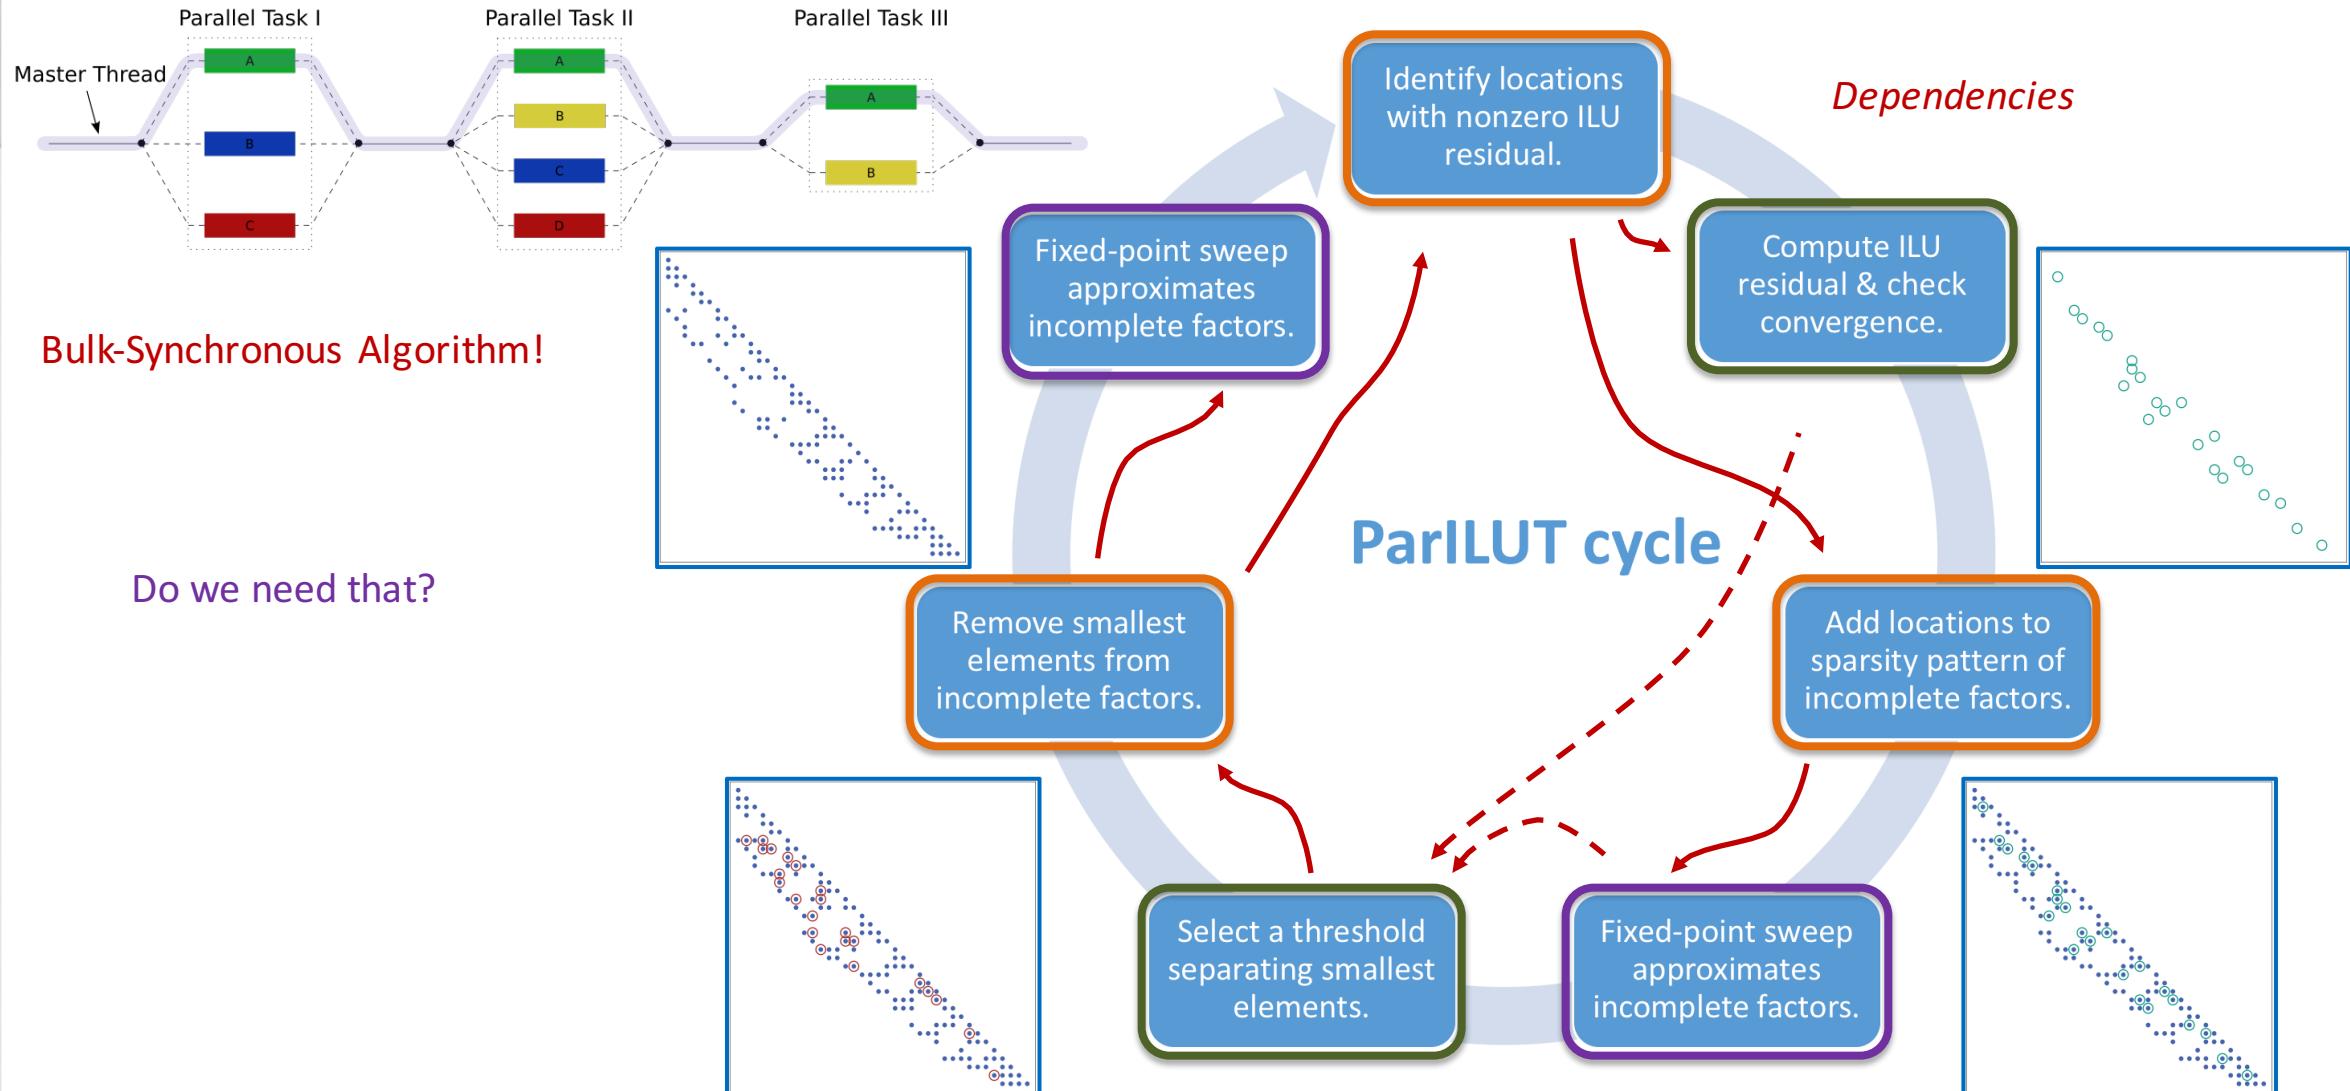
Is this a future-oriented algorithm?



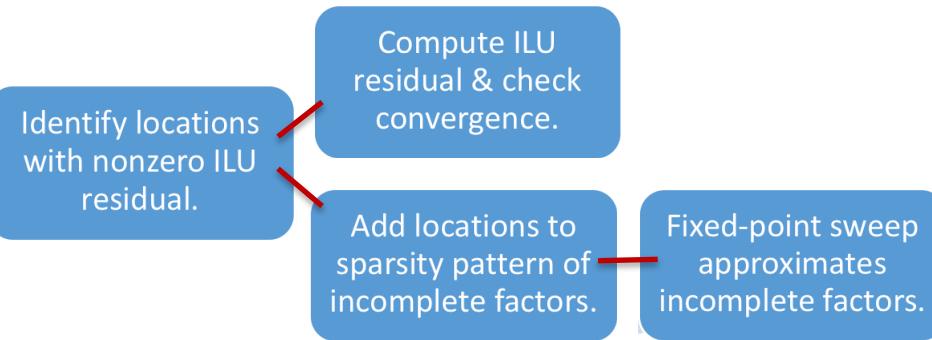
Is this a future-oriented algorithm?



Is this a future-oriented algorithm?



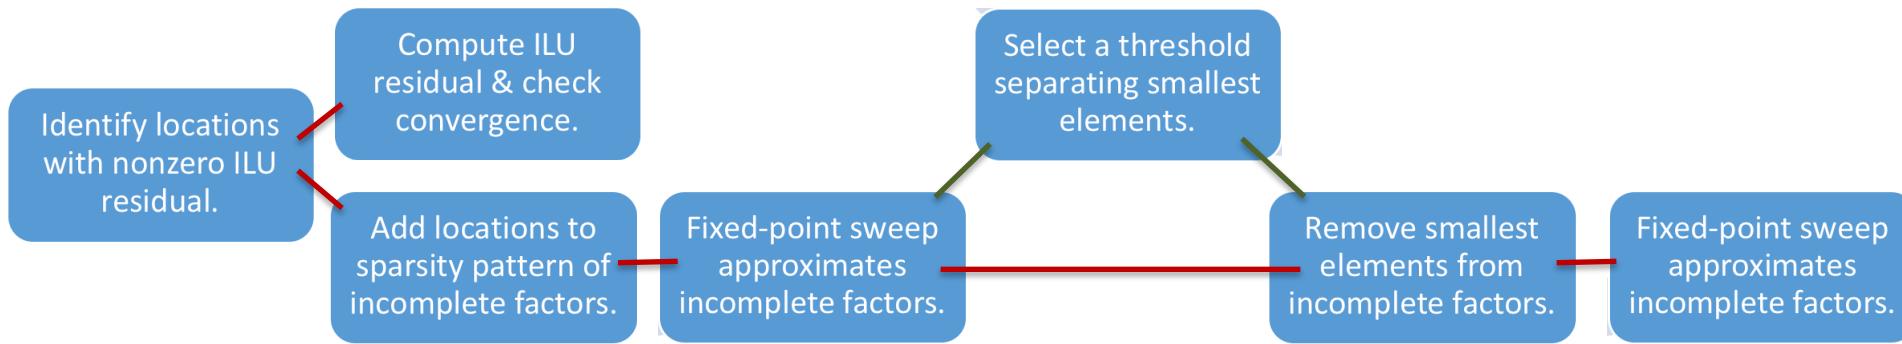
Is this a future-oriented algorithm?



Strong dependency – we can not start before finished.

Weak dependency – if we start before: +/- few nonzeros.

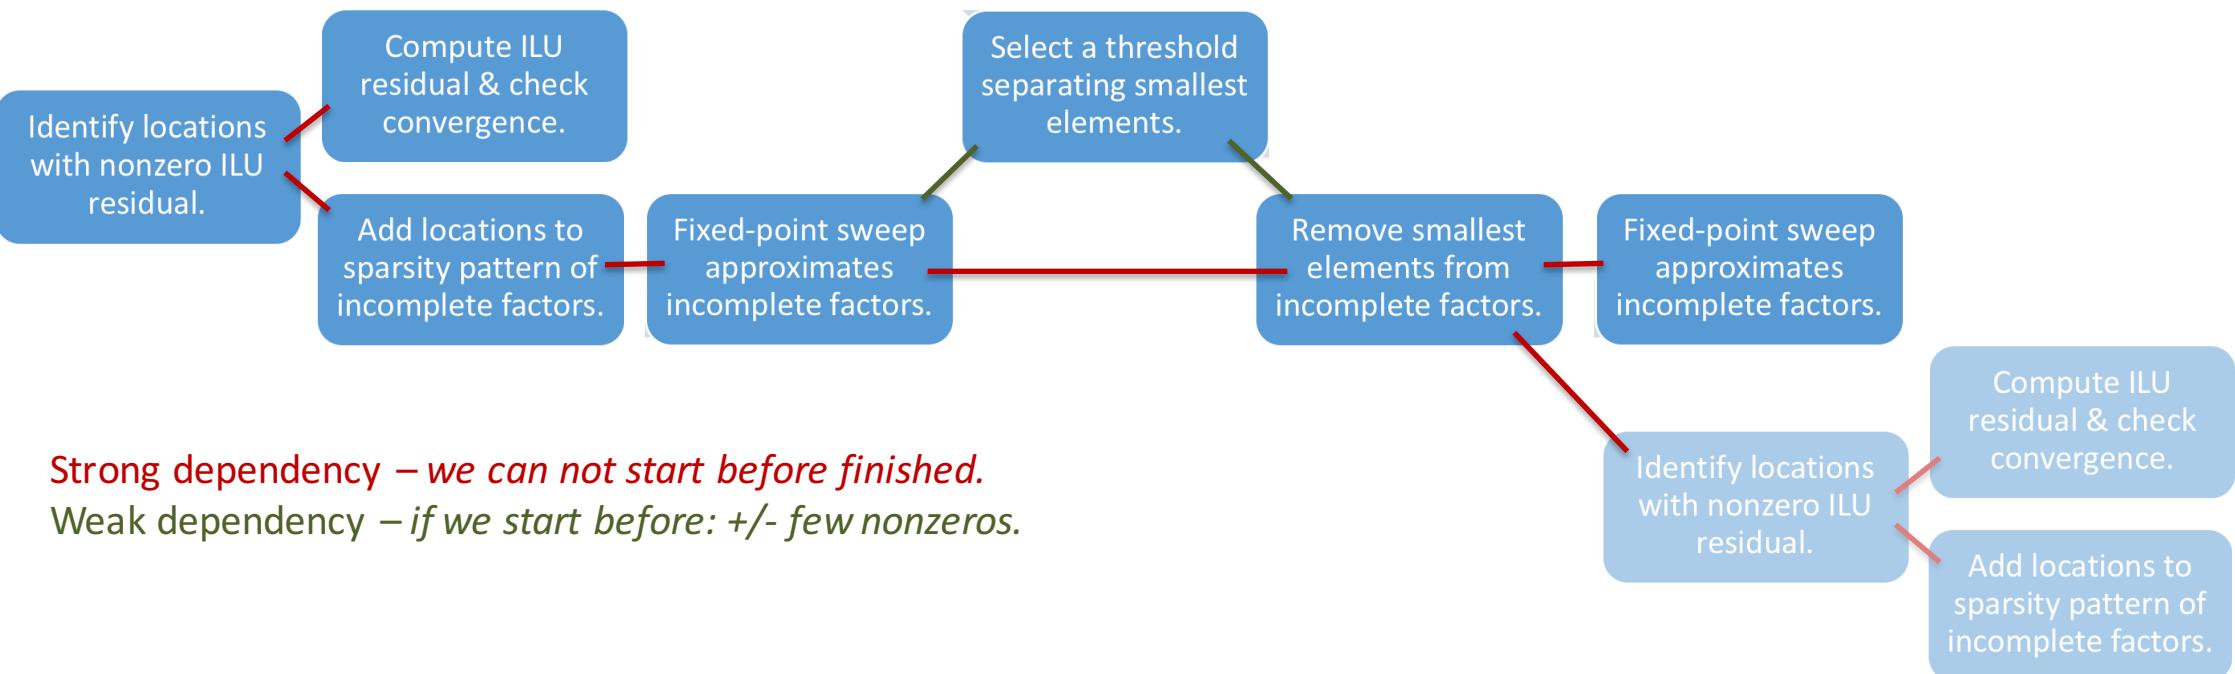
Is this a future-oriented algorithm?



Strong dependency – we *can not start before finished*.

Weak dependency – if we start before: +/- few nonzeros.

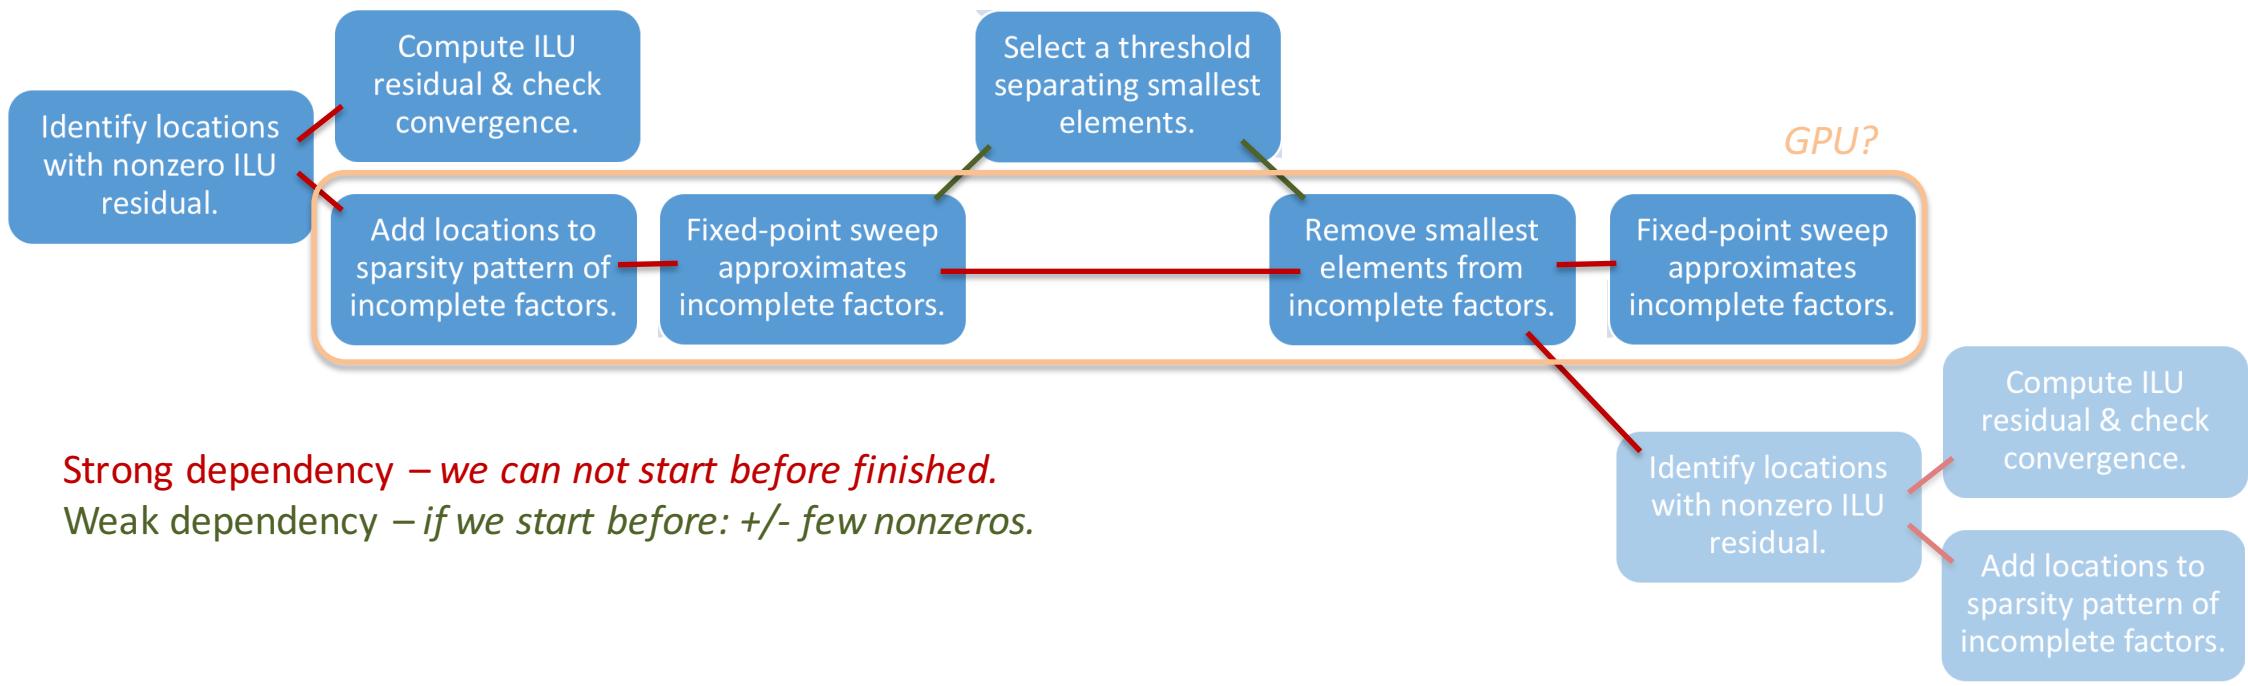
Is this a future-oriented algorithm?



Strong dependency – we can not start before finished.

Weak dependency – if we start before: +/- few nonzeros.

Is this a future-oriented algorithm?



Strong dependency – we can not start before finished.

Weak dependency – if we start before: +/- few nonzeros.

Excellent candidate for hybrid hardware?

Asynchronous execution?